

Graph Databases

- Nodes/Vertices → entities
- edges → relationships
- properties are key-value pairs associated with a specific node or relationship
- ER vs Graph
 - ↳ nodes = entity set
 - ↳ edges = relationship sets
 - each row is a node/edge
- lots of joins ⇒ use a graph DB
- Schema flexibility: new entity types, properties, relationships ... can emerge w/o significant impact on existing data model
 - ↳ the model does not need to be completely developed ahead
- can accommodate the concept of relationship between entities more efficiently than a relational model
- Cypher - specifies what data to retrieve, not how
 - ↳ "select":

```
MATCH(e:Employee)
RETURN e
```

 - match any node node of type employee, return the entire node
 - ↳ "Project":

```
MATCH(e:Employee)
RETURN e.email, e.phone
```

 - returns only specific fields
 - ↳ ordering output:

```
MATCH(e:Employee)
RETURN e.email, e.phone
ORDER BY e.email
```
- "Select" with a condition:
 - find emp info of janet:

```
WHERE e.ename='Janet'
RETURN e
```
 - Search for multiple employees at once:

```
MATCH(e:Employee)
WHERE e.ename in ['Janet','Steve']
RETURN e
```
 - pattern match names starting with s:

```
WHERE e.ename =~ 'S*'
```
 - employees w/o a department:

```
MATCH(e:Employee)
WHERE NOT (e)-[:works_in]-()
RETURN e
```
 - employees not mentoring anyone:

```
MATCH(e:Employee)
WHERE NOT (e)-[:mentors]-()
RETURN e
```
 - employees not mentored by anyone:

```
MATCH(e:Employee)
WHERE NOT (e)<-[:mentors]-()
RETURN e
```
 - employees not mentoring or mentored:

```
MATCH(e:Employee)
WHERE NOT (e)-[:mentors]-()
RETURN e
```
- NULL: a non-existent property is treated as null:


```
MATCH(e:Employee)
WHERE e.job is NULL
RETURN e
```
- "Insert": - CREATE(d:Department{dname:'PR',did=12})
 - <-[works_in]-(e:Employee{ename:'jane'})
 - or: CREATE (e:Employee{ename:'jane'})-[works_in]->(d:Department{dname:'PR',did=12})

Cypher

- in cypher, any query can return data
- ↳ "Insert": add a new relationship between existing nodes:


```
MATCH(n1:Emp{eid='101}),(n2:Emp{eid:201})
CREATE (n1)-[:manages]->(n2);
```
- ↳ Delete: delete relationships before/while deleting the node


```
MATCH(e:Emp{empid:201})-[r:WORKS_IN]->(d:Department{deptid:12}) DELETE e,r
```

 - or:

```
MATCH(e:Employee{empid:201})
DETACH DELETE e
```

 - ↳ deletes edges as well as node
- ↳ "Joins"/traversals:
 - find employees working in HR:


```
MATCH(e:Employee)-[w:WORKS_IN]->(d:Department{dname:"HR"}) RETURN e
```

 - ↳ or:

```
MATCH(e:Employee)-[w:WORKS_IN]->(d:Department) WHERE d.dname = "HR"
RETURN e
```
 - return info on paul + nodes he has relationships with:


```
MATCH(e:Emp)--(n) WHERE e.ename = 'Paul' RETURN e,n
```

 - match(e:Emp)-[]-(n) WHERE e.ename = 'Paul' RETURN e,n
 - return info on paul + nodes of outgoing relationships:


```
match(e:Emp)-[r]->(n) WHERE e.name='Paul' RETURN e,r,n
```

 - return names of all employees steve is managing:


```
MATCH(e:Employee)-[:MANAGES]->(n:Employee)
WHERE e.ename = 'Steve'
RETURN n.ename
```
 - return info of people who are managed or mentored by katie:


```
MATCH(e:Employee)-[:MANAGES|MENTORS]->(n)
WHERE e.ename = 'Katie' RETURN n
```
 - return info of people neither managed nor mentored by katie:


```
MATCH(e:Employee), (n:Employee)
WHERE e.ename = 'Katie' AND NOT (e)-[:MANAGES|MENTORS]->(n) RETURN n;
```
 - return info of people who report to managers managed by steve
 - ↳ MATCH(e:Emp)-[:manages]->()-[:manages]->(n)
 - WHERE e.ename = 'Steve'
 - RETURN n
 - ↳ Or:

```
MATCH(e:Employee)-[:MANAGES*2]->(n)
WHERE e.ename = 'Steve'
RETURN n
```
 - multi-depth traversals:
 - ↳ all the way: $(e)-[*]->(n)$
 - ↳ up to depth of 5: $(e)-[*..5]->(n)$
 - ↳ 3+ edges: $(e)-[*3..]->(n)$
 - ↳ 3-5 edges: $(e)-[*3..5]->(n)$
 - ↳ 3-5 edges: $(e)<-[*3..5]->(n)$
 - ↳ 3-5 edges: $(e)-[*3..5]->(n)$
 - employees under steve who are not managers:


```
MATCH(e:Emp)-[:manages]->(e2:Emp)
WHERE e.name = 'Steve' AND NOT (e2)-[:manages]->()
RETURN e2;
```



```

MATCH(e:Employee) -[:MANAGES]->() -[:MENTORS]->(n)
WHERE e.ename = 'Steve'
RETURN n

```

```

MATCH(n:Employee)
WHERE NOT (:Employee{ename:'Steve'})
-[:MANAGES]->()-[:MENTORS]->(n) RETURN n
  MATCH(e:Employee), (n:Employee)
WHERE e.ename = 'Steve' AND NOT
(e)-[:MANAGES]->()-[:MENTORS]->(n) RETURN n;

(Janet)-[:FRIEND_OF]->(Sue)
(Sue)-[:FRIEND_OF]->(Janet)
MATCH (e1:Employee(ename:'Janet'))
-[:FRIEND_OF*]->(e2:Employee) RETURN e1,e2

```

```

MATCH
(n:Employee{ename:'Janet'})-[:MANAGES]->
(e1:Employee)-[:MANAGES]->(e2:Employee),
(n)-[:FRIEND_OF]->(e2) RETURN n,e1, e2

```

```

      MATCH (b:Employee)-[:MANAGES]->
(m:Employee),(m)-[:MENTORS]->
(e1:Employee),(m)-[:MENTORS]->
(e2:Employee)
WHERE e1 <> e2 RETURN DISTINCT b

```

```

CREATE CONSTRAINT ON (e: Employee)
ASSERT e.eid IS UNIQUE
CREATE CONSTRAINT ON (e: Employee)
ASSERT exists(e.ename)
CREATE CONSTRAINT ON ()-[m:MENTORS]-()
ASSERT exists(m.skill)

```

All the people who has a potential risk for Huntington but not Lebers:

```

MATCH(s:Subject)<-[:motherof|fatherof *]-(:Subject)-[:has]->(d:Disorder {id:1}),(d2:Disorder {id:2})
WHERE NOT(s)<-[:mohter_of*]-(:Subject)-[:has]->(d2)
RETURN s

```

Create Siblings:

```

MATCH(s1:subject)<-[:fatherof]-(f1:subject), (s1)<-
[:motherof]-(m1:Subject),
(s2:Subject)<-[:fatherof]-(f1), (s2)<-[:motherof]-(m1)
WHERE S1.id < S2.id
CREATE (s1)-[:Sibling_of]->(s2)

```

friends of Sally who went to the same school as her:

```

MATCH(p:Person {name:'Sally'})-[:STUDIED_AT]->
(s:School),(p2:Person)-[:STUDIED_AT]->(s)
,(P)-[:FRIEND_OF]-(P2)
RETURN p2;

```

People who went to the same school and graduated the same year, but are not immediate friends:

```

MATCH(p1:Person)-[s1:STUDIED_AT]->(s:School)
,(p2:Person )-[s2:STUDIED_AT]->(s)
WHERE s1.grad_year = s2.grad_year AND NOT (p1)-[:FRIEND_OF]-(p2)
Return p1;

```

Movies for John that he has not yet liked but is of the same Genre as the other movies that he has liked

```

MATCH(p:Person {name:'John'})-[:LIKED]-(m:Movie)-[:BELONGS_TO]->(g:Genre)
,(m2: Movie)-[:BELONGS_TO]->(g)
WHERE NOT (p)-[:LIKES]->(m2)
RETURN m2;

```

Create fanclubs:

```

MATCH(b:Band)
WHERE b.name in ['Eagles', 'Enigma']
CREATE (f1:Fanclub {city:'Montreal'})-[:ASSOCIATED_TO]->(b)
,(f2:Fanclub {city:'Ottawa'})-[:ASSOCIATED_TO]->(b);
MATCH(p:Person)-[:LIKES]->(b:Band), (f:Fanclub)-[:ASSOCIATED_TO]->(b)
WHERE p.city = f.city
CREATE (p)-[:MEMBER_OF]->(f);

```

Assignment 3

- A map-reduce workflow that takes a person and their friends, $(P, (F_1, F_2, \dots))$ and gives common friends for each set of friends.

mapper: input $(P, \text{list of friends})$

For each friend F in list of friends

```
if ( $P \in F$ ) output  $((P, F), (\text{list of friends}) - F)$ 
else output  $((F, P), (\text{list of friends}) - F)$ 
```

reducer: input is a pair of friends and corresponding list of friends: $((P_1, P_2), ((\text{list 1}), (\text{list 2})))$

output $((P_1, P_2), \text{common friends intersection})$

- A map-reduce workflow that produces provinces w/ over 100 patients over 60 who are in the hospital and the total number of patients. [Hospital (Hname, province)]

Patient (Hname, age, Hname) Hname ref hospital]

mapper: For each tuple (Hname, age, Hname) of Patient
if (age > 60) output $((Hname, p))$

For each tuple (Hname, province) of Hospital
output $((Hname, (H, \text{province}))$

reducer: input $((Hname, (p, \dots, p, (H, \text{province})))$

numpatients = # P

output $((\text{Province from } H \text{ tuple}, \text{num patients}))$

mapper 2: input $((\text{Province}, \text{num patients}))$

output identity: $((\text{Province}, \text{num patients}))$

reducer 2: input $((\text{Province}, (\text{num patients 1}, \dots, \text{num patients n}))$

total patients = sum all numpatients

if (total patients > 100)

output $((\text{Province}, \text{total patients}))$

- P6: 10 movies w/ max # user ratings

Ratingsgrp = GROUP ratings BY movieid;

ratingspermovie = FOREACH ratingsgrp GENERATE group AS movieid, COUNT(\$1) AS numratings;

Ratingspermovieord = ORDER ratingspermovie

BY numratings DESC;

top10movies = LIMIT ratingspermovieord 10;

top10moviesinfo = JOIN top10movies BY movieid,
movies BY movieid;

Topmovies = FOREACH top10moviesinfo GENERATE title

AS title, numratings AS numratings;

Topmoviesord = ORDER topmovies BY numratings DESC;
DUMP topmoviesord;

movies per genre 2015 - 2016

Movies20156 = FILTER movies BY year == 2015 OR
year == 2016;

Moviegenres20156 = JOIN movies20156 BY movieid,
Moviegenres BY movieid;

Geresandmovies = GROUP moviegenres20156 BY
(genre, year);

Genremoviecount = FOREACH geresandmovies
GENERATE FLATTEN(group), COUNT(\$1)
AS nummovies;

Orderedgenres = ORDER genremoviecount BY genre, year;
DUMP orderedgenres

- Years in which # movies were less than the previous yr
Moviesperyear = GROUP movies BY year
Yearcount = FOREACH moviesperyear GENERATE
Group AS year, COUNT(\$1) AS nummovies;
Yearcount2 = FOREACH yearcount GENERATE
Year+1 AS curyear, nummovies;
Joinyears = JOIN yearcount BY year, yearcount2
BY curyear;
Moviecounts = FOREACH joinyears GENERATE
Yearcount2::curyear, yearcount::nummovies
AS currnummovies, yearcount2::nummovies
AS prevnummovies;
Badyears = FILTER moviecounts
BY currnummovies < prevnummovies;
DUMP badyears;
Orderbadyears = ORDER badyears BY curyear;
DUMP orderbadyears;

Quiz 5 - Trans. and CC.

It is safe for a program to crash without leaving inconsistent data in the DB if there are no changes to the DB or all changes are complete.

Crashes and aborts should not be present in a system after recovery.

Before the DB sends a commit confirmation to the app it must record the commit from the transaction in its log, flush the log buffers (log tail) to the disk, and the app must initiate the transaction commit.

Data buffers are flushed when the memory buffer pool needs to make room for new pages.

A transaction T updates X then aborts, then the DBS crashes:

↳ It is possible in some situations that it is not necessary for the DBS during the recovery to apply the before image of X.

↳ writes made by T may still be in memory buffers, not disk
so it's like the update never happened which is good

T updates X and commits, then the DBS crashes:

↳ It is possible in some situations that it is not necessary for the DBS during recovery to apply the after image for X

↳ the changes may already be updated in disk so no redo would be necessary.

Conflict equivalent schedules:

w2(b), r1(a), w1(a), r3(a), r2(a), r1(b), r3(a), c3, c2, c1
r1(a), w1(a), w2(b), r3(a), r2(a), r1(b), r3(a), c1, c2, c3
r1(a), w1(a), r3(a), r3(a), c3, w2(b), r1(b), r2(a), c1, c2, c3

Conflict equivalent serializable schedule

r3(a), w2(b), r1(b), w2(c), r4(c), r2(a), r1(b), w3(b), r4(b)
, c2, c1, c4, c3 => T2, T1, T3, T4 (flatten conflict graph)

Submission order → ZPL Schedule

r2(a), w1(b), r1(c), r3(c), w2(c), c3, w2(b), c2, r1(a), c1
(r2(a), w1(b), r1(c), r3(c), c3, r1(a), c1, w2(c), w2(b), c2)

Continue next page...

Quiz 5 - cont.

T1	T2	T3
1 A=read1(A)	1	1
2 A=A+100	2	2
3	3	3 C=read3(C)
4	4	4 C=C-300
C=read1(C)	5	5
6 A=read2(A)	6	6 write3(C)
7	7	7 commit3
8	8	9
C=C-100	9	10
10 write1(A)	10	11
11 write1(C)	11	12
12	12	13 B=read2(B)
B=B-200	13	14
14 write2(B)	14	15
15 commit1	15	16 A=read2(A)
16 A=read2(C)	16	17
17 C=read2(C)	17	18 A=A+200
18	18	19 write2(A)
19	19	20 commit2
20	20	

↳ lost update: write3(c)
 un-repeatable read: read2(A)
 no dirty reads
 no dirty writes

Practice final

- we can enforce the constraint that rating can only be null or 1-10.
- If we create an index on psid of parent, a type II indirect index will not necessarily take less space than type I.
- a search for a sitter by a range on hourly rate and a \geq condition on minPreferredAge would not benefit from clustering the table on minPreferredAge.
- The efficiency of a page nested loop join cannot be increased by adding more buffer frames (assuming the total memory buffer frames are still less than the total number of pages that either relation has).
- The height of a B+ tree index is dependent on the size of the data entries.
- Adding an extra attribute that is not part of the search key of an index will not reduce the number of page pointer entries that can fit into an intermediate node of an index structure.
- In a graph database, a node/vertex is similar to an entity in ER SQL:
- Name + # of sitters who have an entry in sitterCalendar for May - August 2017:

```
SELECT sname, sphone
FROM sitters
WHERE ssid IN (SELECT ssid FROM sitterCalendar WHERE
               availDate BETWEEN '2017-05-01' AND '2017-08-31')
```

• SSid, sname, sphone of sitters w/ at least 5 ratings > 7:

```
SELECT ssid, sname, sphone
FROM sitters s, babySitting b
WHERE s ssid = b ssid AND b rating > 7
GROUP BY s ssid, s name, s phone
HAVING COUNT(*) >= 5
```

• psid of parents w/ more than one child, at least one is female:

```
SELECT psid FROM child WHERE psid IN
              (SELECT psid FROM child WHERE cgender='female')
GROUP BY psid HAVING COUNT(*) > 1
```

Cypher:

- find preferred sitters of people in melanies network

```
MATCH(p1:Parent {name='melanie'}) -[:knows*]->
      (p:parent), (p)-[:prefers]->(b:babysitter)
RETURN b
```

- add a knows rel between parents and preferred sitters

```
MATCH(p:parent)-[:prefers]->(b:babysitter)
CREATE (p)-[:knows]->(b), (b)-[:knows]->(p)
```

P16:

- given a psid, get info on female sitters who've at least 10x
 femaleSitters = filter sitter by gender = 'female';
 parentsBabySitting = filter babySitting by psid = 123456789;
 jnd = join femaleSitters by ssid, parentsBabySitting by ssid;
 grpD = group jnd by (ssid, name, phone);
 smmd = foreach grpD generate(\$0), COUNT(\$1) as
 bookingCount;
 fltrsMMD = filter smmd by bookingCount \geq 10;
 srtd = order fltrsMMD by name;
 dump srtd;

• Map Reduce:

- count of # pages referencing a url
 Map: for each tuple (referencingURL, referencedURL_LIST)
 for each referencedURL in referencedURL_LIST
 If referencedURL != referencingURL
 Output (referencedURL, referencingURL)

Group and shuffle will take in each (referencedURL,
 referencingURL) and output tuples (referencedURL,
 referencingURL_LIST)

Reduce: for each tuple(referencingURL,referencingURL_LIST)
 Aref= {}, counter = 0

For each referencingURL in referencingURL_LIST
 If referencingURL not in aref
 Add referencing url to aref

Counter++

Output (referencedURL, counter)

- produce weighted output of previous result

Map: For each tuple(referencedURL, referenceCount)

 output(referencedURL, ('RefCount',referenceCount))

 For each tuple(URL, popValue) of popularity

 output (URL, ('popularity', popValue))

Reduce: For each tuple (URL, tupleList)

 refTuple = NULL

 popTuple = NULL

 For each tuple t in tupleList

 If t.rel == 'RefCount'

 refTuple = t;

 Else If t.rel == 'popularity'

 popTuple = t;

 If popTuple == NULL

 output(URL, refTuple.referenceCount)

 Else output(URL, refTuple.referenceCount *
 popTuple.popValue)

Practice final cont

Query evaluation:

```
sitter(ssid:INT ,sname:VARCHAR(40) ,sphone:CHAR(12) ,sgender:CHAR(6)
, hourlyRate:INT ,preferredMinAge:INT)
parent((psid):INT , pname:VARCHAR(40), pphone:CHAR(12))
babySitting(reservationId:INT, ssid:INT ,availDate:DATE, psid:INT
,cname:VARCHAR(40) ,starthr:INT, totalhrs:INT, addnlChildren:INT, rating:INT)
--> (ssid , availdate) references sitterCalendar
--> (psid , cname) references child
```

INT takes 4 Bytes, FLOAT takes 8 bytes, all VARCHAR fields take half of the size specified. DATE field is 10 Bytes
All pages are 4K size (you can use 4000 for calculations)

sitter has 1,000 records spread across 20 data pages. 20% of the babysitters are *male*
parent has 20,000 records.
babySitting has 1,000,000 records spread across 20,000 pages. availDate has 10,000 distinct values. totalhrs have values in the range 1-20.

You can make the following assumptions. There are indexes on all primary keys (unclustered). Root and intermediate pages of B+-tree indices are always in main memory. No other page is assumed to be in main memory at the begin of a query execution. There are around 10 buffer frames available. (you can assume another couple of frames are available if it helps your computation to be easier)

```
SELECT s.ssid, s.sname,
       SUM(s.hourlyRate*b.totalhrs*(1+0.50*addnlChildren))
FROM sitter s, babySitting b
WHERE s.ssid = b.ssid
      AND b.availDate BETWEEN '2017-03-01' AND '2017-03-05'
GROUP BY s.ssid, s.sname
```

process + I/O costs:

Steps

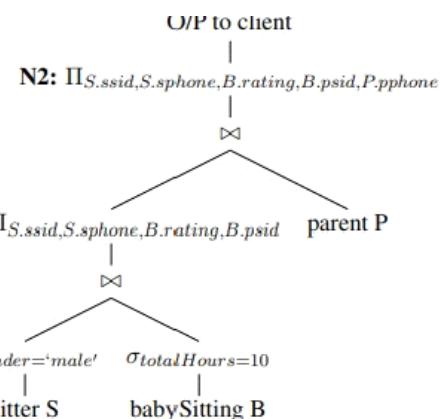
1. First read the 20,000 pages of babySitting, use the selection on availdate and pipeline it to projection and store only the required fields (ssid, totalhrs, addnlChildren) of the selected records in the memory buffers. The width of the output tuple is $4+4+4 = 12$, and can be done into 2 frames \Rightarrow temp1. The number of records in temp1 will be $1,000,000 \times 5/10,000 = 500$. Total IO is 20,000 reads.
2. Block Nested with temp1 as outer and sitter as the inner relation. Temp1 is in memory so join read cost = 20 (for sitter). Compute the amount earned for this babysitting reservation save only required fields (ssid, sname, amt), width = $4+20+8 = 32$... takes about 4 frames (no need to write can be held in memory). Total IO = 20 reads
3. Do an in memory sorting based on ssid, sname. This is then grouped (in memory) - pipelined to sum - and pipeline to project - and send to client directly. Total IO is therefore only $20,000 + 20 = 20,020$

decrease I/O with an index:

Create a clustered index on availDate of babySitting

For selection operator on availDate use this index and , read 1 leaf page + 10 data pages , project and store only required fields in memory (ssid, totalhrs, addnlChildren) (2 frames) \rightarrow temp1 = 11 reads

Rest of the steps same as before - results in $11 + 20 = 31$ I/O



3. N1 processes about 10,000 records (20% of the 50,000 qualified from babySitting), both entering and leaving. Incoming tuple size is either 108B or 104B, leaving tuple size is 24B
- N2 processes about 10,000 records too (because only half of the parents will be needed), both entering and leaving. Incoming tuple size is either 60B or 56B and outgoing tuple size is 36B

Lost update: read-write-write pattern

↳ $r_1 \rightarrow w_2 \rightarrow w_1$
or $r_2 \rightarrow w_1 \rightarrow w_2$

