

# The Relational Model

# Database Models


- The entity-relationship is a semantic model (quite rich in its expressiveness) but it is not used as a model in any DBS. *good as high level for clients*
- **Relational model**
  - Most common
  - Vendors: IBM DB2, Microsoft SQL Server, Oracle, SAP Hana, SAP Sybase, etc...
  - Open-source: PostgreSQL, MySQL, SQLite, Derby, MonetDB,
- “Legacy systems” have older models
  - E.g., IBM’s IMS (hierarchical), CODASYL network model

# Database Models

- Recent and future competitors:
  - object-oriented model: ObjectStore, Versant
  - Semi-structured (document-based): XML, JSON
    - *MongoDB*, *CouchDB*
  - Key-value
    - *Riak*, *Voldemort*, *Amazon's Dynamo*
  - Column-store (key – basic relational)
    - *Google's BigTable*, *Hadoop's Hbase*, *Cassandra*
  - Integrated: object-relational, XML+relational ...

# Database Engines in Use

**Number of systems per category, January 2022**




© 2022, DB-Engines.com

[https://db-engines.com/en/ranking\\_categories](https://db-engines.com/en/ranking_categories)

# Database Engines in Use

**Ranking scores per category in percent, January 2022**




© 2022, DB-Engines.com

[https://db-engines.com/en/ranking\\_categories](https://db-engines.com/en/ranking_categories)

# Database Engines in Use

Complete trend, starting with January 2013



[https://db-engines.com/en/ranking\\_categories](https://db-engines.com/en/ranking_categories)

relational  
is very stable

# Definitions

- **Relational Database:** a set of relations
- **Relation:** Consists of two parts:
  - **Schema:** specifies name of relation, plus a set of attributes, plus the domain/type of each attribute
    - E.g., `Students (sid:int, name:string, login:string, faculty:string, major:string)`
  - **Instance:** set of tuples (all tuples are distinct).
  - Compare with entity set and entity; or with object class and object instance
- A relation can be seen as a **table:**
  - Column headers = attribute names, rows = tuples/records, columns/fields = attribute values
- If clear from context, we say instead of “instance of a relation” simply “relation”
- **Database Schema:** collection of relation schemas

# Example of *Students* Relation

*instance is a row*

Column headers = attributes

sid	name	login	faculty	major
53666	Bartoli	bartoli@cs	Science	Software Engineering
53688	Chang	chang@eecs	Eng	Software Engineering
53650	Chang	chang@math	Science	Computer Science
...				

- All rows are distinct (set-oriented)
- Rows are not ordered (a permutation of rows represents still the same table)
- Columns are per definition not ordered but in practice we often assume a fixed order
  - with this, a single tuple can be represented as

- (53666, Bartoli, bartoli@cs, Science, Software Engineering)

8

*interact w/ database  
using column names*



# Relational DDL and DML

- **Data Definition Language (DDL)**: defines the schema of a database
- **Data Manipulation Language (DML)**: “manipulates” the data, i.e., the instances of the relations  
(interact w/ database)
  - – Insert, update, delete tuples
  - – “Query” the relations: retrieve tuples that fulfill certain criteria (hence, often called “query language”)
    - The Relational Model offers simple and powerful querying of data with precise semantics
    - Data-centric – use column names eg
- *Physical data independence*
  - User only sees relations ←
  - User does not need to know how data is stored
    - Mapping to files
    - Rows vs. columns
  - User does not need to know how queries are executed

# The SQL Query Language

- Developed by IBM (system R) in the 1970s
- Need for a standard since it is used by many vendors
- Standards:
  - SQL-86
  - ...
  - SQL-99 / SQL3 (adds object-relational features)
  - SQL:2003 (adds XML features)
  - ...
  - SQL:2011
- Used to define relations (Data definition)
- Used to write and query data (Data manipulation)
- Standard is very slow to develop
- Many dialects in various DBS

# Concepts vs. Details

- We will discuss the main concepts of SQL (data definition and data manipulation) in class
- To figure out the details, students have to use the handbook
  - In particular: DBS specific dialects, variations...

# SQL Data Types

- All attributes must have a data type.
- SQL supports several basic data types
  - Just as any programming language does
  - Notation somewhat “old-fashioned” as language very old
- Character and string types
  - CHAR(n) denotes a character string of fixed length (containing trailing blanks for padding if necessary). *- fixed sized*
  - VARCHAR(n) denotes a string of up to n characters (between 0 and n characters). *- variable sized*
  - SQL permits reasonable coercion between values of character-string types
- Integer Types
  - INT or INTEGER (names are synonyms)
  - SHORTINT

# Data Types (contd.)

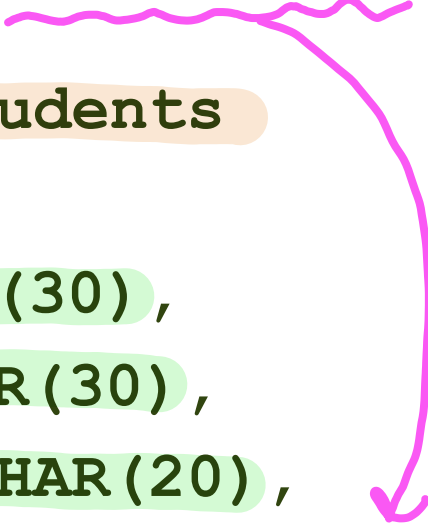
- Floating point numbers
  - **FLOAT** or **REAL** (names are synonyms)
  - **DOUBLE PRECISION**
  - **DECIMAL(n,d)**: real number with fixed decimal point. Value consists of n digits, with the decimal point d positions from the right.
- Dates and time:
  - **DATE**: has the form **'YYYY-MM-DD'**
  - **TIME**: has the form **'15:00:02'** or **'15:00:02.5'**
  - May be compared and converted to string types

} very important
- Bit strings
- **User defined domains**
  - New name for a data type
  - Possibility to define restrictions on values of domain ( $< 10$ )

# Data Definition: Table Creation

- Defines all attributes of the relation
- The type/domain of each attribute is specified
- DBMS enforce correct type whenever a tuple is added or modified
- SQL is case insensitive
- It is possible to define default values

```
CREATE TABLE Students
(
  sid INTEGER,
  name VARCHAR(30),
  login VARCHAR(30),
  faculty VARCHAR(20),
  major VARCHAR(20) DEFAULT 'undefined' )
```



# Data Manipulation: Insert



# Data Manipulation: Insert

- Insert a single tuple

```
INSERT INTO Students
```

```
VALUES (53666, 'Bartoli', 'bartoli@cs',  
        'Science', 'Software Engineering')
```

} Same  
order  
as  
creation

- can contain all or only a subset of attributes:

- Not indicated attributes have special NULL value

```
INSERT INTO Students (sid,name,faculty)
```

```
VALUES (53688, 'Chang', 'Eng')
```

← specify  
different  
orders

```
INSERT INTO Students (sid,name,major)
```

```
VALUES (53650, 'Chang', 'Computer Science')
```

sid	name	login	faculty	major
53666	Bartoli	bartoli@cs	Science	Software Engineering
53688	Chang	NULL	Eng	undefined
53650	Chang	NULL	NULL	Computer Science



# Data Manipulation: Delete

- Can delete all tuples satisfying some condition

**DELETE**

**FROM** Students

**WHERE** name = 'Chang'

sid	name	login	faculty	major
53666	Bartoli	bartoli@cs	Science	Software Engineering
53688	Chang	NULL	Eng	NULL
53650	Chang	NULL	NULL	Computer Science

deletes entire row

# Data Manipulation: Update

- Can update all tuples satisfying some condition

```
UPDATE Students
```

```
SET major = 'Software Engineering'
```

```
WHERE sid = 53688
```

sid	name	login	faculty	major
53666	Bartoli	bartoli@cs	Science	Software Engineering
53688	Chang	NULL	Eng	Software Engineering
53650	Chang	NULL	NULL	Computer Science

# Querying the Data

- Find the names and major of all students in the Faculty of Science

```
SELECT name, major
```

```
FROM Students
```

```
WHERE faculty = 'Science'
```

select \* y  
from students  
gives all rows

sid	name	login	faculty	major
53666	Bartoli	bartoli@cs	Science	Software Engineering
53688	Chang	NULL	Eng	NULL
53650	Chang	NULL	Science	Computer Science

Bartoli	Software Engineering
Chang	Computer Science


select \*  
from students  
where ...

Much more of that  
in a few weeks...  
for all roles

# Data Definition: Destroying Relations

## **DROP TABLE** Students

- Destroys the relation Students. The schema information and the tuples are deleted.



# Data Definition: Altering Tables

**DROP TABLE Students**

- Destroys the relation Students. The schema information *and* the tuples are deleted.

**ALTER TABLE Students**

**ADD COLUMN firstYear:integer**

- The schema of students is altered by adding a new field;
- every existing tuple in the current instance is extended with a **null** value in the new field.

# Integrity Constraints (ICs)

- Integrity Constraints must be true for any instance of the database;
  - Domain definitions (INT or FLOAT) are already a form of constraint
  - Database designer specifies ICs when schema is defined
  - DBMS checks whether ICs remain true whenever relations are modified.
    - If modification violates an IC, the DBMS disallows the modification (throws an error message)
  - Of course, DBMS can only check what is specified in the schema

# Not Null

1C

```
CREATE TABLE Students
(sid CHAR(9) ,
name VARCHAR(30) NOT NULL ,
login VARCHAR(30) ,
faculty VARCHAR(20) ,
major VARCHAR(20)
    DEFAULT 'undefined' )
```

**requires an attribute to always have a proper value**

# Primary Key Constraints

- A set of fields is a **key candidate** for a relation if
  - No two distinct tuples can have same values in all key fields, and
  - This is not true for any subset of the key.
    - Minimum subset of attributes that fulfill uniqueness property ← *unique row identifier for combined key*
- If there are two or more keys, one of the candidates is chosen to be the **primary key**.
- The primary key attributes of a tuple may not be NULL.
- E.g. **sid** is a key for **Students**. (What about **name**?).
- Example of **combined keys**:
  - `Location(building, roomNo, capacity)`
  - Attributes `building` and `roomNo` together build the primary key

*want minimum subset*



# Primary and Candidate Keys in SQL

- Possibly many candidate keys exist, one of which is chosen as the primary key

- Each student has a unique id.
- Each student has a unique login

```
CREATE TABLE Students
```

```
(sid CHAR(9) PRIMARY KEY,
```

```
login VARCHAR(30) NOT NULL UNIQUE,
```

```
name VARCHAR(20),
```

```
...)
```

*/ unique by definition*

*← other candidate*

- Defining primary keys that have more than one attribute

```
CREATE TABLE Location
```

```
(building VARCHAR(20),
```

```
roomNo INT,
```


```
capacity INT,
```

```
PRIMARY KEY(building, roomNo)
```

*) last line*

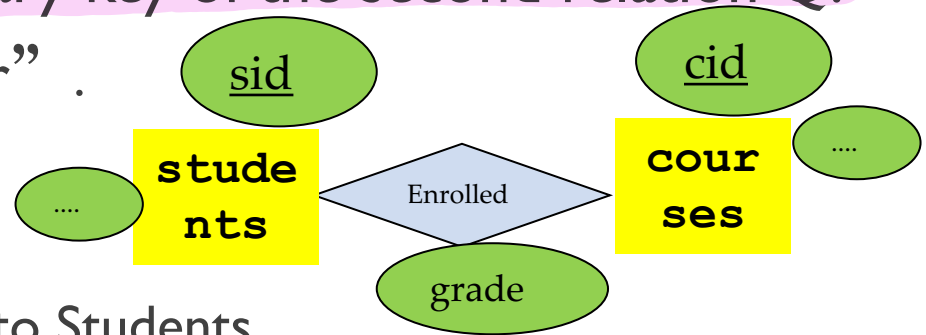


# Note

- Specifics of individual database systems
  - E.g., DB2 required the following
    - If primary key, then also NOT NULL must be defined
  - If there are error messages 
    - Consult the handbooks of the resp. database system
- **Multi-Set** *good for importing data to clean*
  - Tables do not require to have a primary key
  - If there is no primary key
    - Two records/tuples can have the same value in all attributes
    - That is, duplicates are allowed
  - Different to Relational Algebra Assumptions
    - Relations are set → no identical tuples in a relation

# Foreign Key

- **Foreign Key**: Set of attributes in one relation  $R$  that is used to “refer” to a tuple in another relation  $Q$ .
  - Must correspond to the primary key of the second relation  $Q$ .
  - Represents a “logical pointer”.
- Examples
  - in relation **Enrolled**,
    - sid is a foreign key referring to Students



Q

<u>sid</u>	name	login	faculty	...
53666	Bartoli	bartoli@cs	Science	...
53688	Chang	chang@eecs	Eng	...
53650	Chang	chang@math	Science	...
...				

R

<u>sid</u>	<u>cid</u>	grade
53666	Topology112	C
53666	Reggae203	B
53650	Topology112	A
53668	History105	B

# Referential Integrity

- **Foreign Key Constraint:** the foreign key value of a tuple must represent an existing tuple in the referred relation
  - A tuple with this value in the primary key must exist in the referred relation
- { – Enrollment may only contain a tuple referring to a student who exists in the Students relation
- If all foreign key constraints are enforced, **referential integrity** is achieved, i.e., no dangling references

# Foreign Keys in SQL

Only students listed in the Students relation should be allowed to enroll for courses

```
CREATE TABLE Enrolled
(
    sid CHAR(9),
    cid VARCHAR(20),
    grade CHAR(2),
    PRIMARY KEY (sid,cid),
    FOREIGN KEY (sid) REFERENCES Students,
    FOREIGN KEY (cid) REFERENCES Courses
)
```

Student

<u>sid</u>	name	login	faculty	...
53666	Bartoli	bartoli@cs	Science	...
53688	Chang	chang@eecs	Eng	...
53650	Chang	chang@math	Science	...
...				

Enrolled

sid	cid	grade
53666	Topology112	C
53666	Reggae203	B
53650	Topology112	A
53668	History105	B

# Enforcing Referential Integrity: Default

- An Enrolled tuple with a sid is inserted but no tuple with this sid exists in Students
  - Disallow insertion
- A Students tuple is deleted
  - Disallow the deletion of a Students tuple to which Enrolled tuples point
- Other options:
  - cascade

*remove  
reference first*