

# Internals of a DBS I

Query Optimization  
And Execution

---

Relational Operators

---


Files and  
Access Methods

---

Buffer Management

---

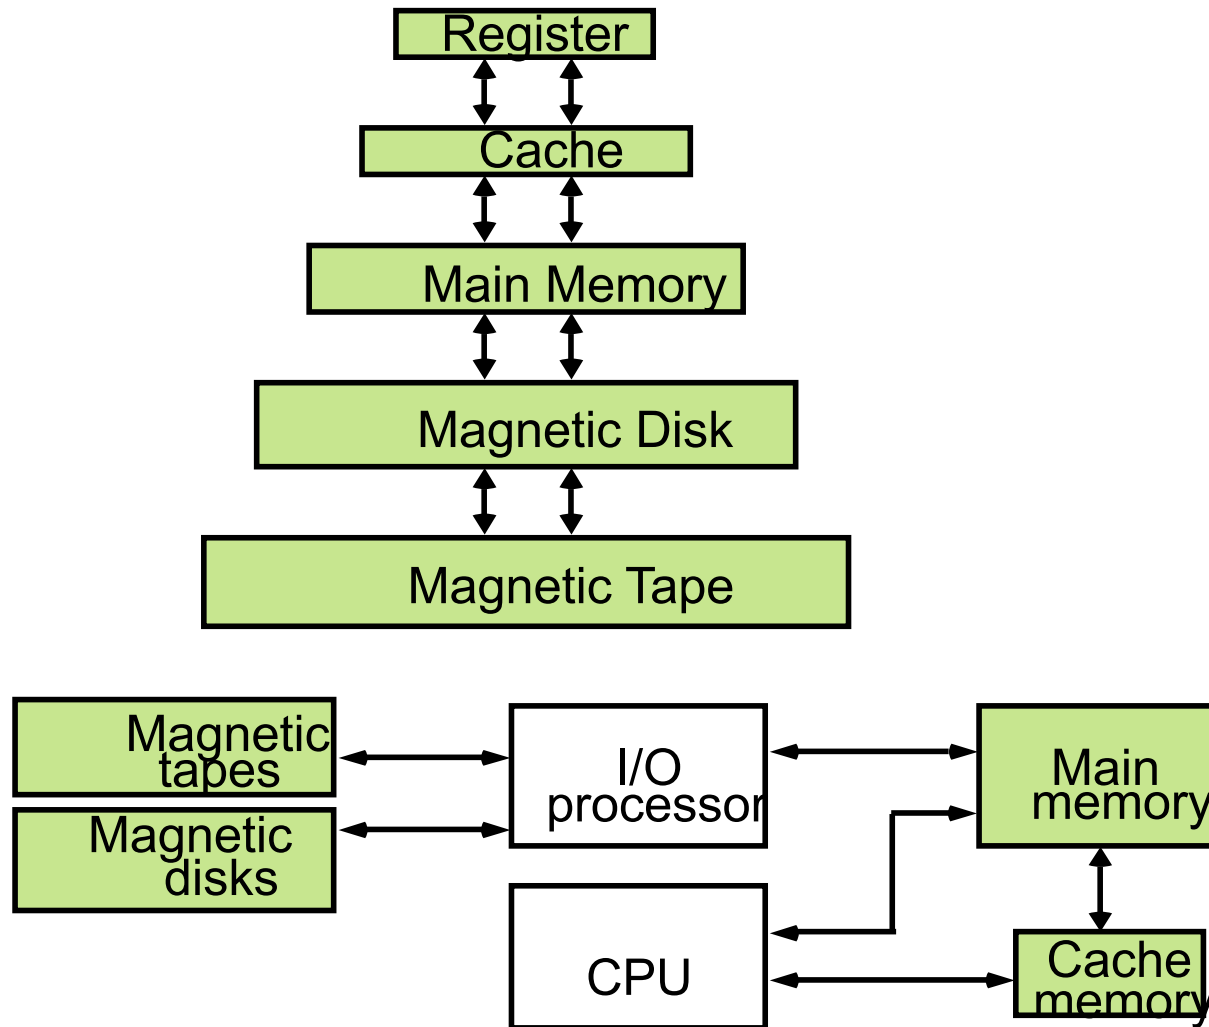
Disk Space  
Management



# The very Essentials of Disk and Buffer Management

# Memory Hierarchy

Memory Hierarchy is to obtain the highest possible access speed while minimizing the total cost of the memory system



# Tape drives



[https://en.wikipedia.org/wiki/9\\_track\\_tape](https://en.wikipedia.org/wiki/9_track_tape)



[https://en.wikipedia.org/wiki/Tape\\_library](https://en.wikipedia.org/wiki/Tape_library)




# Disks and Main Memory

- DBMS stores information persistently on ( “hard” ) disks.
- Unit of transfer main-memory/disk: *disk blocks* or *pages*.



# Why Block/Page Concept



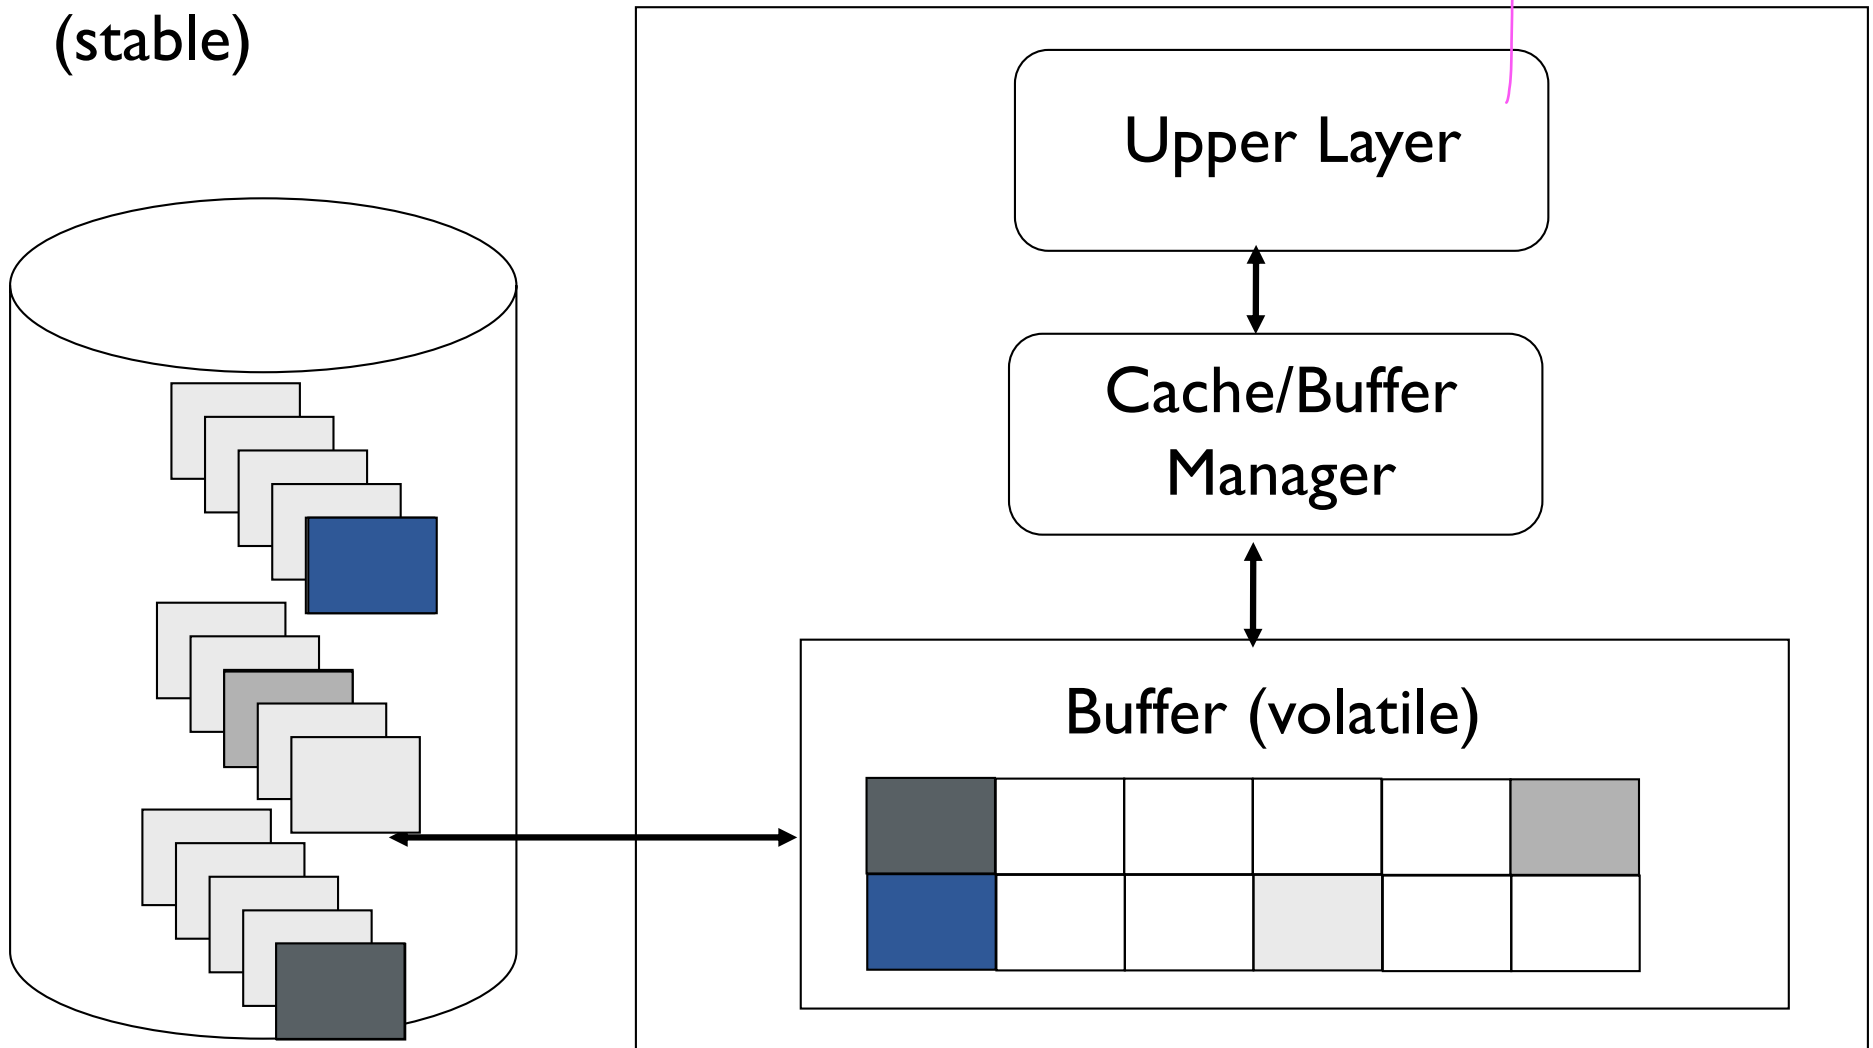
# Disks and Main Memory

- DBMS stores information persistently on ( “hard” ) disks.
- Unit of transfer main-memory/disk: *disk blocks* or *pages*.
- Time to read/write data block:
  - 2- 10 msec for random data block (main factor is seek time)
  - If blocks are sequentially on disk, each additional block only 1 msec
  - Compare main memory access: in nanoseconds
  - SSD: 0.1 ms (still slower than main memory access and calculations...)
- Basic operations
  - **READ**: transfer data from disk to main memory (RAM).
  - **WRITE**: transfer data from RAM to disk.
- Why disks?
  - Cheaper than Main Memory
  - Higher Capacity
  - Main Memory is volatile

Secondary  
Storage  
(stable)

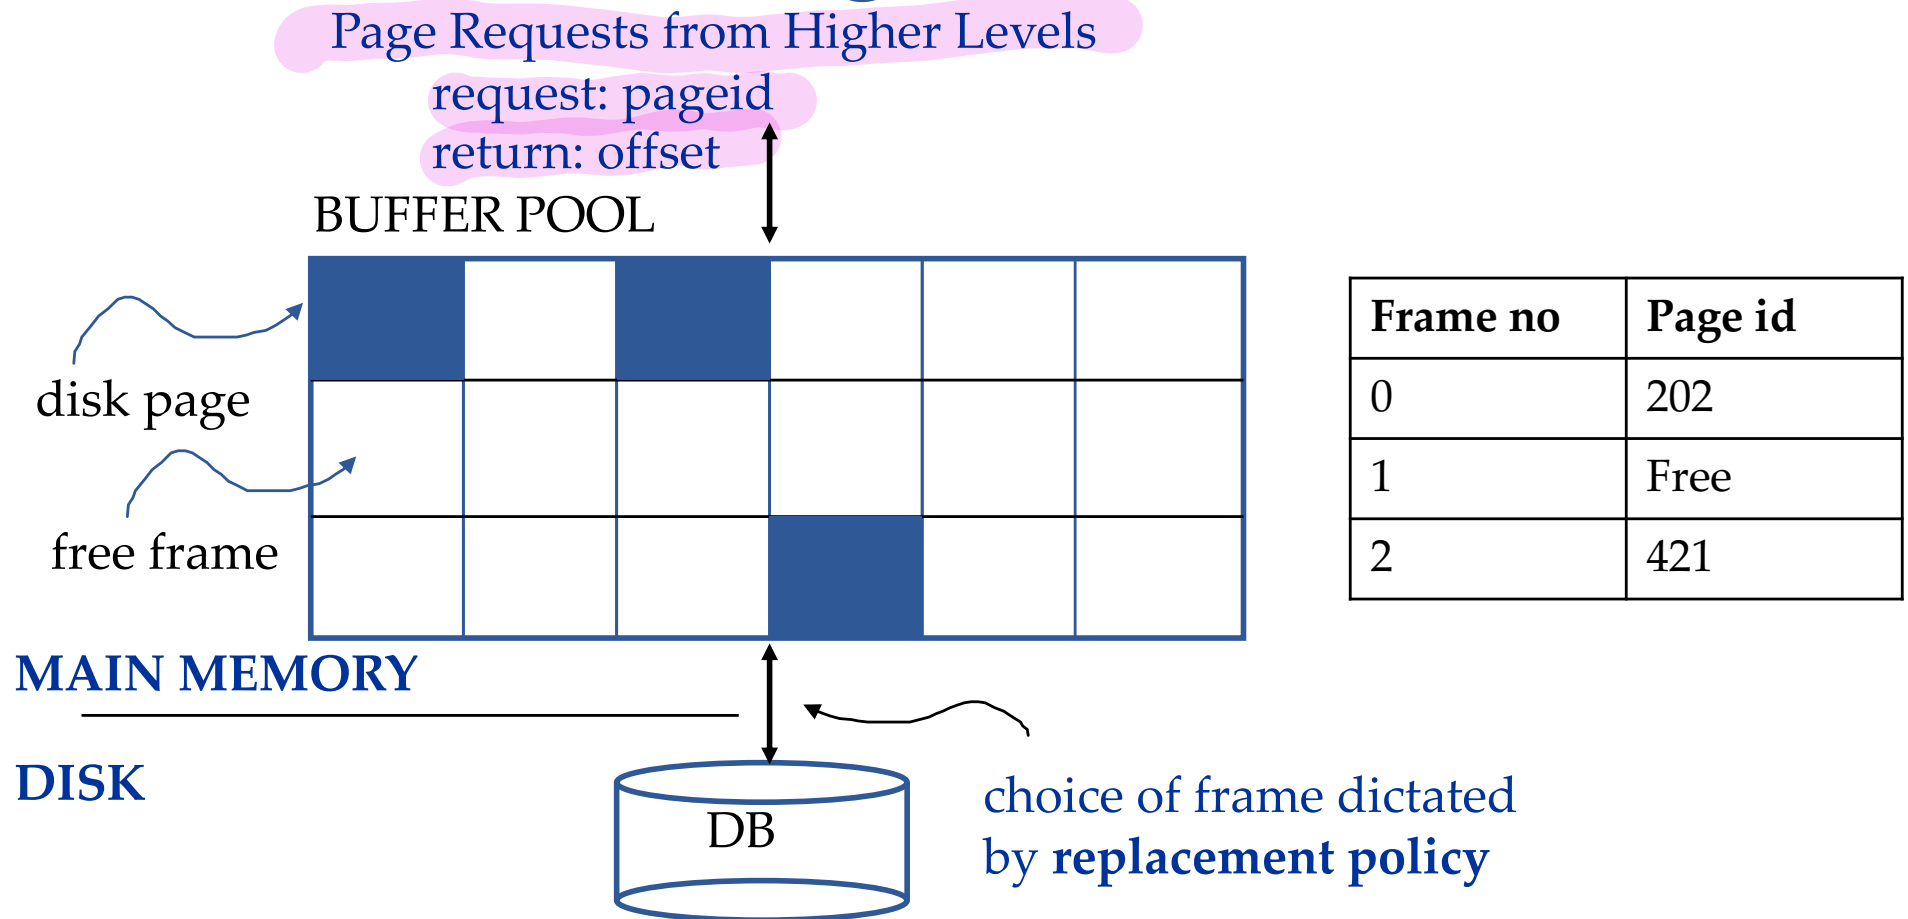
# Architecture

*must know  
page id to  
access*





# Buffer Management in a DBMS



- ❑ Data must be in RAM for DBMS to operate on it!
- ❑ Table of <frame#, pageid> pairs is maintained.
- ❑ Some more information about each page in buffer is maintained

# Loading a page from disk...

□ If requested page is not in pool:

- ☆ If there is an empty frame

  - Choose empty frame

- ☆ Else (no empty frame)

  - Choose a frame for *replacement*

  - If frame is dirty (current page was modified), write it to disk

- ☆ Read requested page into chosen frame

☆ How does the buffer manager know whether a page is dirty?

↳ dirty marker

# Page Pins

Frame no	Page id	Pin counter
0	202	2
1	431	0
2	...	...

- Pin counter for each frame
- When buffer manager returns a page (offset) to a request from upper layer (after possibly loading the page from disk):
  - Buffer manager increases pin counter of corresponding frame
- When upper layer has finished operations on page
  - Upper layer informs buffer manager (release page) and whether page was updated
  - Buffer manager decreases pin counter of frame (and sets dirty bit if modified)
- When loading a page from disk:
  - Replacement frame must have “pin counter” of 0
- Frame is chosen for replacement by a replacement policy:
  - Only unpinned page can be chosen (pin count = 0)
  - Least-recently-used (LRU), Clock, MRU etc.


pin count = number  
of upper layer  
progs using page

# DBMS vs. OS File System

OS does disk space & buffer mgmt: why not let OS manage these tasks?


- Differences in OS support: portability issues
- Some limitations, e.g., files can't span disks.
- Buffer management in DBMS requires ability to:
  - **pin a page** in buffer pool, **force a page** to disk (important for implementing CC & recovery),
  - adjust **replacement policy**, and **pre-fetch pages** based on access patterns in typical DB operations.

# Record Format: Fixed Length



- Length of field (attribute) depends on type
- Works with fixed-length types
- Strings: padding — to reach max value for string
- Offset of each field easy to calculate

# Variable Length

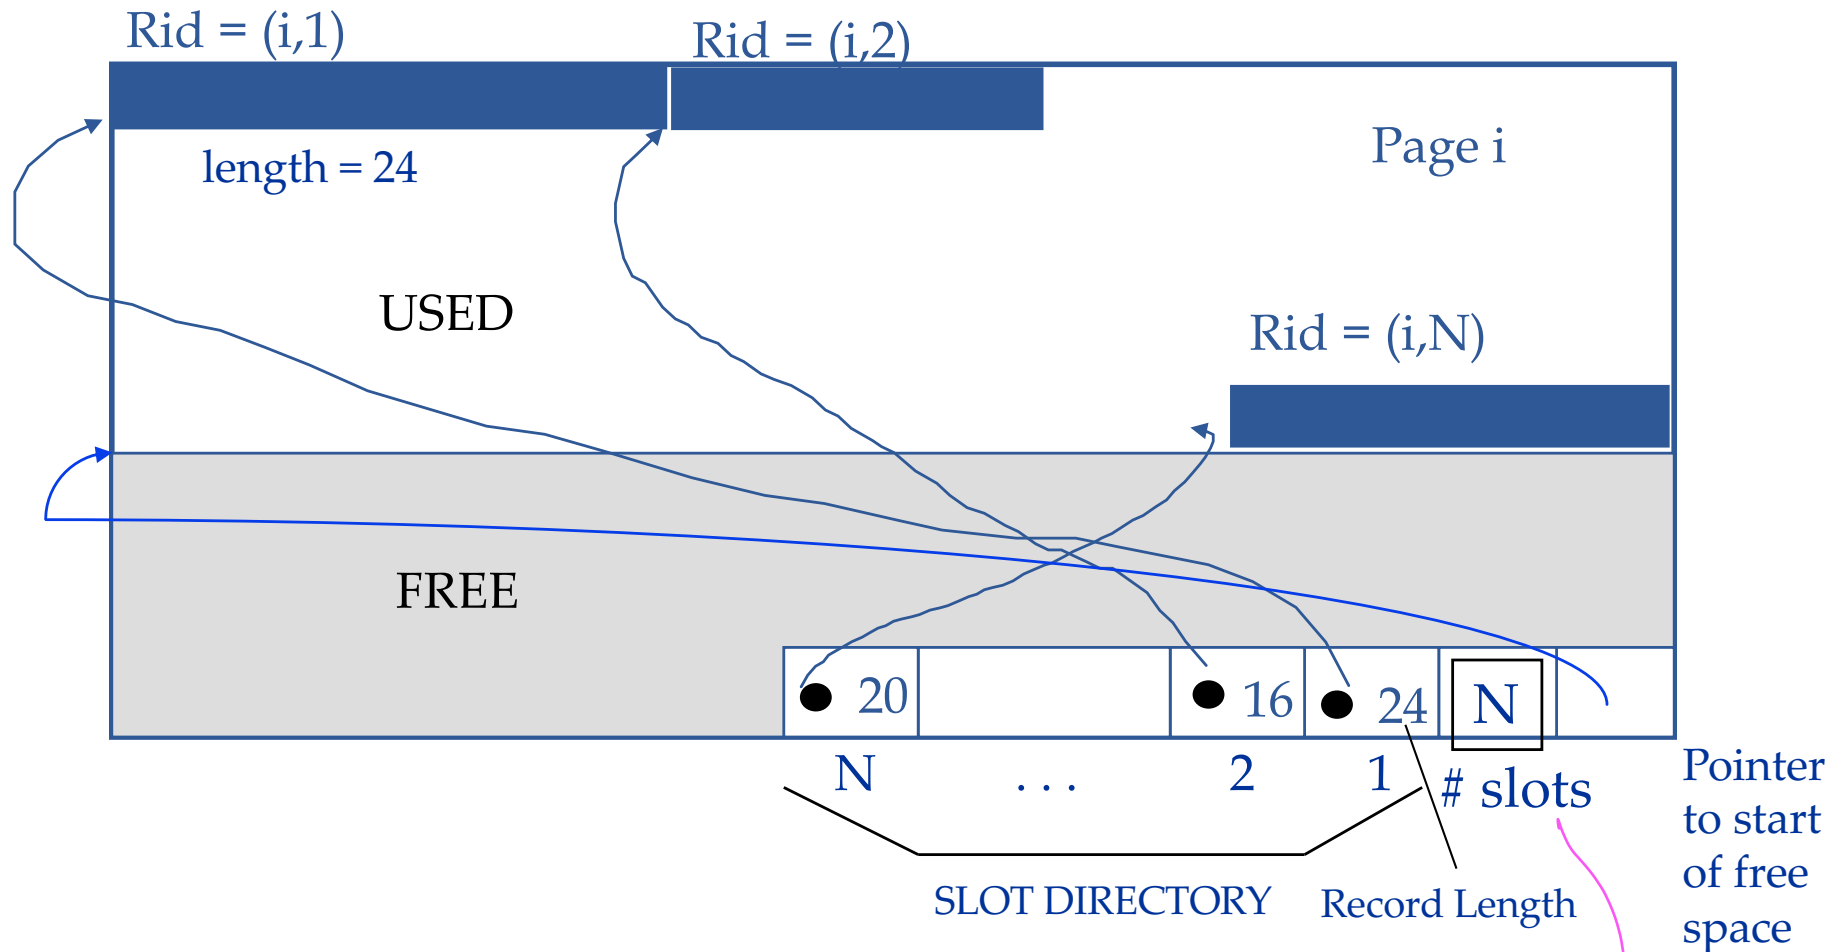


- ➡ Second offers direct access to  $i$ 'th field, efficient storage of nulls; small directory overhead

trade off for  
pointer/meta  
info



# Page Formats: Variable Length Records



➡ **Record id (rid)** = internal identifier of a record:  
 <page id, slot #>.


➡ Can move records on page without changing rid;

unique semantic identifier  
 #records  
 just change where pointer points

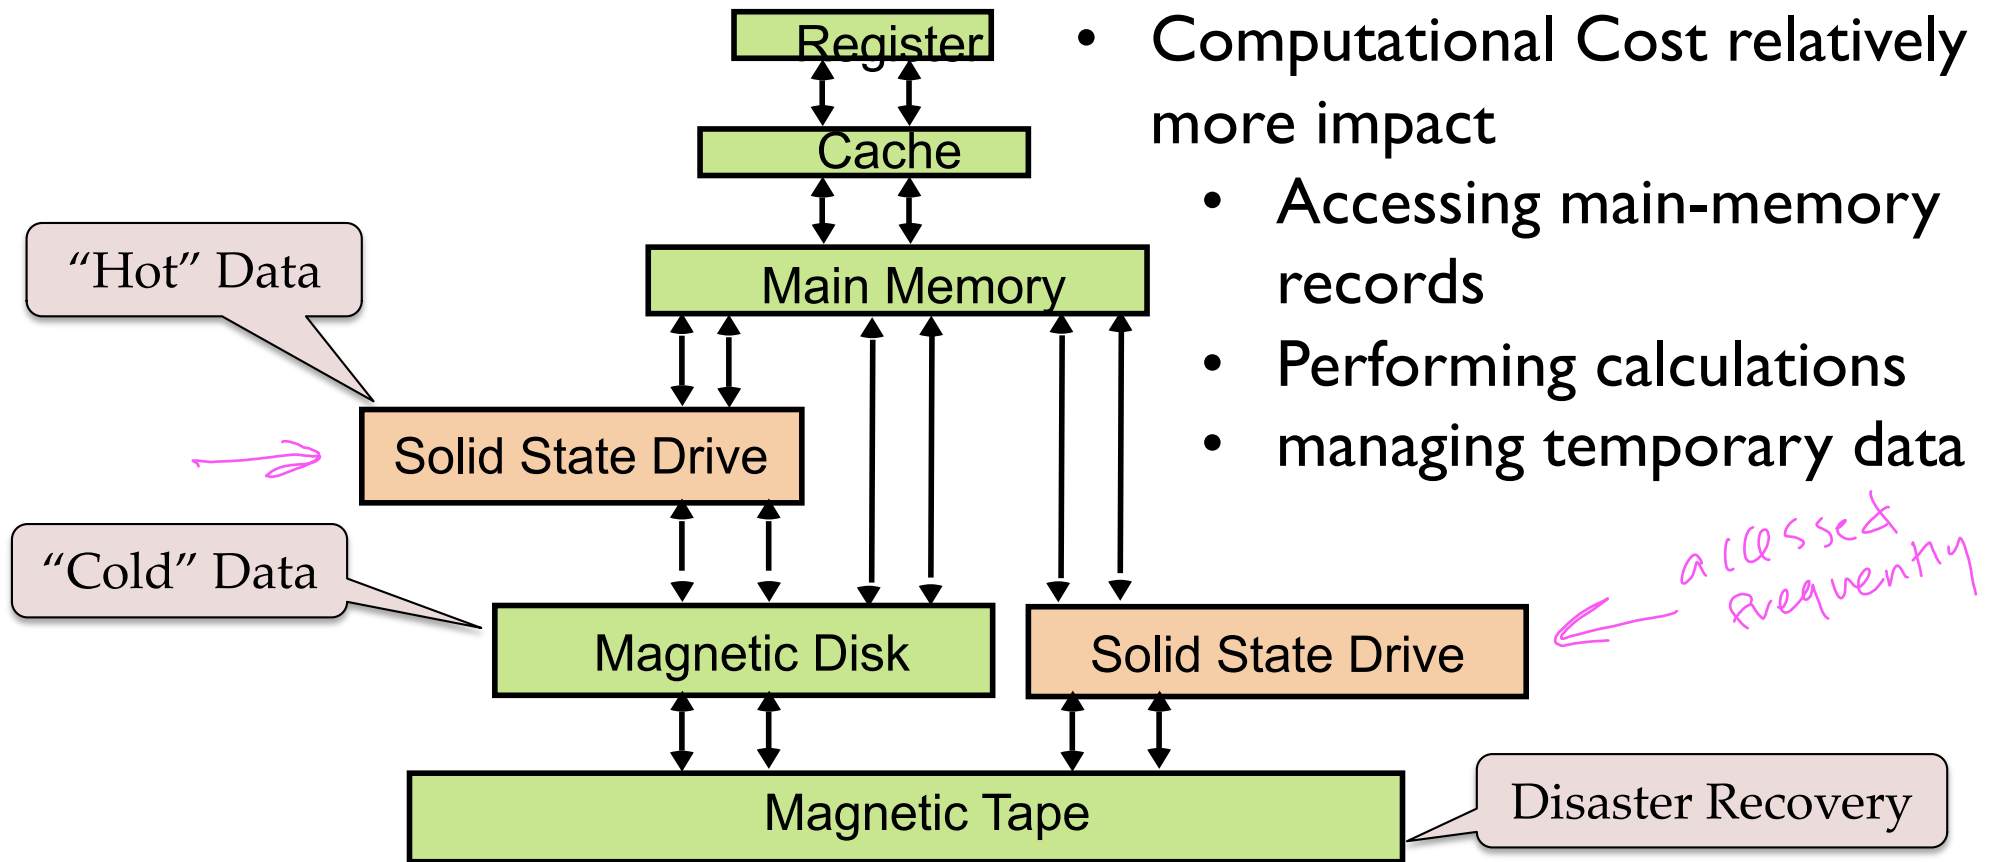
# Summary and Assumptions

- Data records distributed across pages
- Data pages need to be in main memory to be accessed
- Page I/O is much more expensive than any computation that then accesses the data in main memory
- Not all data fits into main memory
- → costs calculated by number of I/O
- → optimize to have as little I/O as possible

# Assumed Distribution





# Trends we ignore: Faster stable storage



# Trends we ignore: memory databases

*all data in main memory*



Secondary Storage (stable)

*— only for persistence / fault tolerance*

Focus on main memory

# Main-memory DBS


- Only I/O for updates (can be highly optimized)
- Optimizing for quick execution becomes main focus
- Layout in main-memory much more important
  - the concept of splitting everything into pages no more relevant
- Main memory algorithms



# Trends we ignore: Column-based database systems

- Store each column of an array of items of the same type
- `SELECT sid from Skaters;`
- `SELECT avg(age) from Skaters;`


# Later: large scale distribution



# Large Scale distribution

- Communication and Coordination Overhead have huge impact on performance