

# Managing and Accessing Records of a Table

# Managing the tuples of a relation

- File: Collection of pages
  - For instance: Pages hold records of one relation
- Management Interface:
- Insert
  - `INSERT INTO Students VALUES (23, 'Bertino', 2016, ... )`
  - Typically returns a rid

*don't  
mix relations  
records  
in a page*

# Retrieving the tuples of a relation

- Get a record by rid
  - Returns record
  - (or rather returns the start position of record in main memory)

More general

- Scan over all records
  - `SELECT * FROM Students`
- Point Query
  - `SELECT * FROM Students WHERE sid = 100`
- Equality Query
  - `SELECT * FROM Students WHERE startyear = 2015`
- Range Search
  - `SELECT * FROM Students`  
`WHERE startyear > 2012 and startyear <= 2014`

Retrieval and management

- `DELETE FROM Students WHERE sid = 100`
- `DELETE FROM Students WHERE endyear < 1950`

} Search then delete

# Managing relation in a file

- As file grows and shrinks, disk pages are allocated and de-allocated.
- To support record level operations, we must:
  - keep track of the *pages* in a file
  - keep track of *free space* on pages
  - keep track of the *records* on a page
- There are many alternatives for keeping track of this.

# Cost Model for Execution

❑ How should we estimate the costs for executing a statement?

★ Number of I/Os ←

★ CPU Execution Cost ←

★ Network Cost in distributed system (ignore for now) ←

★ Assumption in this course

★ I/O cost >>> CPU cost ←

★ Real systems also consider CPU ←

❑ Simplifications

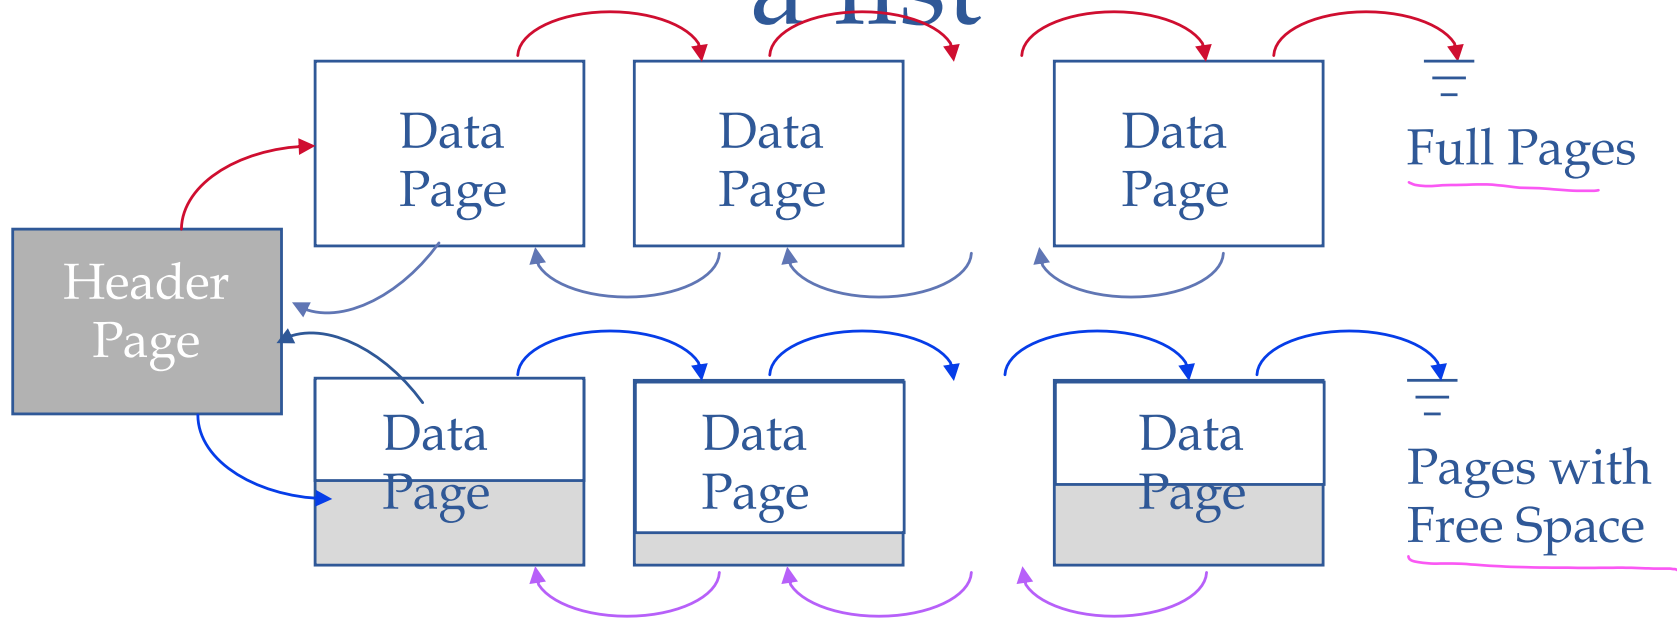
★ only consider disk reads (ignore writes -- assume read-only workload)

★ only consider number of I/Os and not the individual time for each read (ignores page pre-fetch)

★ Average-case analysis; based on several simplistic assumptions.

☞ Good enough to show the overall trends!

# Unsorted heap file implemented as a list



- The header page id and Heap file name must be stored someplace.
- Each page contains 2 'pointers' plus data.

# Heap File

- ☆ Linked, unordered list of all pages of the file
- ☆ How well does it support the different operations?

- insert

- ▲ Insert in any free page
- ▲ Cost is low (insert anywhere) ←

- scan retrieving all records (SELECT \*)?

- ▲ go from page and page and return each tuple
- ▲ Not much optimization possible
- ▲ But tuples can be stored in a compact way ←

- Point query on unique attributes

- Go from page to page, look at each record and return once record is found
- ▲ have to read on avg. half the pages to return one record

- range search or equality search on non-primary key

- ▲ have to read all pages to return subset of records.

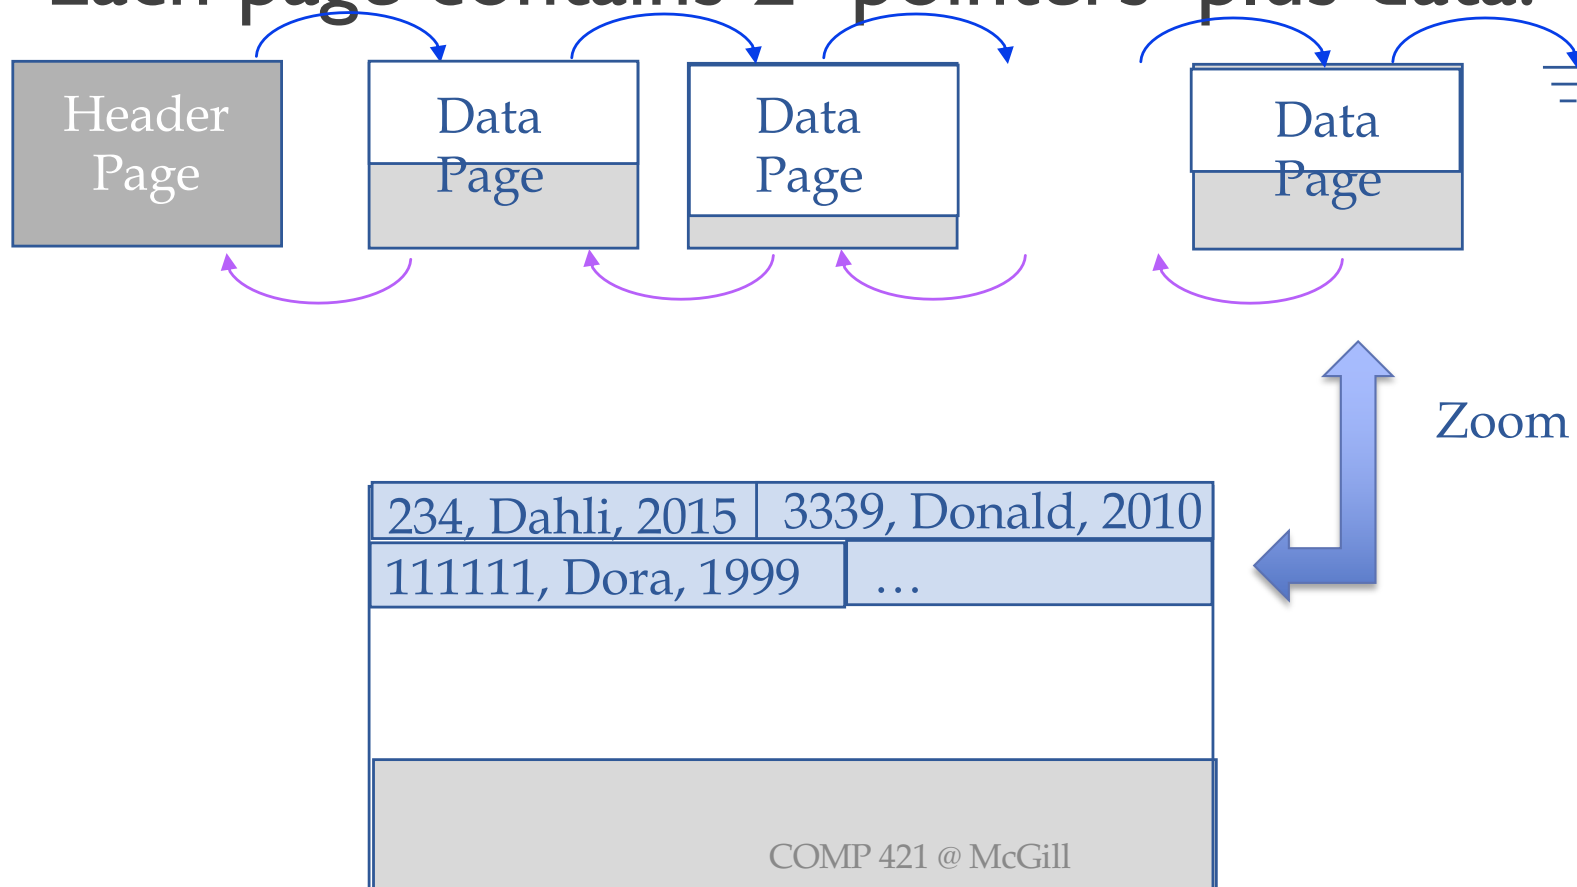
- delete/update

- 7 ▲ same as for equality/range search -- depends on WHERE clause

have to load all pages one at a time  
- find specific record (could be anywhere)

# Sorted file

- Records are sorted by one of the attributes (e.g., name).
- Each page contains 2 'pointers' plus data.





# Sorted File

- insert
  - ▲ have to find proper page
  - ▲ Algorithm to find proper page? → binary search in  $\log_2(\text{number-of-pages})$
  - ▲ overflow possible
  - ▲ Keep empty space on each page → less compact than unsorted heap
- scan retrieving all records (SELECT \*)?
  - ▲ you have to retrieve all pages anyways
- Equality/point search on sort attribute
  - ▲ find first qualifying page with binary search in  $\log_2(\text{number-of-pages})$
- range search on sort attribute
  - ▲ find first qualifying page with binary search in  $\log_2(\text{number-of-pages})$ ; adjacent pages might have additional matching records
- delete/update
  - ▲ finding tuple same as equality/range search depending on WHERE clause
  - ▲ update itself might lead to restructuring of pages
- Sorted output: (ORDER BY)
  - ▲ good if on sorted attribute
- ▲ Search/sort on attributes other than sort attribute
  - ▲ Similar to unsorted heap