

Equality Joins

```
SELECT  *  
FROM    Users U, GroupMembers GM  
WHERE   U.uid = GM.uid
```

p1 _u	uid	uname	experience	age
	123	Dora	2	13
p2 _u	132	Bug	8	60
	267	Sakura	7	15
	111	Cyphon	8	35

p1 _g	uid	gid	stars
	123	G1	2
p2 _g	132	G1	5
	132	G2	3
p3 _g	132	G3	1
	123	G2	4
	111	G4	2

Join Cardinality Estimation

- $| \text{Users} \bowtie \text{GroupMembers} | = ?$ *cardinality of a join*
 - Join attribute is primary key for Users, foreign key in GroupMember
 - Each GroupMember tuple matches exactly with one Users tuple
 - Result: $|\text{GroupMembers}|$
- $| \text{Users} \times \text{GroupMembers} | = ?$ *cross product*
 - Result: $|\text{Users}| * |\text{GroupMembers}|$
 - Cross product is always the product of individual relation sizes
- For other joins more difficult to estimate

Cardinality Estimation

- $| \text{Users} \bowtie \sigma_{(\text{stars} > 3)}(\text{GroupMembers}) | = ?$
 - Result: $| \sigma_{(\text{stars} > 3)}(\text{GroupMembers}) |$
 - Assuming 1-5 stars, uniform distribution for stars $\frac{2}{5}$
 - $\text{Red}(\sigma_{(\text{stars} > 3)}(\text{GroupMembers})) = 0.4$ *reduction factor*
 - Result: $0.4 * |\text{GroupMembers}|$
- $| \sigma_{(\text{experience} > 5)}(\text{Users}) \bowtie (\text{GroupMembers}) | = ?$
 - Assume 1-10 experience levels, uniform distribution for experience
 - $\text{Red}(\sigma_{(\text{experience} > 5)}(\text{Users})) = 1/2$ *reduction factor 5/10*
 - Result: $1/2 * |\text{GroupMembers}|$

Simple Nested Loop Join

- For each tuple in the *outer* relation Users U we scan the entire *inner* relation GroupMembers GM.

foreach tuple u in U do

foreach tuple g in GM do

if u.uid == g.uid then add <u, g> to result

	uid	uname	experience	age
p1 _u	123	Dora	2	13
	132	Bug	8	60
p2 _u	267	Sakura	7	15
	111	Cyphon	8	35

	uid	gid	stars
p1 _g	123	G1	2
	132	G1	5
p2 _g	132	G2	3
	132	G3	1
p3 _g	123	G2	4
	111	G4	2

Simple Nested Loop Join

- For each tuple in the *outer* relation Users U we scan the entire *inner* relation GroupMembers GM.

```
foreach tuple u in U do
  foreach tuple g in GM do
    if u.uid == g.uid then add <u, g> to result
```

- Cost: $\text{UserPages} + |\text{Users}| * \text{GroupMemberPages} = 500 + 40,000 * 1000$!

- NOT GOOD
- We need page-oriented algorithm!

*tuple oriented
algorithm*

Page Nested Loop Join

- For each page p_u of Users U , get each page p_g of GroupMembers GM
 - write out matching pairs $\langle u, g \rangle$, where u is in p_u and g is in p_g .

For each page p_u of Users U

for each page p_g of GroupMembers GM

for each tuple u in p_u do

for each tuple g in p_g do

if $u.uid == g.uid$ then add $\langle u, g \rangle$ to result

	uid	uname	experience	age
$p1_u$	123	Dora	2	13
	132	Bug	8	60
$p2_u$	267	Sakura	7	15
	111	Cyphon	8	35

	uid	gid	stars
$p1_g$	123	G1	2
	132	G1	5
$p2_g$	132	G2	3
	132	G3	1
$p3_g$	123	G2	4
	111	G4	2

Page Nested Loop Join


- For each page p_u of Users U , get each page p_g of GroupMembers GM
 - write out matching pairs $\langle u, g \rangle$, where u is in p_u and g is in p_g .

```
For each page  $p_u$  of Users  $U$ 
  for each page  $p_g$  of GroupMembers  $GM$ 
    for each tuple  $u$  in  $p_u$  do
      for each tuple  $g$  in  $p_g$  do
        if  $u.uid == g.uid$  then add  $\langle u, g \rangle$  to result
```

- Cost: $UserPages + UserPages * GroupPages = 500 + 500 * 1000 = 500,500$

Still not great

Block Nested Loop Join

- For each *block of pages* bp_u of Users U , get each *page* p_g of GroupMembers GM
 - write out matching pairs $\langle u, g \rangle$, where u is in bp_u and g is in p_g .
- *block of pages* bp_u and one page of GM must fit in main memory 
 - For each block of pages bp_u
 - ① • Load block into main memory
 - ② • Get first page from GM
 - Do all the matching between users in bp_u and group members in first page
 - ③ • Get second page from GM (into the same frame the first one was in before)
 - Do all the the matching between users in bp_u and group members in second page
 - ⋮
 - ...
 - ④ • Get last page from GM (into again that frame reserved for GM)
 - ...
- Cost: $UserPages + UserPages / |bp_u| * GroupMemberPages$

Block Nested Loop

	uid	uname	experience	age
$p1_u$	123	Dora	2	13
	132	Bug	8	60
$p2_u$	267	Sakura	7	15
	111	Cyphon	8	35

	uid	gid	stars
$p1_g$	123	G1	2
	132	G1	5
$p2_g$	132	G2	3
	132	G3	1
$p3_g$	123	G2	4
	111	G4	2

Block Nested Loop Join

- Examples depending on available main memory:
- 51 Buffer Frames:
 - $500 + 500/50 * 1000 = 500 + 10,000$
user pages *10 blocks* *join relation*
- 501 Buffer Frames
 - $500 + 500/500 * 1000 = 500 + 1000$
 - Special case: outer relation fits into main memory!!

Index Nested Loops Join

- For each tuple in the *outer* relation Users U we find the matching tuples in GroupMembers GM through an index
 - Condition: GM must have an index on the join attribute

foreach tuple u in U do

find all matching tuples g in GM through index
then add all $\langle u, g \rangle$ to result

	<u>uid</u>	uname	experience	age
p1 _u	123	Dora	2	13
	132	Bug	8	60
p2 _u	267	Sakura	7	15
	111	Cyphon	8	35

	<u>uid</u>	<u>gid</u>	stars
p1 _g	123	G1	2
	132	G1	5
p2 _g	132	G2	3
	132	G3	1
p3 _g	123	G2	4
	111	G4	2

Index Nested Loops Join

foreach tuple u in U do

find all matching tuples g in GM through index

then add all $\langle u, g \rangle$ to result

- Index MUST be on the inner relation (in this case GM).
- Cost: $\text{OuterPages} + \text{CARD}(\text{OuterRelation}) * \text{cost of finding matching tuples in inner relation}$
- In example of previous page:
 - Index on uid on GM is clustered:
 - $500 + 40.000 * (1 \text{ leaf page} + 1 \text{ data pages})$
 - Index on uid on GM is not clustered:
 - $500 + 40.000 * (1 \text{ leaf page} + 2.5 \text{ data pages})$ (on average 2.5 tuples in GM per user)

Index Nested Loops Join

- Switch inner and outer if index is on uid of Users
- Note: uid is primary key in User

– Only one tuple matches!

foreach tuple g in GM do

find the one matching tuple u in U through index

then add <g, u> to result

- Cost: $1000 + 100.000 * (1 \text{ leaf page} + 1 \text{ data page})$


	<u>uid</u>	<u>gid</u>	stars
p1 _g	123	G1	2
	132	G1	5
p2 _g	132	G2	3
	132	G3	1
p3 _g	123	G2	4
	111	G4	2

43

	<u>uid</u>	uname	experience	age
p1 _u	123	Dora	2	13
	132	Bug	8	60
p2 _u	267	Sakura	7	15
	111	Cyphon	8	35

Block Nested Loop vs. Index

- Best Case for Block Nested Loop (if outer relation fits in main memory)
 - OuterPages + InnerPages
- Index Nested Loop:
 - OuterPages + Card(Outer) * matching tuples Inner
- Index Nested Loop wins if:
 - InnerPages > Card(Outer) * matching tuples Inner
 - E.g., if Outer is the result of a selection that only selected very few tuples
 - $\sigma_{(\text{age} > 60)}(\text{Users}) \bowtie (\text{GroupMembers})$



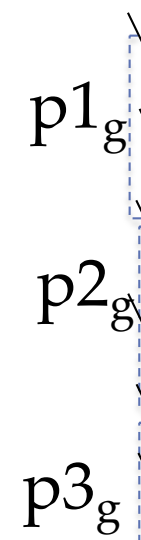
Sort-Merge Join

- Sort U and GM on the join column, then scan them to do a “merge” (on join col.), and output result tuples.
 - In loop:
 - Assume the scan cursors currently points to U tuple u and GM tuple g . Advance scan cursor of U until $u.uid \geq g.uid$ and then advance scan cursor of GM so that $g.uid \geq u.uid$. Do this until $u.uid = g.uid$.
 - At this point, all U tuples with same value in uid (*current U group*) and all GM tuples with same value in uid (*current GM group*) match; output $\langle u, g \rangle$ for all pairs of such tuples.
 - Then resume scanning U and GM.
- U is scanned once; each GM group is scanned once per matching U tuple. (Multiple scans of an GM group are likely to find needed pages in buffer.)

Example of Sort-Merge Join



<u>uid</u>	uname	experience	age
111	Cyphon	8	35
123	Dora	2	13
132	Bug	8	60
267	Sakura	7	15



<u>uid</u>	<u>gid</u>	stars
111	G4	2
123	G1	2
123	G2	4
132	G1	5
132	G2	3
132	G3	1

Cost of Sort-Merge Join

- Relations are already sorted: *(on the disk)*
 - – $\text{UserPages} + \text{GroupMemberPages} = 500 + 1000$
- Relations need to be sorted – simple way:
 - – Sort relations and write sorted relations to temporary stable storage
 - – Read in sorted relations and merge
 - Costs: assuming 100 buffer pages
 - both Users and GroupMembers can be sorted in 2 passes (Pass 0 and 1): $4 * \text{UserPages} + 4 * \text{GroupPages}$
 - Final merge: $500 + 1000 = (\text{UserPages} + \text{GroupPages})$
 - Total: $5 * \text{UserPages} + 5 * \text{GroupPages}$

Cost of Sort-Merge Join

- Relations need to be sorted – use pipelining to combine last sort pass and join
 - Sorting performs Pass 0 reads and writes each of the relations: $2 * \text{UserPages} + 2 \text{ GroupPages}$
 - Pass 1 reads data, sorts and then performs merge in pipeline fashion (ignore details): $\text{UserPages} + \text{GroupPages}$
 - Total: $3 * \text{UserPages} + 3 * \text{GroupPages} = 4,500$