

Query Evaluation

Query Evaluation

- Processing of SQL queries
 - *Parser* translates into internal query representation
 - The parser parses the string and checks for syntax errors
 - Checks existence of relations and attributes
 - Replaces views by their definitions
 - The *query optimizer* translates the query into an efficient execution plan
 - Many different execution plans exist;
 - Choosing an efficient execution plan is crucial
 - The *plan executor* executes the execution plan and delivers data

Query Evaluation

- Query decomposition
 - Queries are composed of **few basic operators**
 - Selection
 - Projection
 - Order by
 - Join
 - Group by
 - Intersection
 - ...
 - Several alternative algorithms for implementing each relational operator
 - Not a single “best” algorithm; efficiency of each implementation depends on factors like size of the relation, existing indexes, size of buffer pool etc.

Basic terminology

- **Access Path**

- The method used to retrieve a set of tuples from a relation

- Basic: file scan
 - index plus matching selection condition (index can be used to retrieve only the tuples that satisfy condition)
 - Partitioning and others

Cost model

- In order to compare different alternatives we must estimate how expensive a specific execution is
- Input for cost model:
 - the query,
 - database statistics (e.g., distribution of values, etc.),
 - resource availability and costs: buffer size, I/O delay, disk bandwidth, CPU costs, network bandwidth,
- Our cost model
 - Number of I/O = pages retrieved from disk
 - assumption that the root and all intermediate nodes of the B+ are in main memory:
 - leaf pages and data pages may not be in main memory!!!
 - A page P **might be retrieved several times** if many pages are accessed between two accesses of P

Basic Query Processing Operations

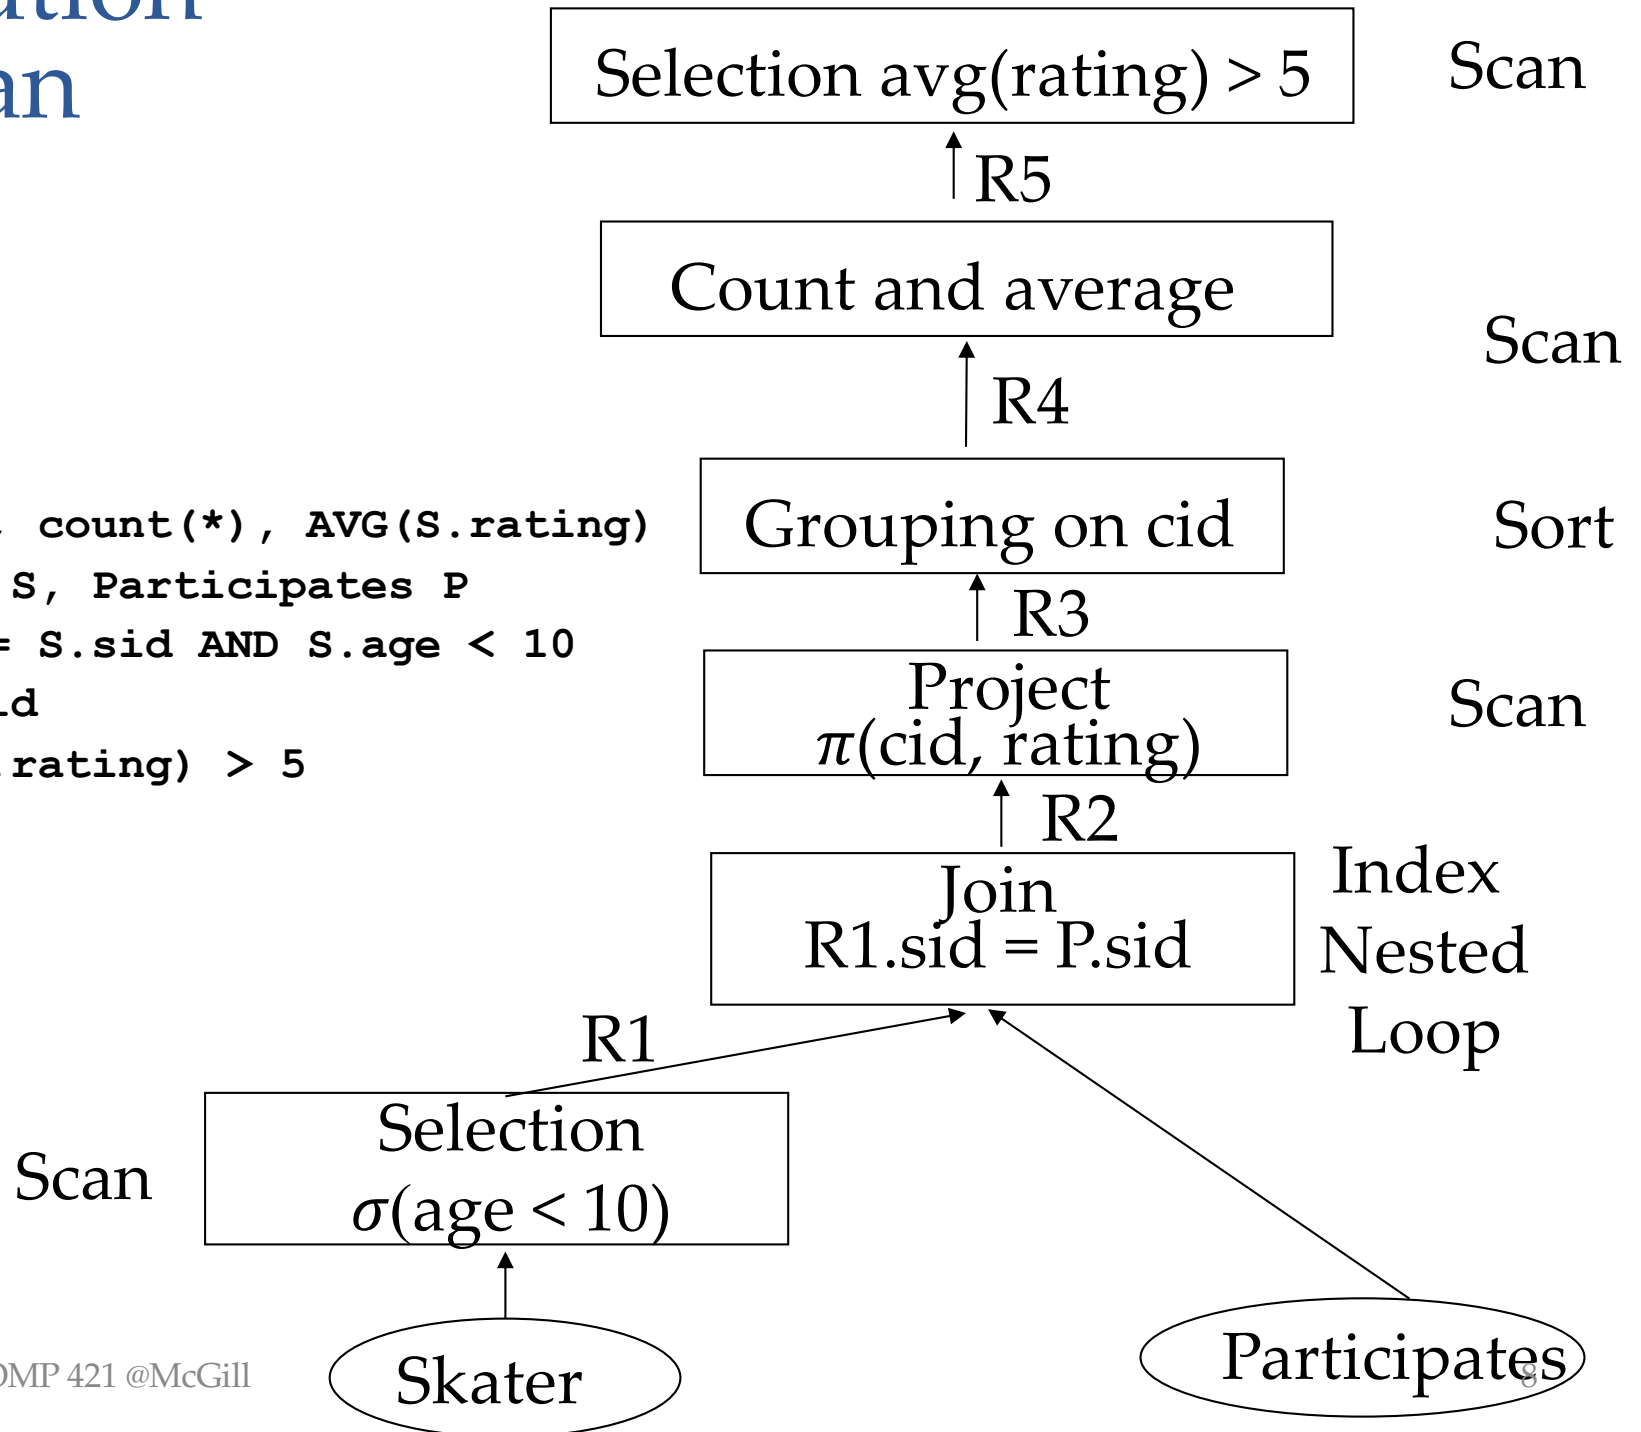
- Selection: σ
 - Scan (without an index)
 - Use an index involving the attributes of the selection
- Projection: π
- Sorting
- Joining
 - Nested loop join
 - Sort-merge join
 - Many others...
- Grouping and duplicate elimination

Concatenation of Operators

- Execution tree:
 - Leaf nodes are the base relations
 - Inner nodes are operators
 - Tuples from base relations (leaves) flow into the parent operator node(s)
 - Output of an operator node flows into the parent node
 - Output of the root flows as result to the client

Execution Plan

```
SELECT P.cid, count(*), AVG(S.rating)
FROM Skaters S, Participates P
WHERE P.sid = S.sid AND S.age < 10
GROUP BY P.cid
HAVING AVG(S.rating) > 5
```



Schema for Examples

Users (uid: int, uname: string, experience: int, age: int)

GroupMembers (uid: int, gid: int, stars: int)

- **Users (U):**
 - 40,000 tuples (denoted as CARD(U))
 - Around 80 tuples per data page
 - 500 data pages (denoted as UserPages)
 - An index on uid has around 170 leaf pages
 - An index on uname has around 300 leaf pages
 - **GroupMembers (GM):**
 - 100,000 tuples (denoted as CARD(GM))
 - Around 100 tuples per data page
 - Total of 1000 data pages (denoted as GroupMemberPages)
 - **Database statistics tables**
 - Keep track of cardinality of tables, number of pages, what indexes, how big an index, etc.
 - keep track of domain of values and their rough distribution
- 9 • E.g., rating values: 1.....10, uniform distribution

Selection

Selectivity / Reduction Factor

- **Reduction Factor** of a condition is defined as
 - $\text{Red}(\sigma_{\text{condition}}(R)) = |\sigma_{\text{condition}}(R)| / |R|$
 - $\text{Red}(\sigma_{\text{experience}=5}(\text{Users})) = |\sigma_{\text{experience}=5}(\text{Users})| / |\text{Users}| = ?$
 - Assume we know that 10,000 users have experience of 5
 - $10.000/40.000 = 0.25$
- If not known, DBMS makes simple assumptions
 - $\text{Red}(\sigma_{\text{experience}=5}(R)) = 1/|\text{different experience levels}| = 0.1$
 - **Uniform distribution assumed**
 - $\text{Red}(\sigma_{\text{age} \leq 16}(R)) = (16 - \min(\text{age}) + 1) / (\max(\text{age}) - \min(\text{age}) + 1) = (16 - 12 + 1) / (61 - 12 + 1) = 5/50 = 0.1$
 - Size of selected range / total size of domain
 - $\text{Red}(\sigma_{\text{experience}=5 \text{ and age} \leq 16}(R)) = ?$
 - Assume uniform and independent distribution
 - $\text{Red}(\sigma_{\text{experience}=5}(R)) * \text{Red}(\sigma_{\text{age} \leq 16}(R))$

Selection

Selectivity / Reduction Factor

- Result sizes: number of input tuples multiplied by reduction factor
- How to know number of different values, how many of a certain value, max, min...:
 - through indices, heuristics, separate statistics (histograms)

Simple Selections

| | |
|---------------|------------------------|
| SELECT | * |
| FROM | Users |
| WHERE | uname LIKE 'B%' |


| | |
|---------------|------------------|
| SELECT | * |
| FROM | Users |
| WHERE | uid = 123 |

- General form: $\sigma_{R.attr \text{ op } value} (R)$
 - No index:
 - Search on arbitrary attribute(s): scan the relation (cost is UserPages=500)
 - Search on primary key attribute: scan on average half of U (cost is UserPages/2=250)
 - Index on selection attributes:
 - Use index to find qualifying data entries, then retrieve corresponding data records.
 - I/O: number of leaf pages with matches + certain number of data pages that have matching tuples
 - Now how much is that??
 - Depends on number of qualifying tuples
- 12 • Depends on type of index

In case of Clustered B+ Tree

Cost:

- Path from root to the left-most leaf **lq** with qualifying data entry:
 - Inner pages in memory
 - One I/O to retrieve **lq** page
- Retrieve page of first qualifying tuple: 1 I/O
- Retrieve all following pages as long as search criteria is fulfilled
- Each data page only retrieved once
- # data pages
 - #matching tuples / tuples per page
 - E.g., if 20% of tuples qualify then roughly 20% of data pages are retrieved




Clustered Index for our Example

- Clustered: matching tuples are in adjacent data pages:
 - # datapages = number of matching tuples / number of tuples per page
 - Example 1: **SELECT * FROM Users WHERE uid = 123**
 - 1 tuple matches because primary key (1 data page)
 - I/O cost: 1 index leaf page + 1 data page
 - Example 2: **SELECT * FROM Users WHERE uname LIKE 'B%'**
 - System estimates the number of matching tuples
 - e.g., around 100 tuples match
 - Since clustered there are all on few pages (say 2 data pages)
 - I/O cost: 1 index leaf page + two data pages
 - Example 3: **SELECT * FROM Users WHERE uname < 'F'**
 - Estimate is that around 10000 tuples match
 - i.e., if 25% of the data, then clustered on approx. 25% of the pages (125 data pages)
 - Cost: 1 leaf pages + 125 data pages
- Note: some systems might retrieve all rids through the index (not efficient).
- 14 In this case, example 3 will read approx. 25% of the leaf pages

In case of non-clustered B+ Tree

Cost:

- Path to first qualifying leaf:
 - One I/O to retrieve **lq** page
- Retrieve page of first qualifying tuple: 1 I/O
- Retrieve page of second qualifying tuple
- ...
- Retrieve next leaf page with qualifying tuple
- Retrieve page of next qualifying tuple
- Sometimes page might have been retrieved previously
 - Might still be in main memory
 - Might have been replaced again



Using an non-clustered Index for Selections

- Non clustered, simple strategy:
 - get one page after the other:
 - Worst case #data pages = #matching tuples
 - Example 1: **SELECT * FROM Users WHERE uid = 123**
 - Same as before
 - Example 2: **SELECT * FROM Users WHERE uname LIKE 'B%'**
 - Estimated that around 100 tuples match
 - In worst case there are on 100 different data pages
 - But even if two are on the same page, when the second record is retrieved the page might already be no more in main memory..
 - cost: 1 index-leaf-page + 100 pages (some pages are retrieved twice)

Using an non-clustered Index for Selections

- Example 3: **SELECT * FROM Users WHERE uname < 'F'**
 - Estimated that 10000 tuples match = 25%
 - Likely that nearly every data page has a matching tuple! (maybe 10 don't???)
 - But as we retrieve tuple by tuple, every retrieval might lead to I/O as page might have been purged from main memory in between
 - cost: 75 leaf pages + 10000 pages (most pages will be retrieved several times)
 - 75 leaves = 25% of all leaf pages
 - Simple scan is faster!!
- Lesson learned:
 - Indices usually only useful with very small reduction factors

Using an non-clustered Index with Sorting for Selections

- Example 3: **SELECT * FROM Users WHERE uname < 'F'**
 - Determine all leaf pages that have matching entries (75 leaf pages)
 - sort matching data entries (rid=pid,slot-id) in leaf-pages by page-id
 - Only fast if the 75 leaf pages with matching entries fit in main memory
 - Retrieve each page only once and get all matching tuples
 - #data pages = #data pages that have at least one matching tuple;
 - worst case is total # of data pages
 - For Example 3
 - Around 10000 tuples match
 - If they are distributed over 490
 - cost: 75 leaf pages + 490 pages (worse than a scan)
 - If they are distributed over 300 pages
 - Cost 75 leaf pages + 300 pages (better than a scan)
 - Note: sorting expensive if leaf-pages do not fit in main-memory

Selections on 2 or more attributes

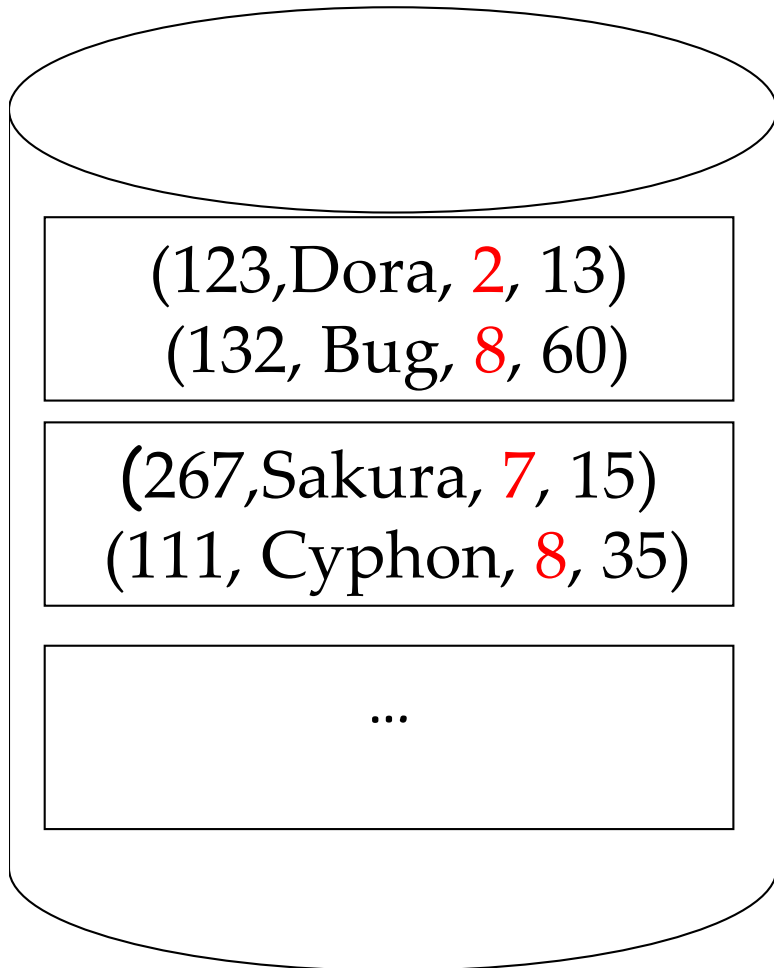
- $A=100 \text{ AND } B=50$ (pages = 500)
 - No index:
 - 500
 - Index on A attribute:
 - Get all tuples for $A=100$ through index, check value of B
 - Cost same as the query for $A=100$
 - 2 indexes; one for each attribute
 - Find rids where $A = 100$ through A index
 - Find rids where $B = 50$ through B index
 - Build intersection of rids
 - Retrieve from data pages all tuples with rids in that intersection
 - In some cases can just use A's index (e.g. when A is unique or has small reduction factor)
 - 1 index with both attributes
- $A=100 \text{ and } B<50$ ($\text{Red}(B<50) = 0.5$)
 - Very low reduction factor for $B<50$, not much use.
- OR

Selections on 2 or more attributes


- $A=100 \text{ OR } B=50$ (pages = 500)
 - No index:
 - 500
 - Index on A attribute:
 - Not useful
 - 2 indexes; one for each attribute
 - Find rids where $A = 100$ through A index
 - Find rids where $B = 50$ through B index
 - Build union of rids
 - Retrieve from data pages all tuples with rids in that union
 - 1 index with both attributes (A,B)
 - Have to read all the leaf pages anyways.

External Sorting

Example Setup



Represented in the following as:




Summary of example:

- 5 pages
- Each with two tuples

External Sorting

Task:

```
SELECT  *  
FROM    Users  
ORDER BY experience
```

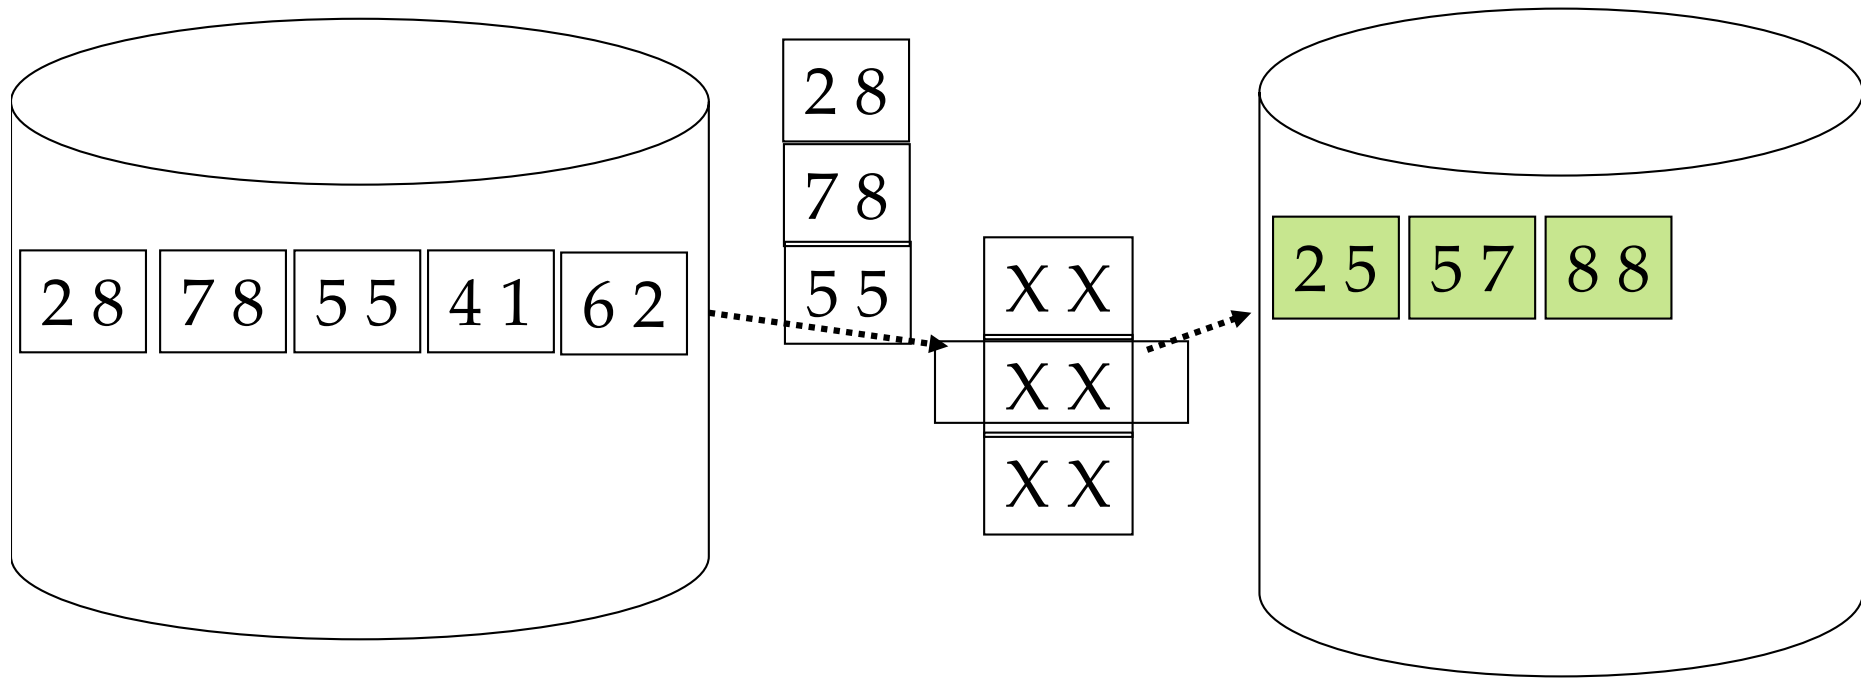


Summary of example:

- 5 pages
- Each with two tuples

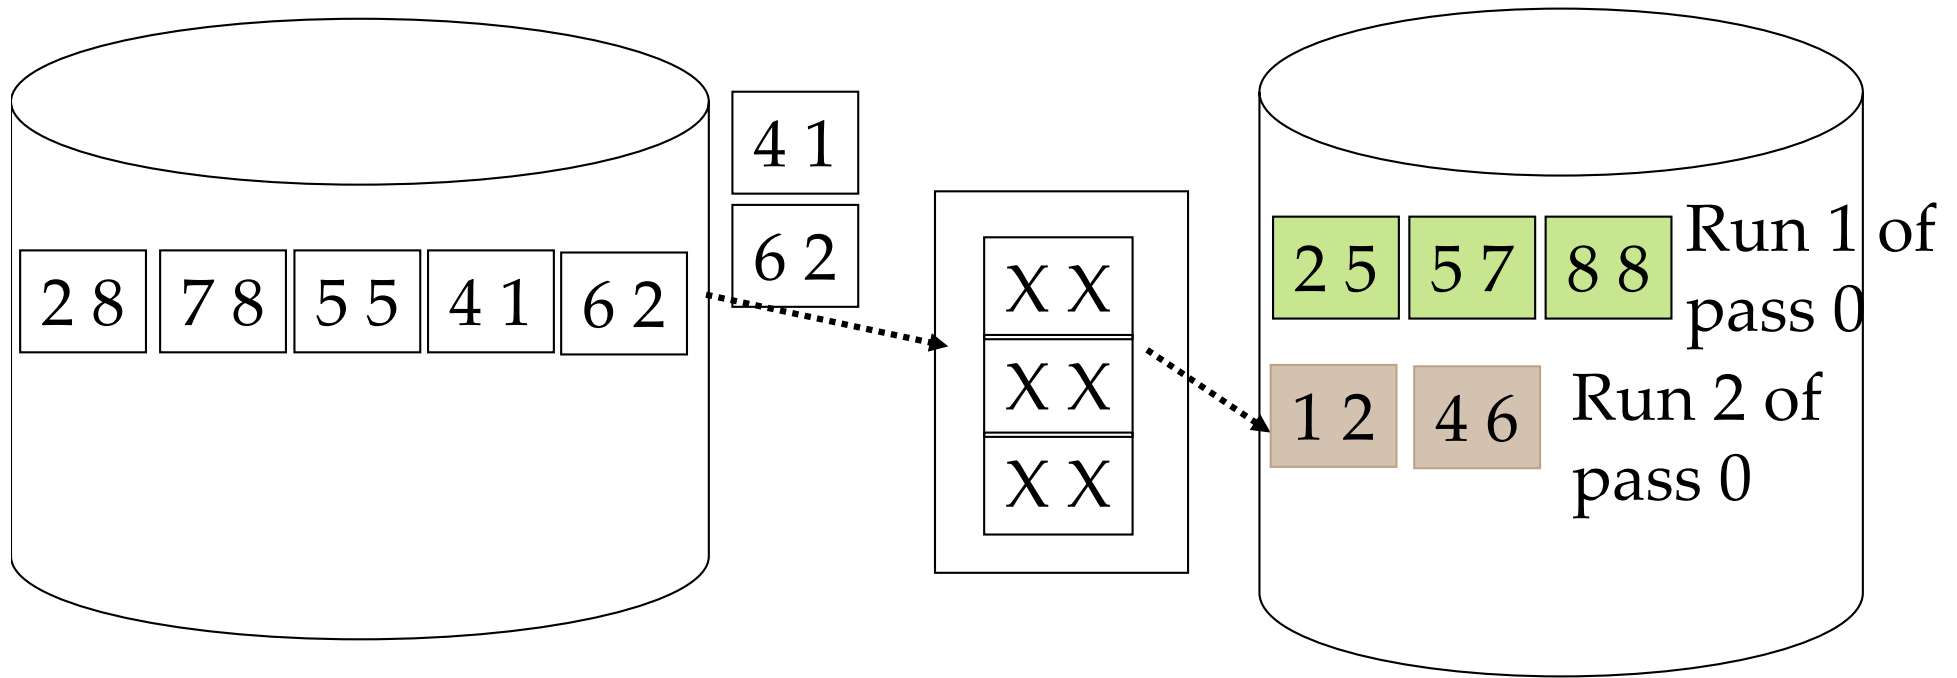
What if only 3 buffer frames available?
That is, data pages do not
fit into main memory

External Sorting - Pass 0




- In example: 5 pages, 3 buffer frames
- N pages, B buffer frames:
 - Bring B pages in buffer
 - Sort with any main memory sort
 - Write out to disk to a temporary file;
 - it's called a run of B pages

External Sorting - Pass 0

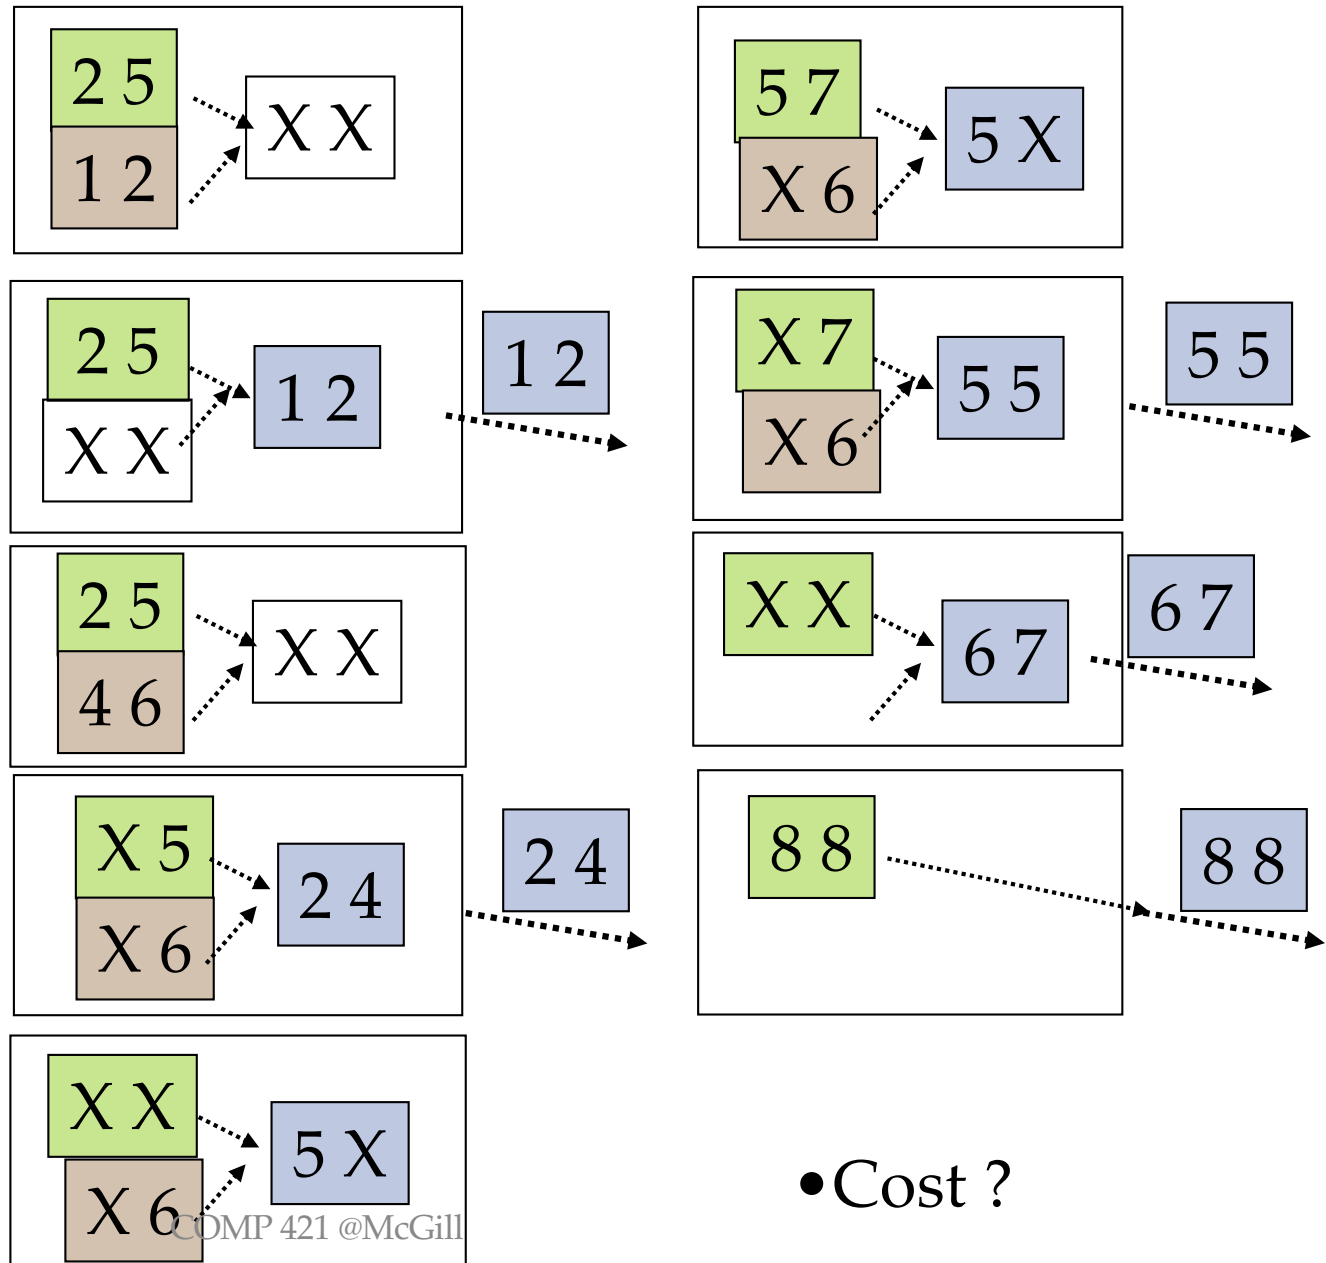


- In example: $5/3 = 2$ runs (round up!)
- In total N/B runs
- Cost ?

External Sorting Pass 1:



- N/B input frames
- One output frame
- Merge Sort the runs of pass 0
- Hopefully $B-1 > N/B$
- Otherwise pass 2...



• Cost ?

Sort

- Sometimes a Pass 2 is needed:
 - Pass 0 created more runs than there are main memory buffers
 - Therefore Pass 1 produces more than one run
 - Take the first $B-1$ runs from pass 0 and merge them to one bigger Pass 1 run
 - Then take the next $B-1$ runs from pass 0 and merge them to one bigger Pass 1 run
 - ...
 - Pass 2 takes the runs of Pass 1 and merges them
 - If there are less than B Pass 1 runs, then this is the final pass
 - Otherwise Pass 3...
- Number of passes: $1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$
- Cost = $2N * (\text{\# of passes})$

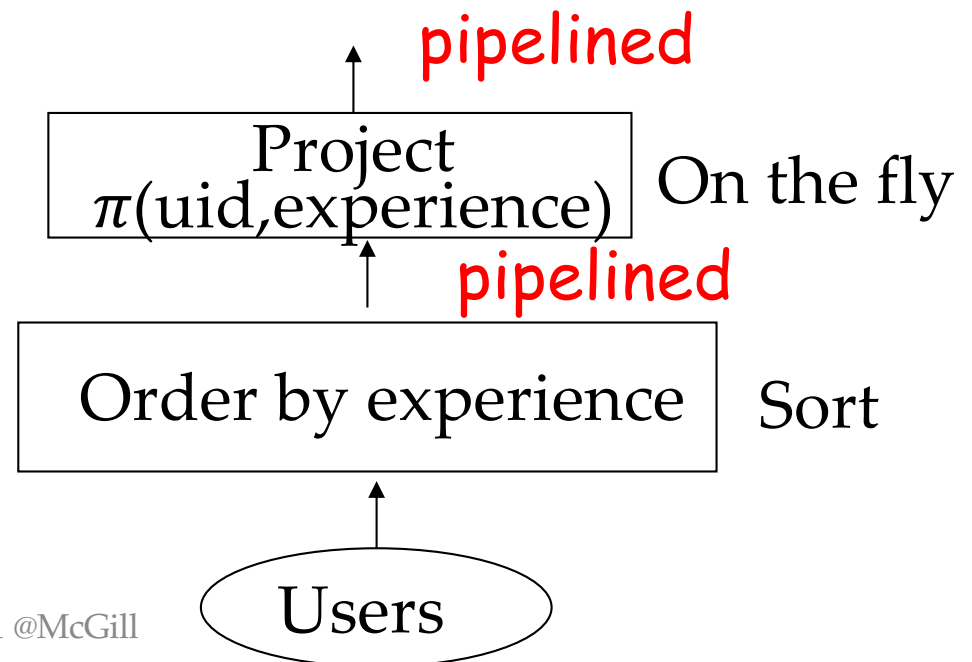
Sort Costs

Number of passes

| N | B=3 | B=5 | B=9 | B=17 | B=129 | B=257 |
|---------------|-----|-----|-----|------|-------|-------|
| 100 | 7 | 4 | 3 | 2 | 1 | 1 |
| 1,000 | 10 | 5 | 4 | 3 | 2 | 2 |
| 10,000 | 13 | 7 | 5 | 4 | 2 | 2 |
| 100,000 | 17 | 9 | 6 | 5 | 3 | 3 |
| 1,000,000 | 20 | 10 | 7 | 5 | 3 | 3 |
| 10,000,000 | 23 | 12 | 8 | 6 | 4 | 3 |
| 100,000,000 | 26 | 14 | 9 | 7 | 4 | 4 |
| 1,000,000,000 | 30 | 15 | 10 | 8 | 5 | 4 |

Sort together with other operators

- I/O Costs:
 - SELECT uname, experience FROM Users ORDER BY experience
 - If everything fits into main memory (Only pass 0 needed):
 - Read number of data pages
 - sort and pipeline result into next operator: $\pi(\text{uname}, \text{experience})$
 - Pass 0 + pass 1 needed
 - Pass 0: read # pages, write # pages (have to write temp. results!)
 - Pass 1: read # pages, sort and pipeline result into next operator
 - $3 * \text{\#pages}$
 - Pass 0 + pass 1 + pass 2 needed
 - $5 * \text{\#pages}$



Sort in real life

- Blocked I/O:
 - use more output pages and flush to consecutive blocks on disk
 - Might lead to more I/O but each I/O is cheaper
- Other optimizations:
 - At write in Pass 0 to disk (if needed):
 - Do projection on necessary attributes → each tuple is smaller → less pages
 - that is, projection is pushed to the lowest level possible

Equality Joins

```
SELECT  *  
FROM    Users U, GroupMembers GM  
WHERE   U.uid = GM.uid
```

| p1 _u | uid | uname | experience | age |
|-----------------|-----|--------|------------|-----|
| | 123 | Dora | 2 | 13 |
| p2 _u | 132 | Bug | 8 | 60 |
| | 267 | Sakura | 7 | 15 |
| | 111 | Cyphon | 8 | 35 |

| p1 _g | uid | gid | stars |
|-----------------|-----|-----|-------|
| | 123 | G1 | 2 |
| p2 _g | 132 | G1 | 5 |
| | 132 | G2 | 3 |
| p3 _g | 132 | G3 | 1 |
| | 123 | G2 | 4 |
| | 111 | G4 | 2 |

Join Cardinality Estimation

- $| \text{Users} \bowtie \text{GroupMembers} | = ?$
 - Join attribute is primary key for Users
 - Each GroupMember tuple matches exactly with one Users tuple
 - Result: $|\text{GroupMembers}|$
- $| \text{Users} \times \text{GroupMembers} | = ?$
 - Result: $|\text{Users}| * |\text{GroupMembers}|$
 - Cross product is always the product of individual relation sizes
- For other joins more difficult to estimate

Cardinality Estimation

- $| \text{Users} \bowtie \sigma_{(\text{stars} > 3)}(\text{GroupMembers}) | = ?$
 - Result: $| \sigma_{(\text{stars} > 3)}(\text{GroupMembers}) |$
 - Assuming 1-5 stars and 100,000 members:
 - 40,000
- $| \sigma_{(\text{experience} > 5)}(\text{Users}) \bowtie (\text{GroupMembers}) | = ?$
 - Assume 1-10 experience levels and 40,000 users and uniform distribution for experience
 - $\text{Red}(\sigma_{(\text{experience} > 5)}(\text{Users})) = 1/2$
 - Result: $1/2 * |\text{GroupMembers}|$

Simple Nested Loop Join

- For each tuple in the *outer* relation Users U we scan the entire *inner* relation GroupMembers GM.

foreach tuple u in U do

foreach tuple g in GM do

if u.uid == g.uid then add <u, g> to result

| | <u>uid</u> | uname | experience | age |
|-----------------|------------|--------|------------|-----|
| p1 _u | 123 | Dora | 2 | 13 |
| | 132 | Bug | 8 | 60 |
| p2 _u | 267 | Sakura | 7 | 15 |
| | 111 | Cyphon | 8 | 35 |

| | <u>uid</u> | <u>gid</u> | stars |
|-----------------|------------|------------|-------|
| p1 _g | 123 | G1 | 2 |
| | 132 | G1 | 5 |
| p2 _g | 132 | G2 | 3 |
| | 132 | G3 | 1 |
| p3 _g | 123 | G2 | 4 |
| | 111 | G4 | 2 |

Simple Nested Loop Join

- For each tuple in the *outer* relation Users U we scan the entire *inner* relation GroupMembers GM.

```
foreach tuple u in U do
  foreach tuple g in GM do
    if u.uid == g.uid then add <u, g> to result
```

- Cost: $\text{UserPages} + |\text{Users}| * \text{GroupMemberPages} = 500 + 40,000 * 1000$!
- NOT GOOD
- We need page-oriented algorithm!

Page Nested Loop Join

- For each page p_u of Users U , get each page p_g of GroupMembers GM
 - write out matching pairs $\langle u, g \rangle$, where u is in p_u and g is in p_g .

```
For each page  $p_u$  of Users  $U$ 
  for each page  $p_g$  of GroupMembers  $GM$ 
    for each tuple  $u$  in  $p_u$  do
      for each tuple  $g$  in  $p_g$  do
        if  $u.uid == g.uid$  then add  $\langle u, g \rangle$  to result
```

| | <u>uid</u> | uname | experience | age |
|--------|------------|--------|------------|-----|
| $p1_u$ | 123 | Dora | 2 | 13 |
| | 132 | Bug | 8 | 60 |
| $p2_u$ | 267 | Sakura | 7 | 15 |
| | 111 | Cyphon | 8 | 35 |

| | <u>uid</u> | <u>gid</u> | stars |
|--------|------------|------------|-------|
| $p1_g$ | 123 | G1 | 2 |
| | 132 | G1 | 5 |
| $p2_g$ | 132 | G2 | 3 |
| | 132 | G3 | 1 |
| $p3_g$ | 123 | G2 | 4 |
| | 111 | G4 | 2 |

Page Nested Loop Join

- For each *page* p_u of Users U , get each *page* p_g of GroupMembers GM
 - write out matching pairs $\langle u, g \rangle$, where u is in p_u and g is in p_g .

```
For each page  $p_u$  of Users  $U$ 
  for each page  $p_g$  of GroupMembers  $GM$ 
    for each tuple  $u$  in  $p_u$  do
      for each tuple  $g$  in  $p_g$  do
        if  $u.uid == g.uid$  then add  $\langle u, g \rangle$  to result
```

- Cost: $UserPages + UserPages * GroupPages = 500 + 500 * 1000 = 500,500$

Block Nested Loop Join

- For each *block of pages* bp_u of Users U , get each *page* p_g of GroupMembers GM
 - write out matching pairs $\langle u, g \rangle$, where u is in bp_u and g is in p_g .
- *block of pages* bp_u and one page of GM must fit in main memory
 - For each block of pages bp_u
 - Load block into main memory
 - Get first page from GM
 - Do all the matching between users in bp_u and group members in first page
 - Get second page from GM (into the same frame the first one was in before)
 - Do all the the matching between users in bp_u and group members in second page
 - ...
 - Get last page from GM (into again that frame reserved for GM)
 - ...
- Cost: $UserPages + UserPages / |bp_u| * GroupMemberPages$

Block Nested Loop

| | uid | uname | experience | age |
|--------|-----|--------|------------|-----|
| $p1_u$ | 123 | Dora | 2 | 13 |
| | 132 | Bug | 8 | 60 |
| $p2_u$ | 267 | Sakura | 7 | 15 |
| | 111 | Cyphon | 8 | 35 |

| | uid | gid | stars |
|--------|-----|-----|-------|
| $p1_g$ | 123 | G1 | 2 |
| | 132 | G1 | 5 |
| $p2_g$ | 132 | G2 | 3 |
| | 132 | G3 | 1 |
| $p3_g$ | 123 | G2 | 4 |
| | 111 | G4 | 2 |

Block Nested Loop Join

- Examples depending on available main memory:
- 51 Buffer Frames:
 - $500 + 500/50 * 1000 = 500 + 10,000$
- 501 Buffer Frames
 - $500 + 500/500 * 1000 = 500 + 1000$
 - Special case: outer relation fits into main memory!!

Index Nested Loops Join

- For each tuple in the *outer* relation Users U we find the matching tuples in GroupMembers GM through an index
 - Condition: GM must have an index on the join attribute

foreach tuple u in U do

find all matching tuples g in GM through index

then add all <u, g> to result

| | <u>uid</u> | uname | experience | age |
|-----------------|------------|--------|------------|-----|
| p1 _u | 123 | Dora | 2 | 13 |
| | 132 | Bug | 8 | 60 |
| p2 _u | 267 | Sakura | 7 | 15 |
| | 111 | Cyphon | 8 | 35 |

| | <u>uid</u> | <u>gid</u> | stars |
|-----------------|------------|------------|-------|
| p1 _g | 123 | G1 | 2 |
| | 132 | G1 | 5 |
| p2 _g | 132 | G2 | 3 |
| | 132 | G3 | 1 |
| p3 _g | 123 | G2 | 4 |
| | 111 | G4 | 2 |

Index Nested Loops Join

foreach tuple u in U do

find all matching tuples g in GM through index

then add all $\langle u, g \rangle$ to result

- Index MUST be on the inner relation (in this case GM).
- Cost: $\text{OuterPages} + \text{CARD}(\text{OuterRelation}) * \text{cost of finding matching tuples in inner relation}$
- In example of previous page:
 - Index on uid on GM is clustered:
 - $500 + 40.000 * (1 \text{ leaf page} + 1 \text{ data pages})$
 - Index on uid on GM is not clustered:
 - $500 + 40.000 * (1 \text{ leaf page} + 2.5 \text{ data pages})$ (on average 2.5 tuples in GM per user)

Index Nested Loops Join

- Switch inner and outer if index is on uid of Users
- Note: uid is primary key in User

– Only one tuple matches!

foreach tuple g in GM do

find the one matching tuple u in U through index

then add <g, u> to result


- Cost: $1000 + 100.000 * (1 \text{ leaf page} + 1 \text{ data page})$

| | <u>uid</u> | <u>gid</u> | stars |
|-----------------|------------|------------|-------|
| p1 _g | 123 | G1 | 2 |
| | 132 | G1 | 5 |
| p2 _g | 132 | G2 | 3 |
| | 132 | G3 | 1 |
| p3 _g | 123 | G2 | 4 |
| | 111 | G4 | 2 |

| | <u>uid</u> | uname | experience | age |
|-----------------|------------|--------|------------|-----|
| p1 _u | 123 | Dora | 2 | 13 |
| | 132 | Bug | 8 | 60 |
| p2 _u | 267 | Sakura | 7 | 15 |
| | 111 | Cyphon | 8 | 35 |

Block Nested Loop vs. Index

- Best Case for Block Nested Loop (if outer relation fits in main memory)
 - OuterPages + InnerPages
- Index Nested Loop:
 - OuterPages + Card(Outer) * matching tuples Inner
- Index Nested Loop wins if:
 - InnerPages > Card(Outer) * matching tuples Inner
 - E.g., if Outer is the result of a selection that only selected very few tuples
 - $\sigma_{(\text{age} > 60)}(\text{Users}) \bowtie (\text{GroupMembers})$



Sort-Merge Join

- Sort U and GM on the join column, then scan them to do a “merge” (on join col.), and output result tuples.
 - In loop:
 - Assume the scan cursors currently points to U tuple u and GM tuple g . Advance scan cursor of U until $u.uid \geq g.uid$ and then advance scan cursor of GM so that $g.uid \geq u.uid$. Do this until $u.uid = g.uid$.
 - At this point, all U tuples with same value in uid (*current U group*) and all GM tuples with same value in uid (*current GM group*) match; output $\langle u, g \rangle$ for all pairs of such tuples.
 - Then resume scanning U and GM.
- U is scanned once; each GM group is scanned once per matching U tuple. (Multiple scans of an GM group are likely to find needed pages in buffer.)

Example of Sort-Merge Join

| <u>uid</u> | uname | experience | age |
|------------|--------|------------|-----|
| 111 | Cyphon | 8 | 35 |
| 123 | Dora | 2 | 13 |
| 132 | Bug | 8 | 60 |
| 267 | Sakura | 7 | 15 |

| <u>uid</u> | <u>gid</u> | stars |
|------------|------------|-------|
| 111 | G4 | 2 |
| 123 | G1 | 2 |
| 123 | G2 | 4 |
| 132 | G1 | 5 |
| 132 | G2 | 3 |
| 132 | G3 | 1 |

Cost of Sort-Merge Join

- Relations are already sorted:
 - $\text{UserPages} + \text{GroupMemberPages} = 500 + 1000$
- Relations need to be sorted – simple way:
 - Sort relations and write sorted relations to temporary stable storage
 - Read in sorted relations and merge
 - Costs: assuming 100 buffer pages
 - both Users and GroupMembers can be sorted in 2 passes (Pass 0 and 1): $4 * \text{UserPages} + 4 * \text{GroupPages}$
 - Final merge: $500 + 1000 = (\text{UserPages} + \text{GroupPages})$
 - Total: $5 * \text{UserPages} + 5 * \text{GroupPages}$


Cost of Sort-Merge Join

- Relations need to be sorted – use pipelining to combine last sort pass and join
 - Sorting performs Pass 0 reads and writes each of the relations: $2 * \text{UserPages} + 2 \text{ GroupPages}$
 - Pass 1 reads data, sorts and then performs merge in pipeline fashion (ignore details): $\text{UserPages} + \text{GroupPages}$
 - Total: $3 * \text{UserPages} + 3 * \text{GroupPages} = 4,500$

Hash Join: first step

- Partition both relations using hash fn **h** (that returns a value between 1 and B-1 (if there are B buffer frames):
 - $h(u.uid) = i \rightarrow$ tuple u of User is in UserPartition i .
 - $h(g.uid) = i \rightarrow$ tuple g of GroupMember is in GroupMemberPartition i .
- Partitioning algorithm for relation U

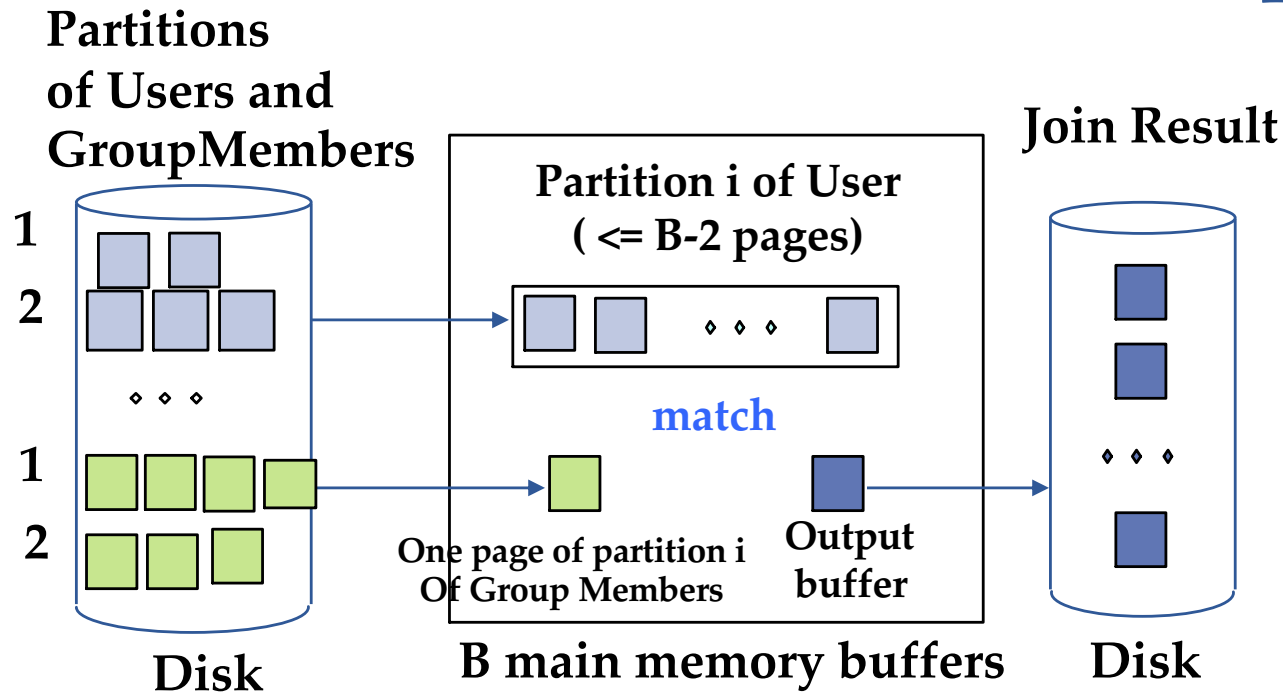
```
For each page of Users U
  for each tuple u in page do
    append u to UserPartition h(u.uid)
```



Hash Join: assumptions

- Assumption for this course:
 - Each partition of User fits into main memory and there are still 2 more buffer frames available
 - Each partition of User smaller than $B-2$
 - This holds if
 - Hash function creates partitions of equal size
 - Partition size for Users is $\text{UserPages} / (B-1) = 500 / B-1$
 - Partition size for GroupMembers is $\text{GroupMemberPages} / (B-1) = 1000 / B-1$
 - $500 / B-1$ (rounded up) $\leq B-2$
 - That is, buffer has at least 24 buffer frames ($B=24$).

Hash-Join: second step



- For u of User, g of GroupMember
 - $u.uid == g.uid$ and u in UserPartition $i \rightarrow g$ in GroupMemberPartition i
 - Thus, for each i , join UserPartition i with GroupMemberPartition i only.
- For each i , $1 < \dots i < \dots$ Number of partitions**
- Load partition i of Users into main memory**
- For each page of partition i of GroupMembers**
- Load page into main memory**
- write all matching tuples (u, g) with $u.uid = g.uid$ to output**

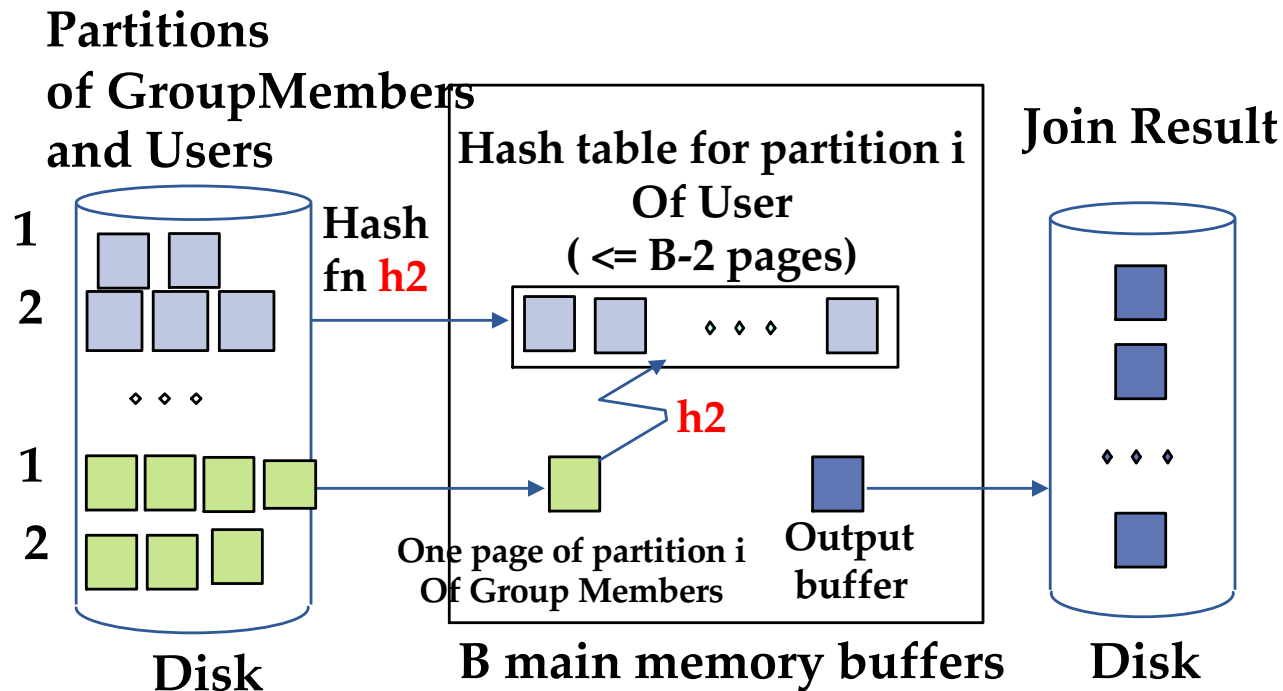
Hash Join: Costs

- Step I:
 - Read and write Users (partition U)
 - Read and write GroupMembers (partition GM)
 - $2 * \text{UserPages} + 2 * \text{GroupMemberPages} = 1000 + 2000$
- Step II:
 - Read Users (partition by partition)
 - Read GroupMembers (page by page in partition order)
 - $\text{UserPages} + \text{GroupPages} = 500 + 1000$
- $3 * \text{Userpages} + 3 * \text{GroupPages}$

Merge-Join vs. Hash-Join

- In our example:
 - Both $3 * \text{UserPages} + 3 * \text{GroupPages}$
- Hash-Join better
 - if one relation really large
 - (sort in merge-join might need another pass)
- Merge Join better
 - if hash partitions of smaller don't fit into main memory
- Merge Join result is sorted
- Hash join good for parallelization
 - Let each pair of partitions be matched on a different node

Hash-join: CPU optimized



- Reorganize partition i of User
 - Partition according to a second hash-function
 - $h2(u.uid) = j \rightarrow u$ is put on page j
- Match tuples of GroupMembers
 - $h2(g.uid) = j \rightarrow$ matching u tuple must be on page j
- Avoids scanning the entire partition of U for each tuple of Group Members

The Projection Operation

```
SELECT  GM.uid, GM.gid  
FROM    GroupMembers GM
```

- Usually done on the fly together with another operation (or pipelined)
 - For instance, while reading in the transaction for a sort, or join etc.
- More complex: **SELECT DISTINCT** name FROM ...
 - Expensive operation
 - Requires sort in order to eliminate duplicates!!
 - Often done at the very end (less tuples) or whenever the relation is sorted for some other reason
 - Database user: use DISTINCT only when really necessary

Set Operations

- Intersection and cross-product special cases of join.
- Union (Distinct) and Except similar;
- For instance: sorting based approach to union:
 - Sort both relations (on combination of all attributes).
 - Scan sorted relations and merge them.

Sort → scan → merge

Aggregate Operations (AVG, MIN, etc.)

- Usually done at the very last step after all selections/joins etc.
- Without grouping:
 - In general, requires scanning the relation.
- With grouping:
 - Sort on group-by attributes, then scan relation and compute aggregate for each group. (Can improve upon this by combining sorting and aggregate computation.)

Execution Plan

- A plan describes how a query is executed
- A Tree (sequence) of basic operators (select, join, project, sort, etc) used to process the query
- For each operator, an indication how it will be executed (index nested loop, sort, index, simple scan....)

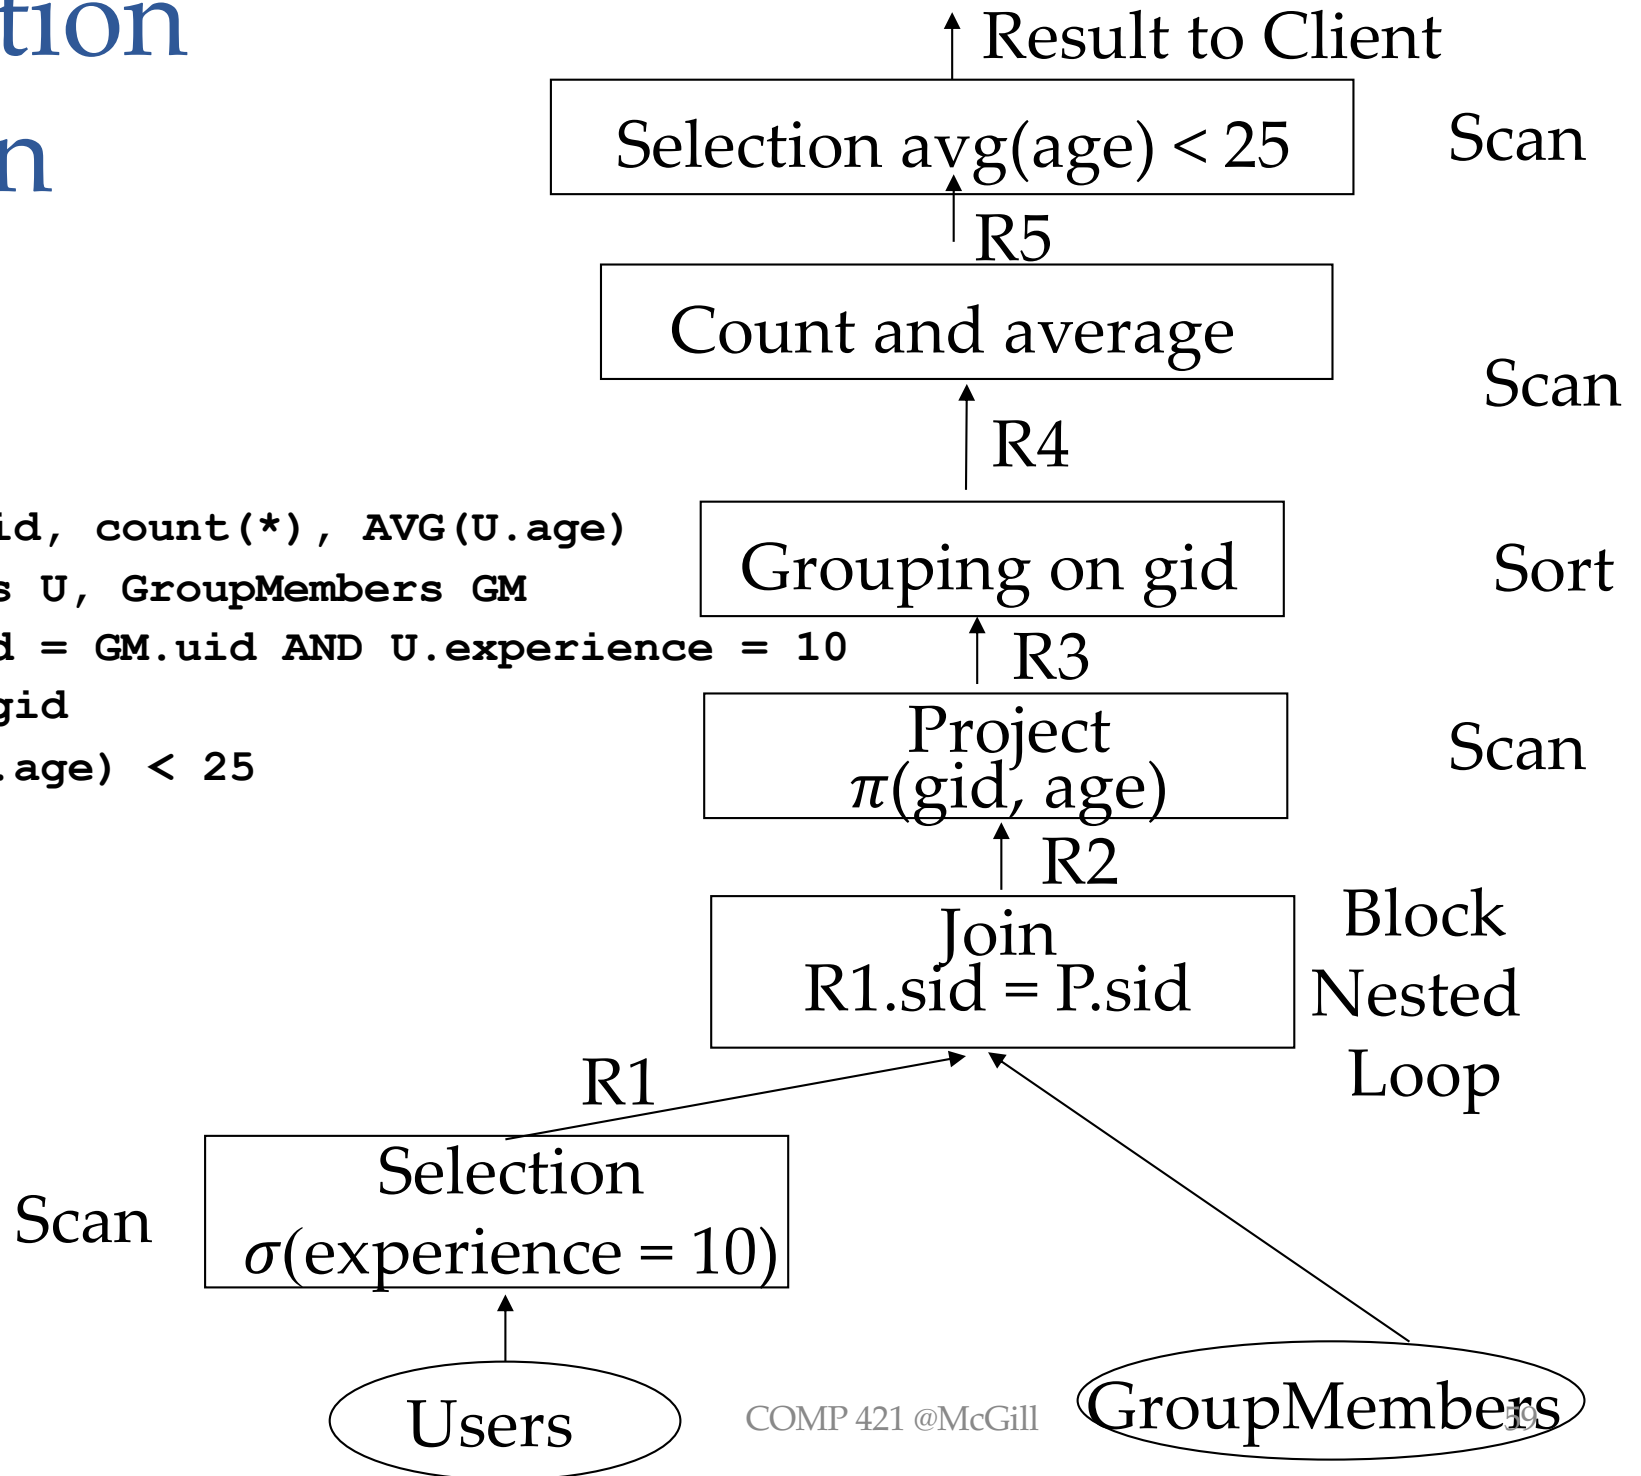
Example: Step by step

```
SELECT  GM.gid, count(*), AVG(U.age)
FROM    Users U, GroupMembers GM
WHERE   U.uid = GM.uid AND U.experience = 10
GROUP BY GM.gid
HAVING  AVG(U.age) < 25
```

- Scan the Users table, determine all tuples with experience = 10 (scan or index)
 - Results in intermediate relation R1
- Join R1 and GroupMembers (several join options)
 - Results in intermediate relation R2
- Eliminate attributes other than gid and age.
 - Results in intermediate relation R3
- Group the tuples of R3 on gid (done by sort).
 - Results in intermediate relation R4
- Scan R4, count and perform average
 - Results in intermediate relation R5
- Return all tuples with $\text{avg}(\text{age}) < 25$.

Execution Plan

```
SELECT  GM.gid, count(*), AVG(U.age)
FROM    Users U, GroupMembers GM
WHERE   U.uid = GM.uid AND U.experience = 10
GROUP BY GM.gid
HAVING  AVG(U.age) < 25
```



Pipelining

- Execution within one operator:
 - As soon as a tuple is determined it is forwarded to next operator
- Parallel execution of operators
- No materialization of intermediate relations if it can be avoided
- Iterator operator uses pipelining
 - Every operator is an iterator
 - Iterator provides the following interface to “parent” operators: **open, getNext, close**
 - Iterator calls interface methods of “children” operators

iterator based systems


Execution Plan

need everything grouped here →

Pipelined: fill a buffer frame and then start join

Selection
 $\sigma(\text{experience} = 10)$

Scan



Join
 $R1.sid = P.sid$

Block
Nested
Loop

Project
 $\pi(gid, age)$

Pipelined

On the fly

Pipelined

Sort

Grouping on gid

↑ R3

Count and average

Pipelined
Scan

Selection $\text{avg}(age) < 25$

↑ Result to Client **Pipelined**

Scan



Equality Join with further restrictions

```
SELECT  *
FROM    Users U, GroupMembers GM
WHERE   U.uid = GM.uid
AND     experience = 10
```

*Making the
outer relation as
small as possible
is good*

- (assume 1% have experience = 10)
- Do selection before join;
- Stream (pipeline) qualifying tuples into join
- Block nested loop (get qualifying tuples until block is full)
 - $\text{UserPages} + \boxed{\text{UserPages} * 0.01 / (B-2)} * \text{GroupMemberPages} = 500 + 1 * 1000 = 1,500$
- Index nested loop
 - $\text{UserPages} + \text{Card}(\text{Users}) * 0.01 * \text{Cost finding GroupMembers} = 500 + 400 * (1 + 2.5) = 500 + 1400 = 1900$

Equality Join with further restrictions

```
SELECT  *  
FROM    Users U, GroupMembers GM  
WHERE   U.uid = GM.uid  
AND     experience = 10
```

- Sort-Merge Join
 - Pipelining cannot be done as sort needed; therefore intermediate relation
 - But in particular case: result of selection fits into main memory; therefore sort of user tuples in main memory
 - Cost: $\text{UserPages} + 3 * \text{GroupPages} = 3500$
- Hash Join
 - As Users after selection fits into main memory, it becomes kind of a hash join with one partition.
 - Thus, I don't need to partition GroupMembers, and hash join becomes identical to Block Nested Loop

63 $500 + 1000$

Two types

Optimization Techniques

- **Algebraic optimization:**

- Use simple rules to perform those operations first that eliminate a lot of tuples
- • Push down selections and projections ←
– build a basic operator tree
- Do not yet consider HOW to execute each operator
- Consider the number of tuples that flow from one operator to the next
- Key issues: statistics

implementations

- **Cost-based optimizations**

brute force

- Consider a set of alternative plans created by algebraic optimization
- Consider for each operator how it could be executed

64

- Key issues: available indexes, available operator implementations

execute operation in an order such that most appropriate access paths (indexes) can be used

Algebraic

Projection, Selection and Join

- ❑ Pushing down selections and projections to the relations to which selection refers
- ❑ Careful with project

```
SELECT u.uname, u.experience
FROM   Users u, GroupMembers g
WHERE  u.uid = g.uid
AND    g.gid = 'G1'
AND    u.age > 50
```

- $\pi_{\text{uname}, \text{experience}}(\sigma_{\text{gid}=\text{G1} \wedge \text{age} > 50}(\text{Users} \bowtie \text{GroupMembers}))$
 - first join, then selection, then project
- $\pi_{\text{uname}, \text{experience}}(\sigma_{\text{age} > 50}(\text{Users}) \bowtie \sigma_{\text{gid}=\text{G1}}(\text{GroupMembers}))$
 - push down selections
- $\pi_{\text{uname}, \text{experience}}(\sigma_{\text{age} > 50}(\text{Users}) \bowtie \pi_{\text{uid}}(\sigma_{\text{gid}=\text{G1}}(\text{GroupMembers})))$
 - push down SOME project

Push Projections

- Pushing down projections will not reduce the number of tuples but the SIZE of the intermediate results
- Be careful not to lose attributes that you need later on
- $\pi_{\text{uname}}(\text{Users} \bowtie \text{GroupMembers}) \neq \pi_{\text{uname}}(\text{Users}) \bowtie \text{GroupMembers}$

Algebraic Optimization (contd)

SELECT S.sname, P.cid

FROM Skaters S, Participates P, Competitions C

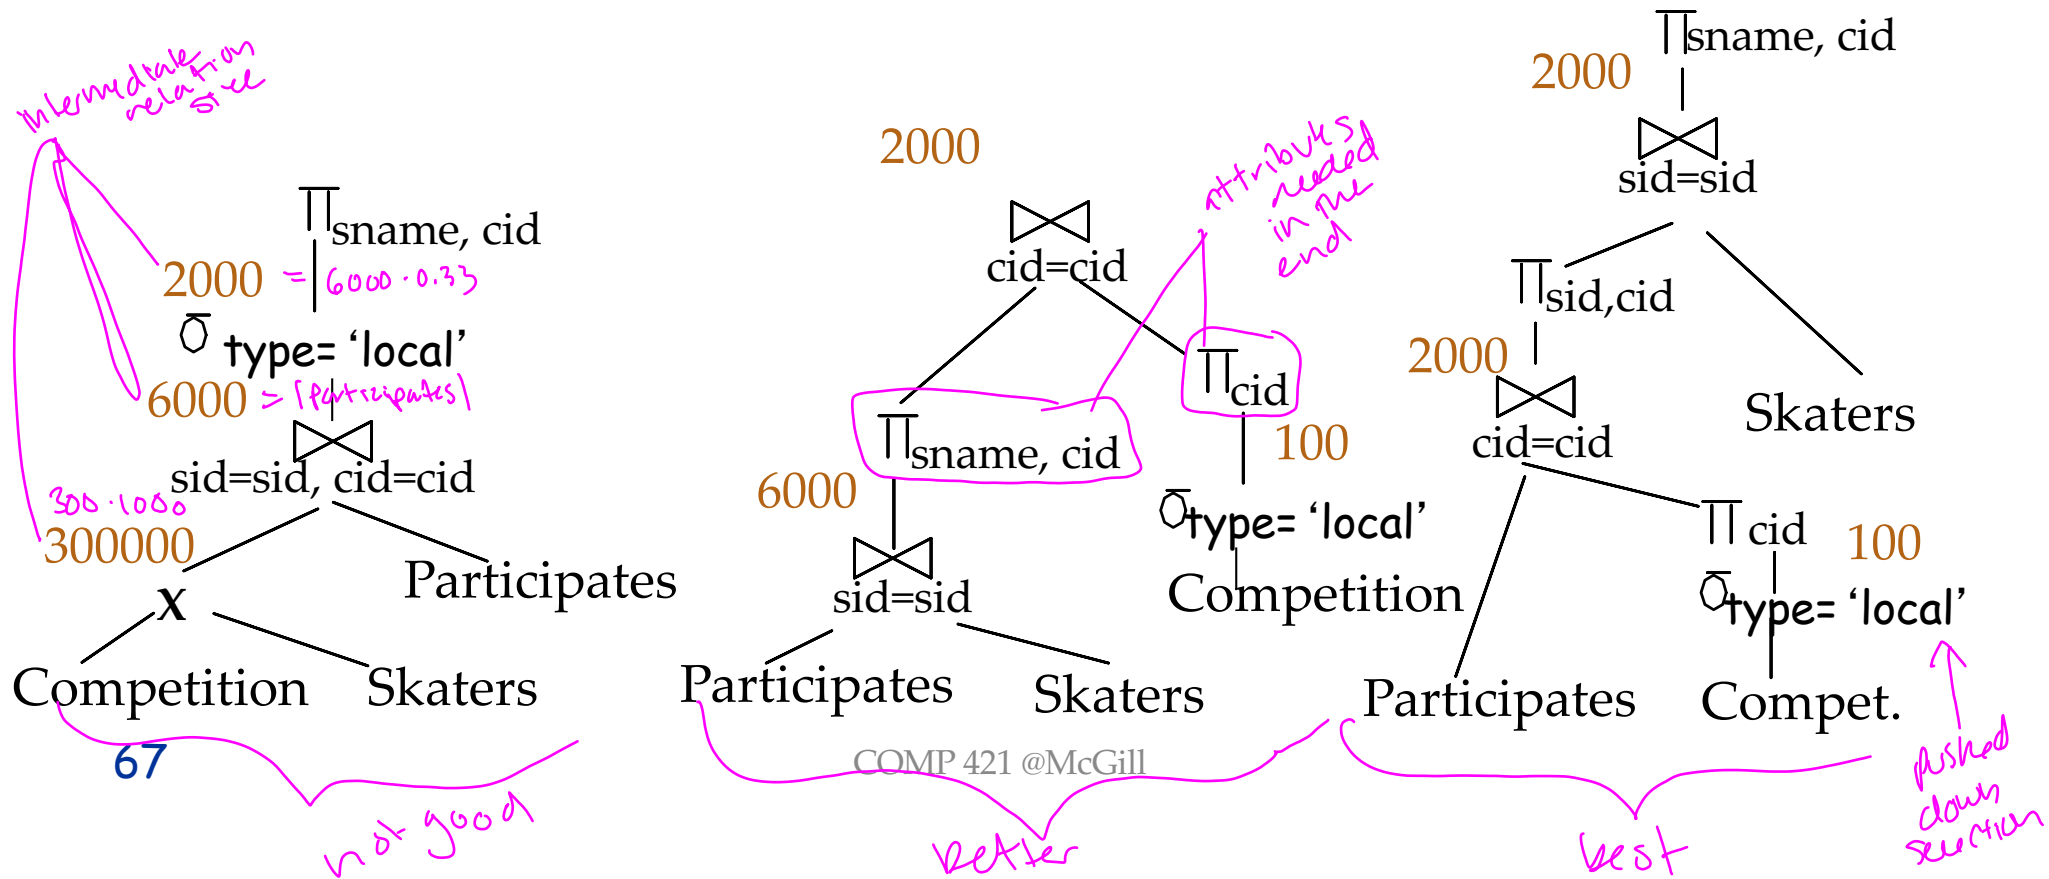
WHERE P.sid=S.sid AND P.cid=C.cid AND C.type= 'local' ;

Skaters (sid: 4 Byte, uname (10 Byte), age (4Byte)): 1000 tuples

Participates (sid: 4 Byte, cid: 4 Byte, rank (4 Byte): 6000 tuples

Competition (cid: 4 Byte, location (10 Byte), type (5 Byte): 300 tuples

type attribute: 3 different values *reduction factor: 33%*



Cost Based Optimization

- Find a plan with low cost
- Dynamic programming in bottom-up fashion for deep join plans :
 - – Pass 1: Find best 1-relation plan for each relation.
 - – Pass 2: Find best way to join result of each 1-relation plan (as outer) to another relation. (All 2-relation plans.) *consider inner and outer*
 - – Pass N: Find best way to join result of a (N-1)-relation plan (as outer) to the N' th relation. (All N-relation plans.)
 - Prune high-costs alternatives
 - Retain alternatives with interesting features
 - Cheapest plan overall, plus ←
 - Cheapest plan for each *interesting order* of the tuples. ←
 - Bottom-up calculation each possible execution plan, **estimate the cost**
 - In class: something between algebraic and cost-based optimization

*example in my
courses
slides*

Nested Queries

- ❑ Select S.sname
From Skaters S
Where S.sid IN (Select P.sid From Participates P
Where P.cid = 103)
 - ☆ Execute subquery first; returns intermediate relation, with distinct sid
 - ☆ Join outer relation and intermediate relation with any join method
- ❑ Select S.sname
From Skaters S
Where exists (Select * from Participates P
where P.cid = 103 and P.sid = S.sid)
 - ☆ Have to execute inner query for each outer tuple; little optimization possible
- ❑ Select S.sname
From Skaters S, Participates P
Where S.sid = P.sid AND P.cid = 103
 - ☆ all optimizations possible
- ❑ smart rewriting system able to transform queries automatically

database
transforming
to this

Why did we learn all this

System Tuning

- ❑ If your application has some standard, well-known queries:
 - ☆ Create appropriate indices to speed up these queries
- ❑ If your application has many many updates and inserts
 - ☆ Be careful with creating indices
 - ☆ Each INSERT will insert new value in each tuple
 - ☆ Updates change some indices
- ❑ SQL **Explain** explains how a query is executed internally

Statistics in DB2

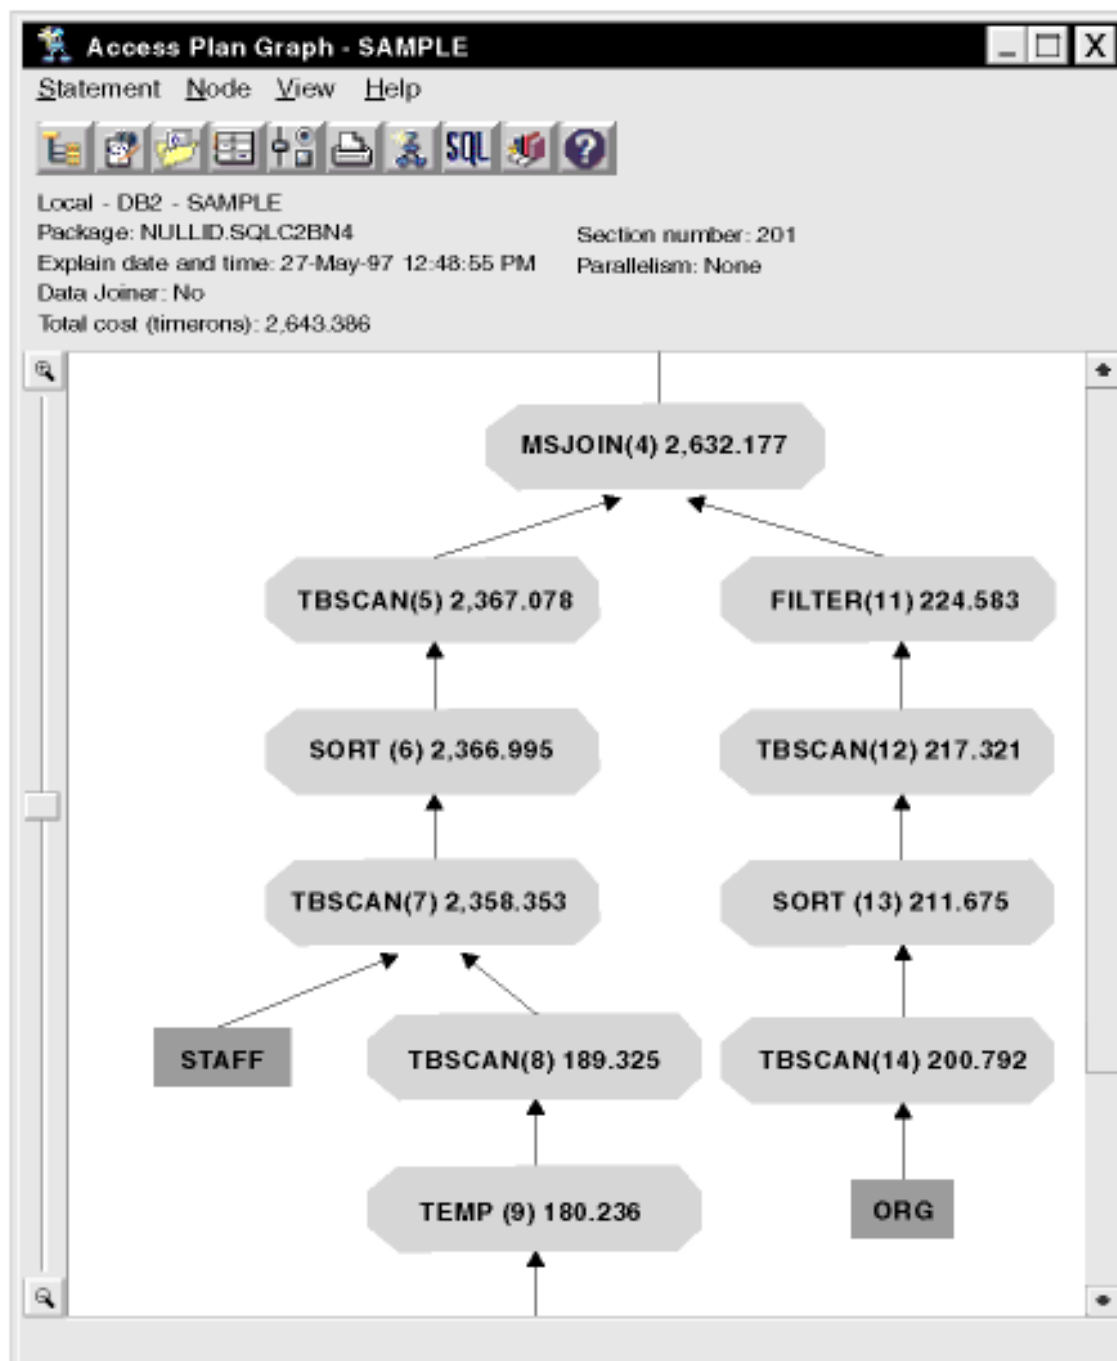
❑ Start **Runstat** to collect information about tables :

★ **SYSCAT.TABLES** and **SYSSTAT.TABLES**:

- Number of pages per table, number of tuples, etc.

★ **SYSCAT.INDEXES** and **SYSSTAT.INDEXES**

- Number of index leaf pages, index levels, degree of clustering, number of distinct values in first column, page fetch estimate for different buffer size, etc.



Operator details - MSJOIN[4]

Local - DB2 - SAMPLE

Level of details: ☐ Overview ☒ Full

| | | |
|-----------------------|------------------------------|------------------------------|
| Cumulative cost | | |
| Total cost | 2,632.177 timerons | |
| CPU cost | 55,853.416 instructions | |
| I/O cost | 15.921 I/Os | |
| First row cost | 2,587.959 timerons | |
| | | |
| Cumulative properties | | |
| Tables | DENISEW.STAFF DENISEW.ORG | |
| | | |
| Input arguments | | |
| Join predicates | Number 3 | Selectivity 0.0024999999 |
| | | Text (Q5.DEPTNUM=Q4.DEPT) |

Save As...

Print...

Close

Help