

declarative languages

Large Scale Data Processing

Adaptation from

Magdalena Balazinska (Univ. of Washington)
Mining of Massive Datasets, by Rajaraman and Ullman
Alan Gates (Yahoo!)
Olston

Declarative Languages

- Map-Reduce framework hides scheduling and parallelization details
- Limited query expressiveness
 - Complex queries difficult to write
- Declarative languages on top of map-reduce
 - Pig Latin (Yahoo!)
 - Like relational algebra
 - Open source
 - HiveQL (Facebook)
 - SQL like language
 - Open source
 - SQL / Tenzing
 - proprietary

Pig (Latin)

- Pig Latin *english based*
 - A higher SQL like language to run complex queries that require several map-reduce jobs
- **Pig**
 - An execution engine
 - Translates Pig Latin programs into graphs of map-reduce jobs
 - Executes them on top of Hadoop
 - An Apache open source project

Example (Alan Gates, Yahoo)

users(name, age), pagelog(url, uname)

Find the top 5 most popular pages for users aged 18-25

SQL:

```
SELECT url, count(*) as clicks
```

```
FROM users U, pagelog P
```

```
WHERE U.name = P.uname
```

```
AND U.age >= 18
```

```
AND U.age <= 25
```

```
GROUP BY URL
```

```
ORDER BY clicks desc
```

```
LIMIT 5
```

```
-- FETCH FIRST 5 ROWS ONLY
```

Map reduce program: 170 lines of code

red = key words
blue = file/input

In Pig Latin

Users = **load** 'users' **as** (name,age); → load file into variable

Filtrd = **filter** Users **by** age >= 18 **and** age <= 25; selection

Pages = **load** 'pages' **as** (uname, url);

Jnd = **join** Filtrd **by** name, Pages **by** uname;

Grpd = **group** Jnd **by** url;

Smmd = **foreach** Grpd **generate** (\$0), **COUNT**(\$1) **as** clicks;

Srtd = **order** Smmd **by** clicks **desc**;

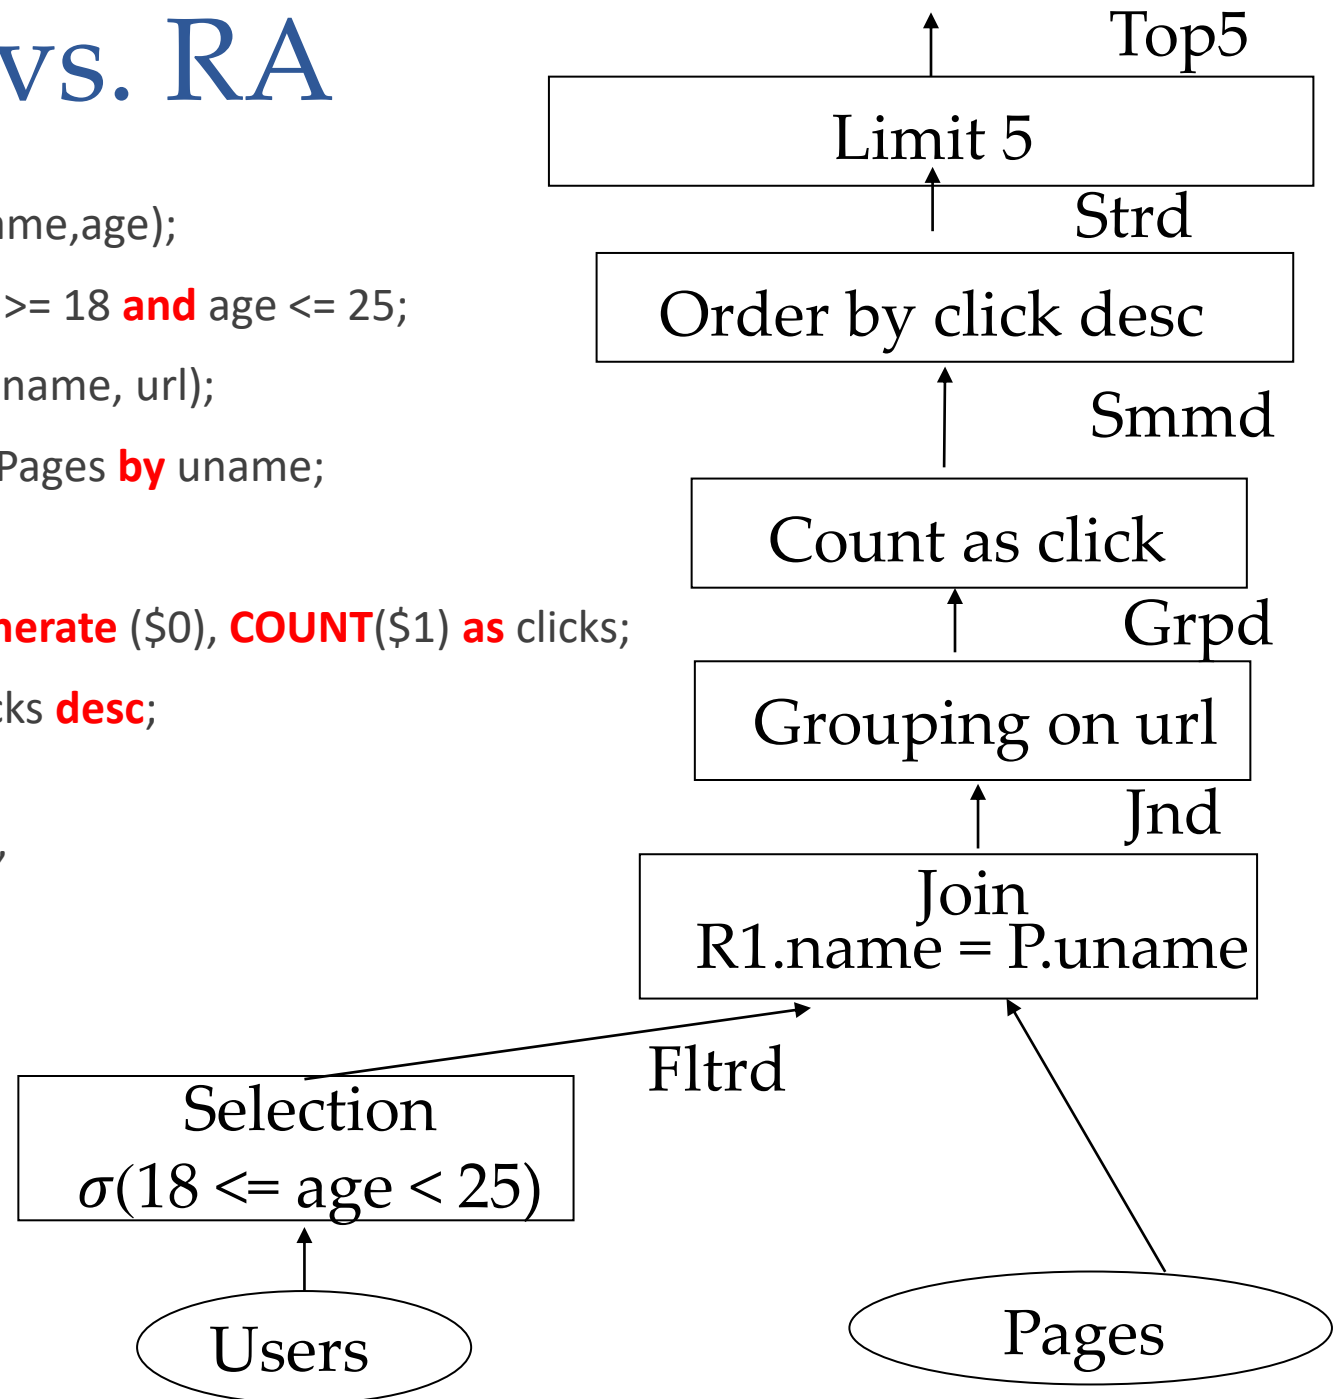
Top5 = **limit** Srtd 5;

store Top5 **into** 'top5sites'

Very similar to Relational Algebra

Pig Latin vs. RA

```
Users = load 'users' as (name,age);  
Fltrd = filter Users by age >= 18 and age <= 25;  
Pages = load 'pages' as (uname, url);  
Jnd = join Fltrd by name, Pages by uname;  
Grpd = group Jnd by url;  
Smmd = foreach Grpd generate ($0), COUNT($1) as clicks;  
Srted = order Smmd by clicks desc;  
Top5 = limit Srted 5;  
store Top5 into 'top5sites'
```



Load and Store

- Users = **load** *'users'* **as** (name,age);
- **Load:** read information from file into a (temp) relation
 - Mostly user defined to translate file format into a relational format
- **Store:** write relation into file
 - Again, usually provided by user

Pig Latin Operators

- Fltrd = **filter** Users **by** age >= 18 **and** age <= 25;
 - Left side: new intermediate relation
 - Right side operation on existing relations
- Operators
 - Selection
 - Res = **filter** R1 **by**:
 - SELECT * FROM R1 WHERE ...
 - By age >= 18; by url matches '*oracle*'
 - Join
 - Res = **join** R1 **by** a1, R2 **by** a2:
 - SELECT * FROM R1, R2 WHERE R1.a1 = R2.a2
 - Order by
 - Res = **order** R1 **by** a1 **desc**
 - SELECT * FROM R1 order by a1 desc

Group By

- Given relation $\text{Rel}(A, B, C)$ with three tuples

(a1, b1, c1)

(a1, b2, c2)

(a3, b3, c3)

- $\text{Grpd} = \text{group Rel by } A;$

- Result relation is $\text{Grpd}(\text{group}, \text{Rel})$

- Attribute 'group' has the same type as attribute A of Rel

- Attribute 'Rel' is a multiset (bag)

- In the given example, Grpd has two tuples, one for each value of A; first attribute of the tuple is the value of A, the second is the set of all tuples of Rel that have this particular value of A

- **dump** Grpd:

41 • (a1, {(a1,b1,c1), (a1,b2,c2)})

• (a3, {(a3, b3, c3)})



For each

- Assume same as before
 - Grpd = **group** Rel **by** A;
 - Result relation is Grpd(group, Rel):
 - (a1, {(a1,b1,c1), (a1,b2,c2)})
 - (a3, {(a3, b3, c3)})
- For each (two options)
 - Smmd = **foreach** Grpd **generate** (\$0), **COUNT**(\$1) **as** c;
 - Smmd = **foreach** Grpd **generate** group, **COUNT**(Rel) **as** c;
- Result relation is Smmd(group, c)
 - Attribute 'group' has the same type as attribute A of Rel ←
 - Attribute c a long
 - **dump** Smmd: *to screen*
 - (a1, 2L) *2 elements*
 - (a3, 1L) *1 element*

Projection and others

- Projection
 - Assume $R1(A, B, C)$
 - $Rel = \text{for each } R1 \text{ generate } A, B;$

Data Model and Flattening

- Supported types: *big data supports*
 - Atomic (string, number...)
 - Tuple (58, 'lilly', 10, 10)
 - Multiset {(58, 'lilly', 10, 10), (33, 'debby', 5, 7)}
 - Further nesting possible: (1, (2,3))
 - Maps (advanced)
- Flattening example
 - Assume $R = \{(1, (2,3))\}$
 - $\text{Res} = \text{foreach } R \text{ generate } \$0, \text{flatten}(\$1)$
make it appear as 1 list instead of nested
 - $\text{Res} = \{(1, 2, 3)\}$
 - Semantics somewhat obscure...
- Sometimes output type not quite clear: try to flatten...

Implementation

- Parser and Query Generator:
 - Everything between load and store translates into one logical plan
- Logical plan:
 - Graph of Hadoop map-reduce jobs
- All statements between two groups → one Map-reduce job

*doesn't execute anything
before dump or store*

Map Reduce Assignment

- Use Pig Latin on top of Hadoop to process (not so) large data set
- We have set up a hadoop cluster with 4 nodes (virtual)
- You have access to that cluster
- You have to write PigLatin Queries
- You have to observe the execution
- Instructions will be on myCourses