


Transactions

Business transactions

- Business transaction: involves exchange between two or more entities (selling, buying, renting, booking).
- In computers: business transactions become **electronic transactions** (same ideas)



Book-keeping

Money Transfer

In Application Program

Transfer(id1,id2,value) {

/* retrieve first balance through SQL */

ResultSet rs = stmt.executeQuery(

"SELECT balance FROM accounts WHERE accountid =" + id1);

while (rs.next()) { bal = rs.getInt("balance");}

/* calculate new value */

bal += value

/* update balance through SQL */

stmt.executeUpdate(

"UPDATE accounts SET balance = " + bal + " WHERE accountid = " + id1)

/* do similar for second balance */

ResultSet rs = stmt.executeQuery(

"SELECT balance FROM accounts WHERE accountid =" + id2);

while (rs.next()) { bal = rs.getInt("balance");}

bal -= value

stmt.executeUpdate(

"UPDATE accounts SET balance = " + bal + " WHERE accountid = " + id1)}

-
1. read(*A*)
 2. $A := A + 50$
 3. write(*A*)
 4. read(*B*)
 5. $B := B - 50$
 6. write(*B*)

Electronic transaction

- An *electronic transaction* encapsulates operations that belong logically together
- *Boundaries* of the transaction have to be defined by the programmer (according to application semantics)
- Contains
 - DBS operations
 - Application code
- Examples:
 - Purchasing elements in shopping cart of online e-commerce side
 - Flight Reservation (might book more than one flight)

*define
logical
units
of
work*

Application vs. DBS

- A user's program
 - interaction with the database (select, update, delete...)
 - own computation
- A database **transaction** is the DBMS's abstract view of a user program:
 - a sequence of **read operations $r(X)$** and **write operations $w(X)$** on **objects (X)** (tuple, relation,...) of the DB
 - **Read:** bring object into main memory from disk, send value to application (same as copy value into program variable)
 - **Write:** bring object into main memory from disk and modify it. Write it back to disk (might be done sometime later)

Online transaction processing (OLTP)

- Database Schemas with 20-30 tables
- Pre-defined application tasks of reasonable size
 - 4-20 SQL operations cover one task = transaction
 - Individual Queries relatively simple (1-4 tables involved).
 - Often point queries (all information from one client /student)
 - Efficient execution possible
 - Update intensive: 15 to 50% of SQL statements are updates
- Examples
 - Banking Teller machines
 - Online banking (more queries than teller machine....)
 - Travel Reservations
 - E-Commerce
 - Stock Exchange
 - Cash Register in Supermarket
- Volumes: up to hundreds/thousands of transaction per minute
- Other application domains
 - More queries (and complex ones) – less updates *analytical processing systems*
 - Much simpler tasks – only one read or one write operation

ACID Properties

- The success of transactions was due to well-defined properties that are provided by the database system:
 - **ATOMICITY:**
 - A transaction is atomic if it is executed in its entirety or not at all *all or nothing*
 - **CONSISTENCY:**
 - a transaction must preserve the consistency of the data *integrity constraints + application programmer*
 - **ISOLATION:**
 - in case that transactions are executed concurrently: The effect must be the same as if each transaction were the only one in the system.
 - **DURABILITY:**
 - The changes made by a transaction must be permanent (= they must not be lost in case of failures)
- The application programmer only has to indicate
 - When a transaction starts ←
 - The sequences of SQL operations ←
 - When the transaction finishes (abort or commit request possible) ←
 - AID are guaranteed by the DBS

Money Transfer

In Application Program

```
Transfer(id1,id2,value) {
```

```
    /* begin txn */
```

```
    /* often automatic */
```

```
    ...
```

```
    /* increase balance 1st account*/
```

```
    ResultSet rs = stmt.executeQuery(
```

```
        "SELECT balance FROM accounts WHERE accountid =" + id1);
```

```
    ...
```

```
    /* decrease balance 2nd account*/
```

```
    ResultSet rs = stmt.executeQuery(
```

```
        "SELECT balance FROM accounts WHERE accountid =" + id2);
```

```
    while ( rs.next() ) { bal = rs.getInt("balance");}
```

```
    bal -= value
```

```
    /* abort if necessary*/
```

```
    if bal < 0
```

```
        connection.rollback();
```

```
    ELSE {
```

```
        stmt.executeUpdate(
```

```
            "UPDATE accounts SET balance = " + bal + " WHERE accountid = " + id1)
```

```
        connection.commit();}
```

```
}
```

first update

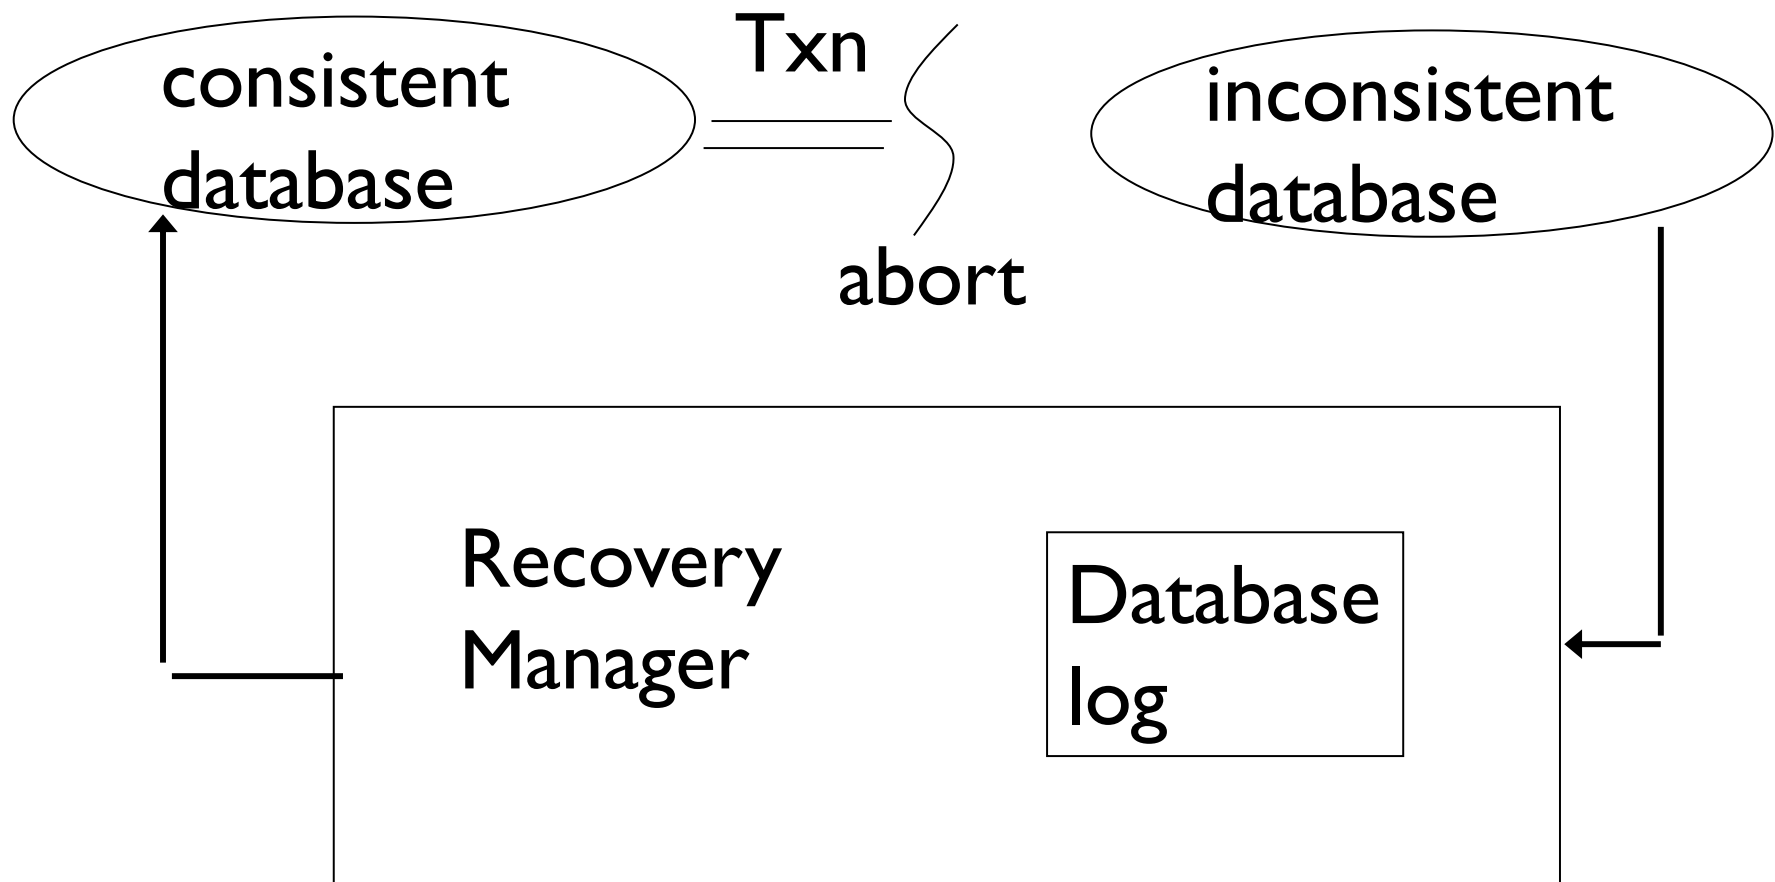
rollback the transaction -undo the update

connection object from jdbc

Atomicity: All or Nothing

- A transaction T might **commit** after completing all its actions.
 - If the user is informed about the commit they can be sure that all changes performed by T are installed in the DB.
- A transaction might **abort** (or be aborted by the DBMS) after executing some actions.
 - In this case the DBMS undoes all modifications so far.
 - After the abort the DB state is as if the transaction had never started.
 - After notification of abort the user knows that none of the transaction's modifications is reflected in the database.
- **Local recovery:** eliminating partial results in case of abort
 - before executing write(x)
 - store “before image” of x (somewhere in main memory → called the log)
 - » log record: txn-id, rid, before-image
- Examples
 - committing transaction: read(x), write(x), write(y), read(z), commit
 - aborting transaction
 - read(x), write(x), write(y) ... write⁻¹(y), write⁻¹(x) abort
 - write⁻¹(y) installs before-image of y taken from log

Atomicity




Durability

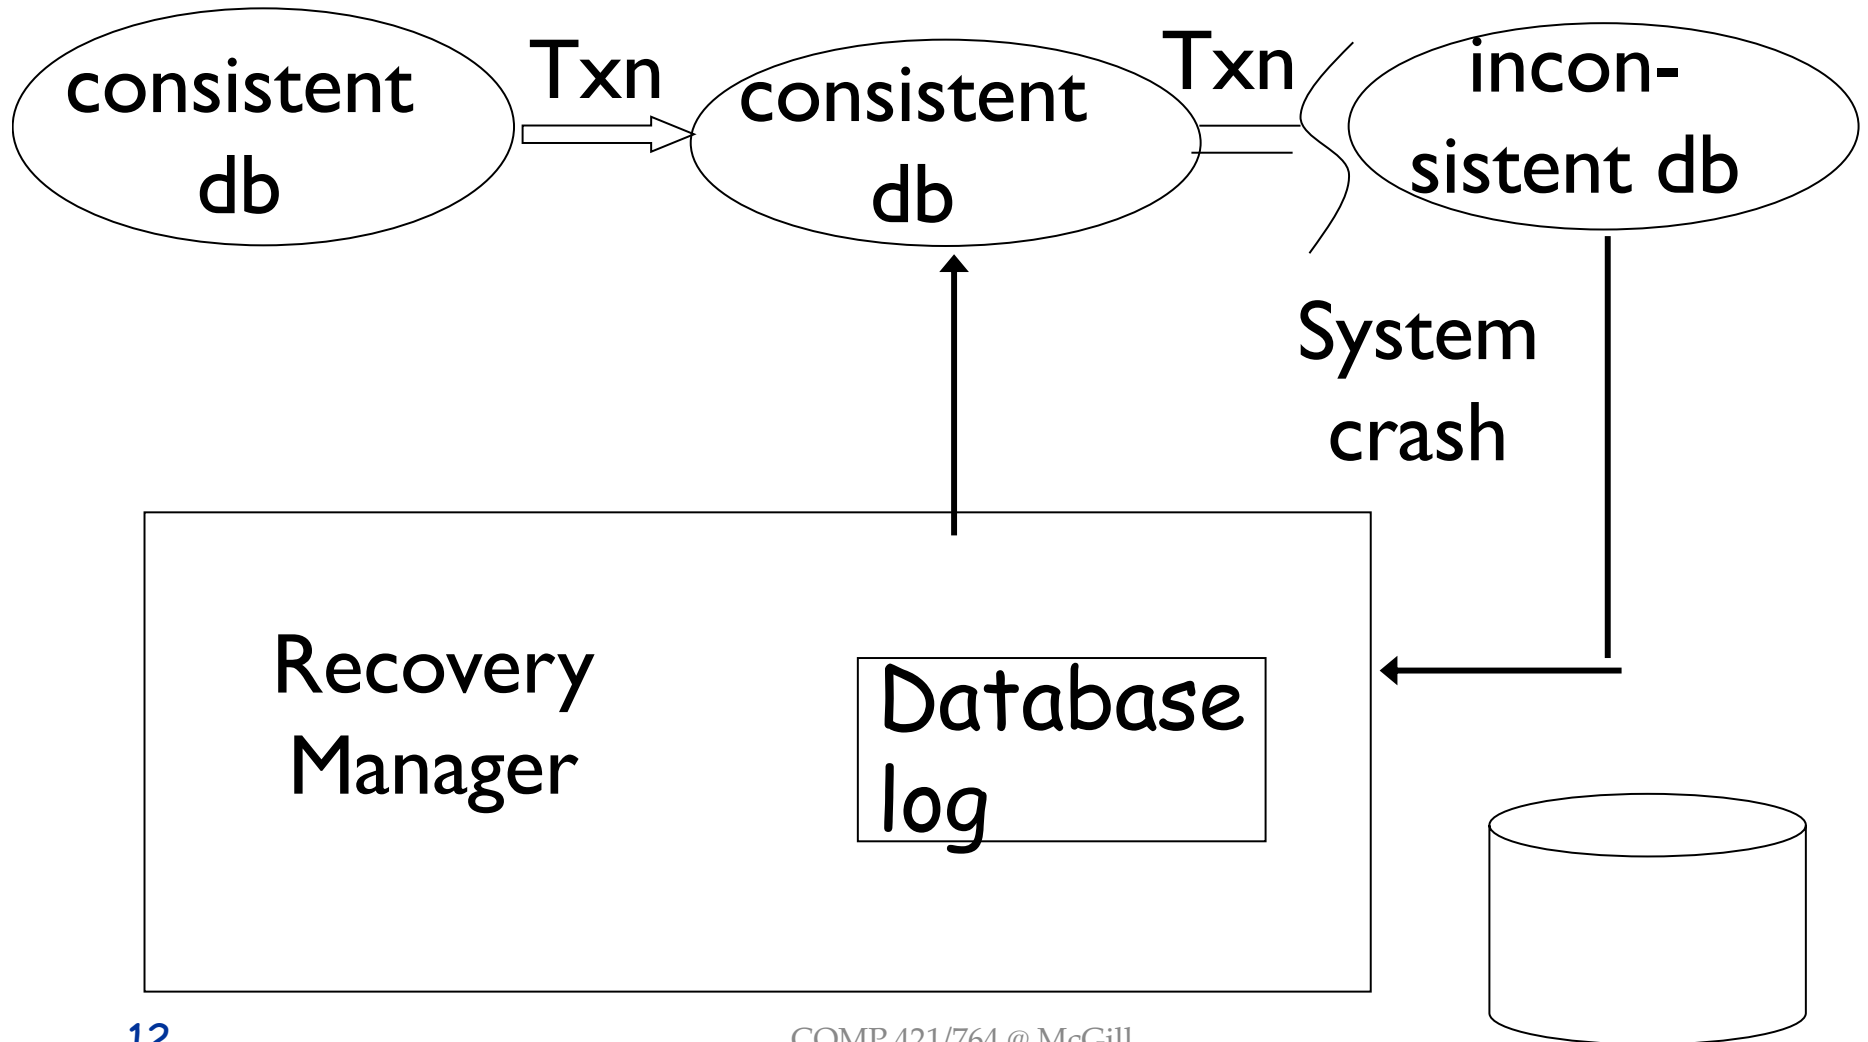
- **Durability**: There must be a guarantee that the changes introduced by a transaction will last, i.e., survive failures
 - once user has commit confirmation the change must persist in database despite possible crash

❖ Desired Behavior after system restarts:

- T1, T2 & T3 should be durable.
- T4 & T5 should be aborted (effects not seen).



Global Recovery



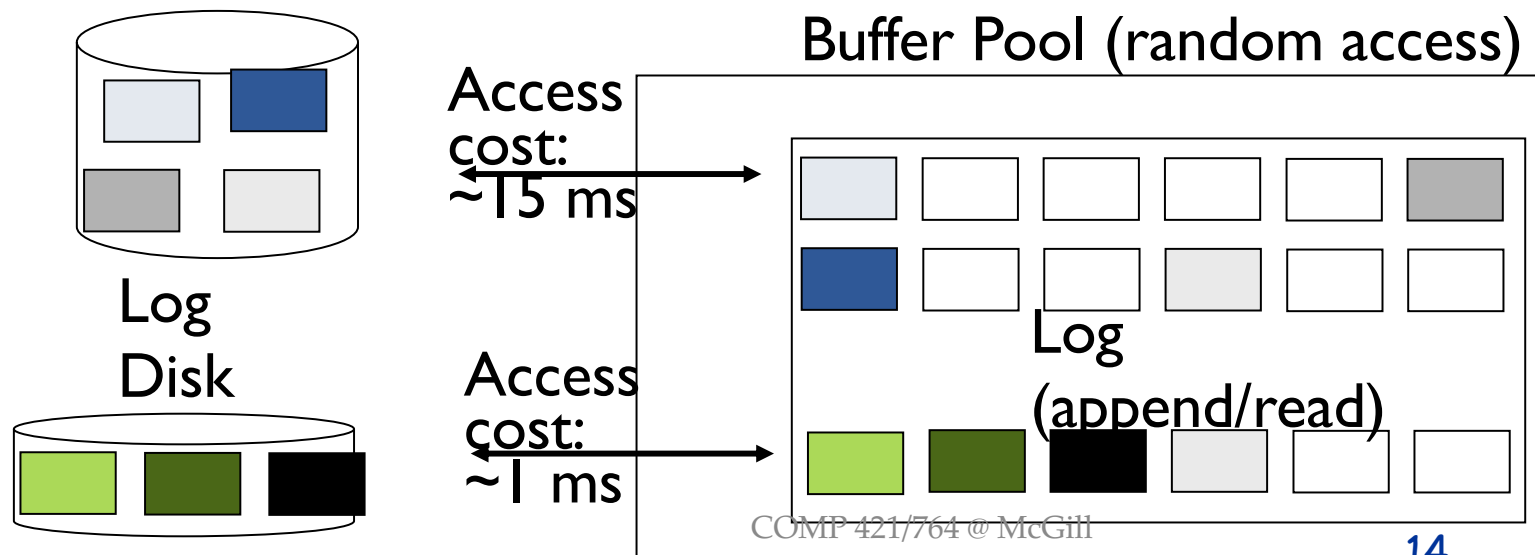
Recovery after System Crash

- restart server and perform **global recovery:**
 - transactions that committed before crash
 - make sure all changes are in database
 - transactions that aborted before crash
 - make sure no changes are in database
 - transactions that are active at time of crash
 - make sure no changes are in database
- Assumption:
 - disk does not crash ← *everything on disk is persistent and safe*

enough info needs to be on disk to recover

Updates vs. logs

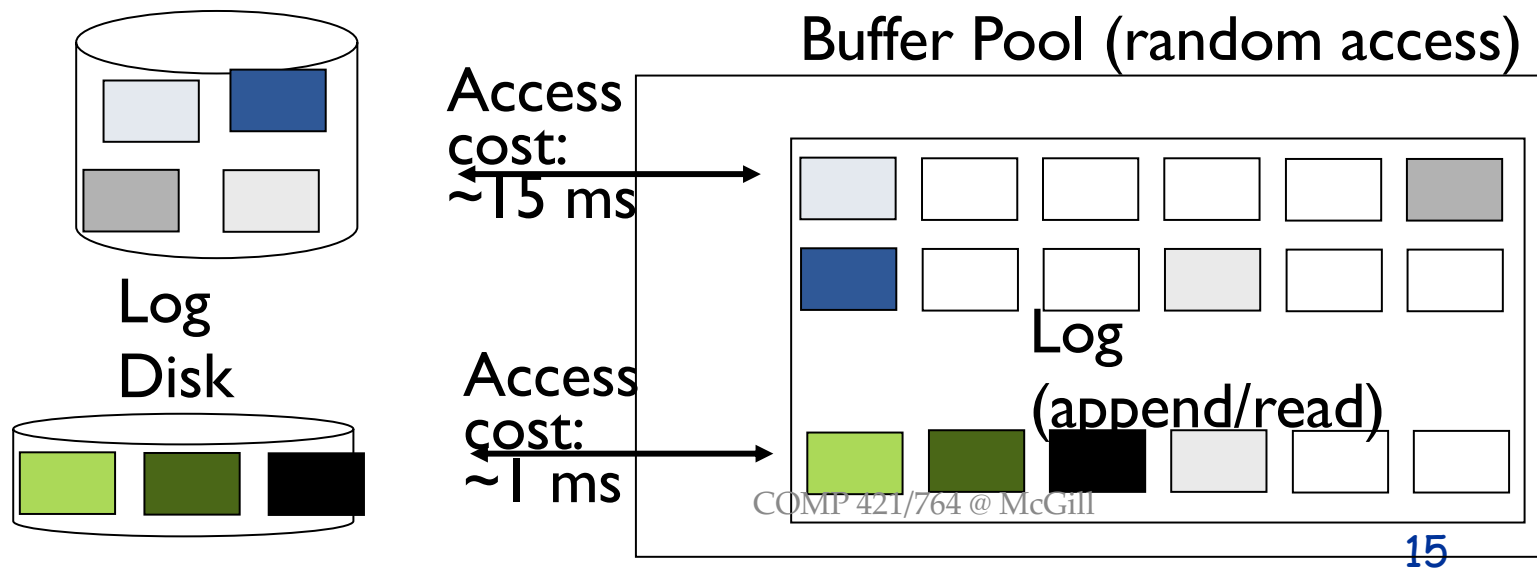
- Each update changes the record on the corresponding page
- Dirty (changed) pages are written to disk whenever the buffer manager decides (replacement policy)
- To be able to undo/redo in case of crash
 - Write log information to a log disk
 - Faster writes than random writes to standard disk



High performance system
↓
log is on a different disk

WAL: write ahead logging

- for each write(x) of transaction T; x residing on DB page P:
 - Append log record with before- and after-image of x to log tail
 - before-image: you can undo changes of active transactions
 - after-image: you can redo changes of committed transactions
- Flush log to disk before flushing the P: then you can undo when you have to abort *all necessary before images are on disk*
- Flush log to disk before commit of update transaction T: then you can redo when something committed (read-only transactions don't create dirty pages and log records – so ignore)
- Log is typically an append; flush means all of the log is flushed in one shot; cheaper than standard flush



UNDO/REDO Recovery

- Recovery:

- Undo: For each transaction T that did not commit

- Find all of T's log records and install before images

→ all records with
aborts or neither
commit nor abort (CNF)

- Redo: For each transaction T that did commit

- Find all of T's log records and install after image

→ all records with
commit + records

- If there were many committed transactions updating x?

- Challenges:

- Doing things in the right order

- How do I know which transactions committed

- Write commit/abort logs, too

- How do I avoid redoing all committed transactions since system start

- Checkpointing

- Write out dirty pages on a regular basis during normal processing

16 • Helps reducing redo phase