

Isolation and Concurrency Control

Isolation

- DBMS can execute transactions concurrently
 - exploitation of resources
 - one transactions performs I/O, the other uses CPU etc.
 - Exploiting multi-core
- Isolation:
 - although transactions execute concurrently, each transaction runs in isolation -- not affected by the actions of other transactions
- Isolation is enforced by a **concurrency control** protocol,
- Concurrency control provides serializable executions
 - Net effect of transactions executing concurrently is identical to executing all transactions one after the other in some serial order

Isolation

Txn 1

Txn 2

consistent
database

inconsistent
database

Txn 1

Txn 2

consistent
database

consistent
database

Concurrency
Control

Serial Execution: Example

- Consider two transactions (X_{acts}/Txn):

T1:
 $A = A + 100$,
 $B = B - 100$
commit

T2:
 $A = 1.06 * A$,
 $B = 1.06 * B$
commit

- T1 transfers \$100 from B's account to A's account.
- T2 credits both accounts with a 6% interest payment.
- Assume $A=B=200$ at beginning
- Serial execution I: T1 executes before T2: Values of A and B?
 - $A = 318, B = 106$ (Sum = 424) ←
- Serial execution II: T2 executes before T1: Values of A and B?
 - $A = 312, B = 112$ (Sum = 424) ←
- Both execution orders make sense (sums are the same)

Concurrent Execution: Example

- Consider two transactions (*Xacts/Txn*):

T1:
 $A = A + 100$,
 $B = B - 100$
commit

T2:
 $A = 1.06 * A$,
 $B = 1.06 * B$
commit

Same as:

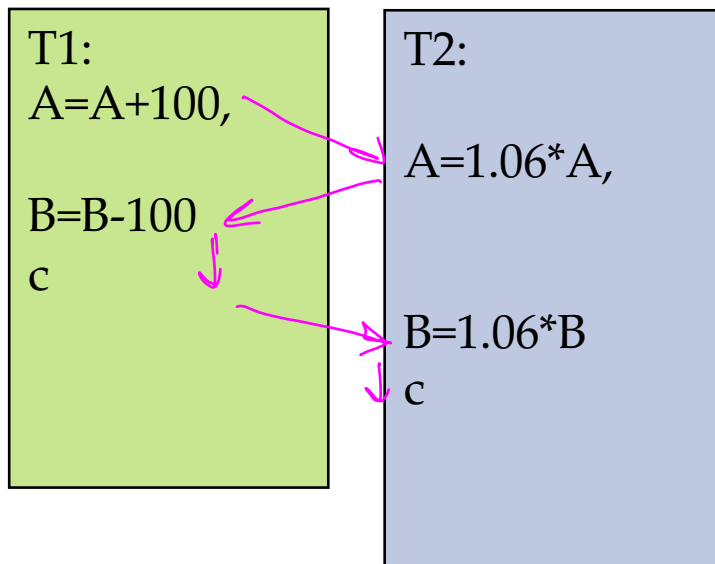
T1:
 $R1(A)$
 $W1(A)$
 $R1(B)$
 $W1(B)$
c1

T2:
 $R2(A)$
 $W2(A)$
 $R2(B)$
 $W2(B)$
c2

- T1 transfers \$100 from B's account to A's account.
- T2 credits both accounts with a 6% interest payment.
- Now assume T1 and T2 are submitted at the same time
 - No guarantee that T1 will execute before T2 or vice-versa
 - The net effect *must* be equivalent to these two transactions running serially in some order.

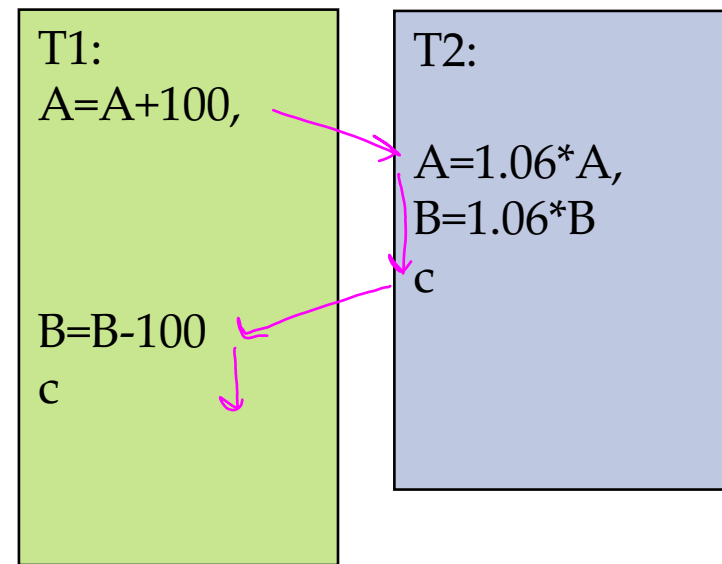
Example (Contd.)

□ Consider two interleavings (**schedules**):



□ good ✓

Same as
executing
one after
the other



□ A bad one: ✗

☆ The 100\$ that are transferred are included twice in the interest rate calculation

318 + 212 bad

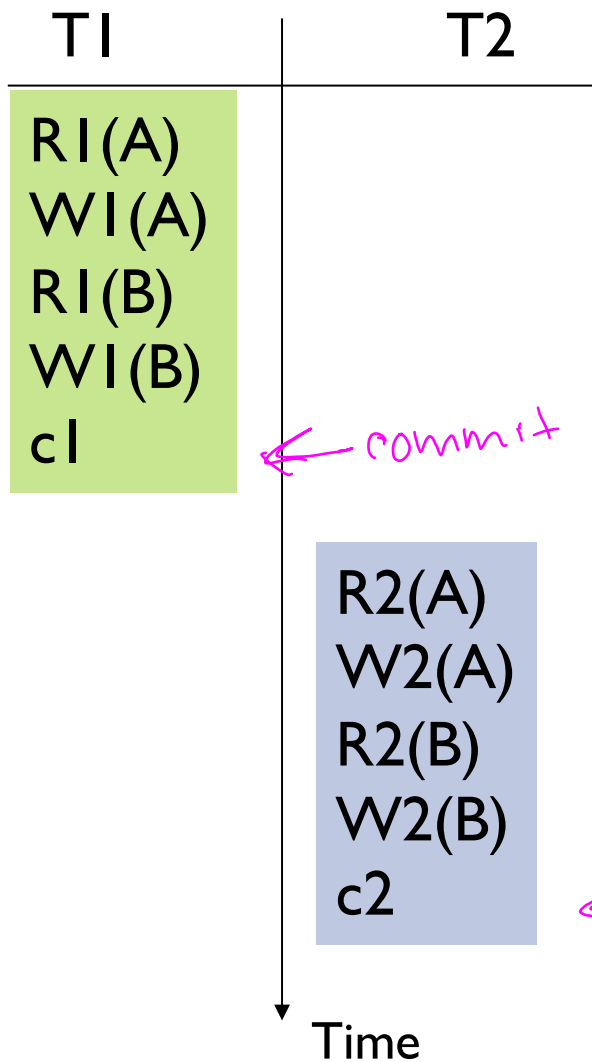
inconsistent state

Schedules


- Transaction:
 - A sequence of read and write operations on objects of the DB (denoted as $r(x)/w(x)$)
 - Each transaction must specify as its final action either commit (c), i.e. complete successfully or abort (a), i.e., terminate and undo all the actions carried out so far.
- Schedule:
 - sequence of actions (read,write,commit,abort) from a set of transactions
 - Reflects how the DBMS sees the execution of operations; ignores things like reading/writing from OS files etc.
- Complete Schedule:
 - Contains commit/abort for each of its transactions.
- Serial schedule:
 - Schedule where transactions are executed one after the other.

Examples

Serial Schedule




Non serial Schedule



Reading Uncommitted Data: Dirty Reads

- ❑ T1: money transfer (as before)
- ❑ T3: sum of all accounts

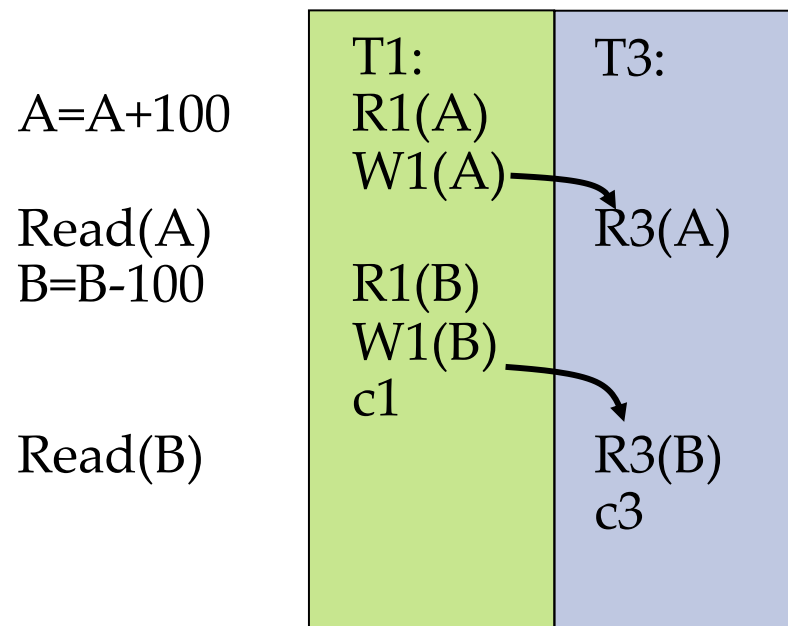


- ❑ The user perspective: *appears...*
 - ☆ T3 executes after T1 because it reads the A value T1 has written;
 - ☆ T3 executes before T1 because it read the B value before T1 has written

- ❑ If T3 reads from T1 before T1 commits, it might read inconsistent data (**Inconsistent or Dirty Reads**)

Reading Uncommitted Data: Dirty Reads

- ❑ T1: money transfer (as before)
- ❑ T3: sum of all accounts




- ❑ The user perspective:
 - ☆ T3 always reads after T1 writes;
 - ☆ It is as if T3 executes serially after T1

- ❑ If T3 reads from T1 before T1 commits, it **might** read inconsistent data (**Inconsistent or Dirty Reads**)
- ❑ But it might also be ok

Unrepeatable Reads

- ❑ T4 reads A twice
- ❑ T5 updates A



- ❑ The user perspective:

- ☆ T4 executes before T5
because it reads A before T5 writes it

- ☆ T4 executes after T5 because
it reads A after T5 writes it

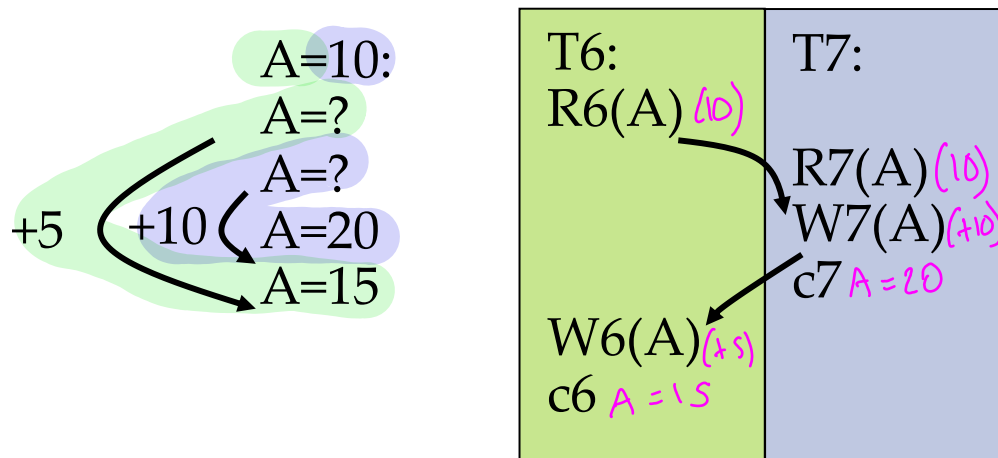
conceptually executes before AND after

- ❑ If T4 reads twice the same data item, but T5 changes the value between the first and second read, then we have **unrepeatable read** situation.

Serious

Lost Update

- T6: $A = A + 5$
- T7: $A = A + 10$
- $A = 10$ at start




- The user perspective:
 - It is as if T7's update has never taken place; it is not reflected in the database
 - If it were reflected the final value of A should be 25 and not 15

❑ Can lead to **lost update**

T7's update was lost

Committed and Aborted Transactions

- If a transaction aborts, all its actions are undone.
- It is as if they were never carried out



- The user perspective:

- • T8 reads a value for A that actually will never exist!
- • Problem even bigger if the read triggered a further change in the database
- • Problem even bigger if A is updated in between; how to undo in this case???

- **Dirty Read** can lead to reading a non-existing value
- **Dirty Write** can mess up the database
 - AVOID AT ALL COSTS!!


Conflicting Operations

- **Conflicting operations:** Two operations conflict if *(from 2 different transactions)*
 - They access the same object
 - Both operations are writes, or one is write and one is read
- Schedules
 - serial schedule: $T1\ T2 = r1(x)\ r1(y)\ w1(y)\ c1\ r2(x)\ w2(x)\ r2(y)\ c2$
 - all operations of T1 are executed before any operation of T2
 - in particular: if T1 and T2 have several conflicting operations, T1's operation is always executed before T2's conflicting operation.
 - Our examples:
 - one operation of T1 was ordered before the conflicting operation of T2
 - another operation of T1 was ordered after T2's conflicting operation

Conflict Serializable Schedules

- □ Two schedules are **conflict equivalent** if:
 - ☆ Involve the same actions of the same (committed) transactions
 - ☆ Every pair of conflicting actions of (committed transactions) is ordered the same way
- □ Schedule S is **conflict serializable** if
 - ☆ S is conflict equivalent to some serial schedule which contains the committed transactions of S
 - ☆ Textbook differentiates between
 - Serializable
 - Conflict-serializable
 - View-serializable
 - ☆ Here:
 - **conflict-serializable = serializable**
 - Ignore view-serializable

Serial Examples



Serial Schedules

Schedules written in a different format:




SL1: r1(x) w1(x) w1(y) c1 r2(z) r2(y) w2(x) c2

SL2: r2(z) r2(y) w2(x) c2 r1(x) w1(x) w1(y) c1

Examples

interleaved

conflict
serializable



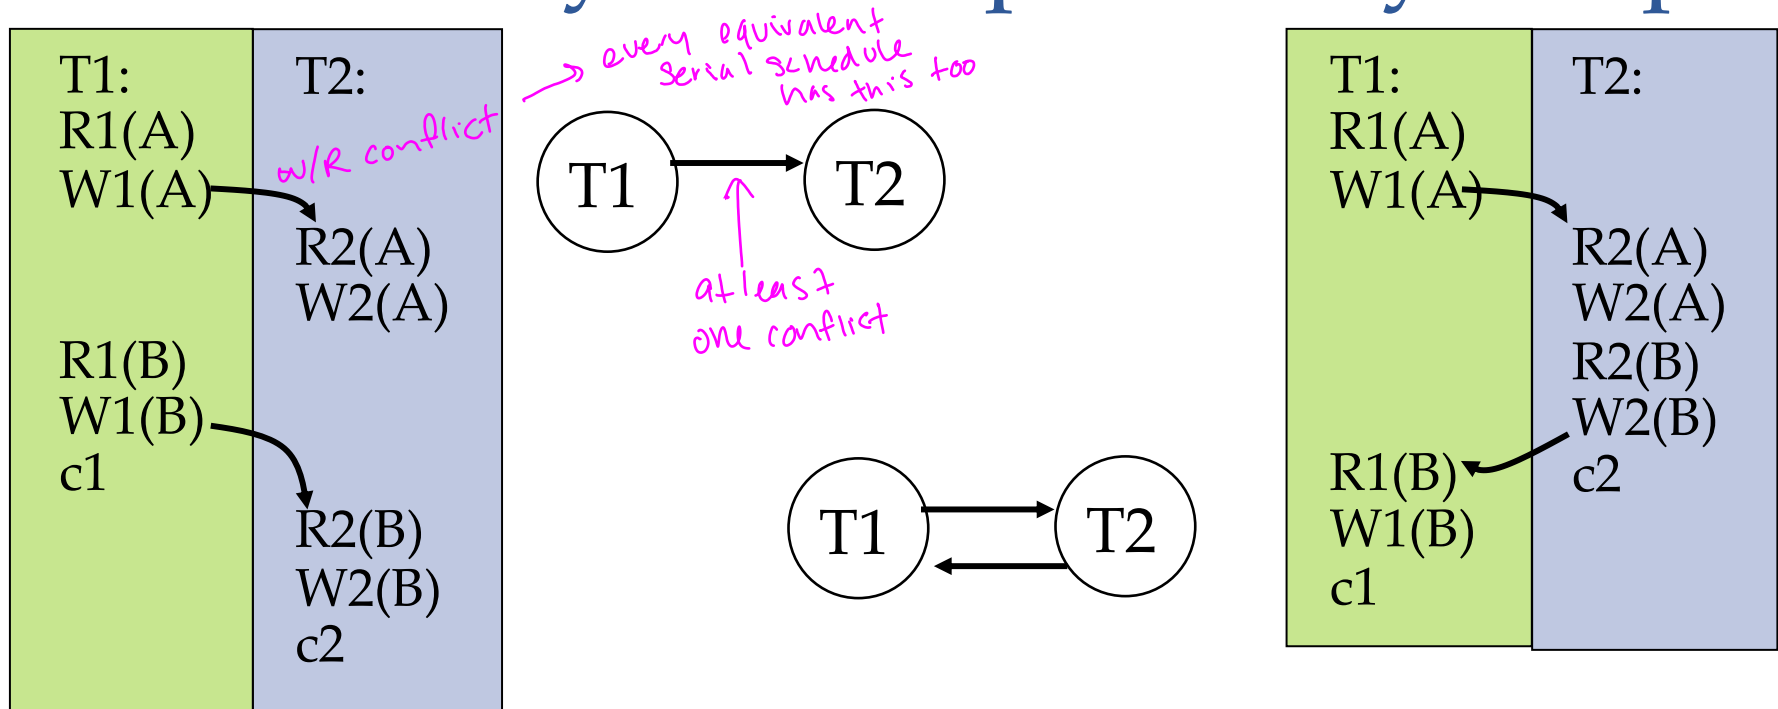
Schedules written in a different format:

S1: r1(x) w1(x) r2(z) r2(y) w2(x) c2 w1(y) c1

S2: r1(x) r2(z) r2(y) w1(x) w1(y) c1 w2(x) c2

S3: r1(x) r2(z) w1(x) w1(y) c1 r2(y) w2(x) c2

Serializability and Dependency Graphs



- Dependency graph / Serialization graph / precedence graph / Serializability graph for a schedule:

- Let S be a schedule $(T, O, <)$

- Each transaction T_i in T is represented by a node
- There is an edge from T_i to T_j if an operations of T_i precedes and conflicts with one of T_j 's operations in the schedule.

Dependency Graphs

□ Theorem: Schedule is conflict serializable if and only if its dependency graph is acyclic


□ Generating an equivalent serial schedule (flatten)

Continue until no nodes are left

Choose a source (i.e. a node without incoming edges)

put the corresponding transaction next in the serial order

Delete the node and all outgoing edges



We want an equivalent serial schedule

Concurrency Control

- Given an execution (schedule) we can test whether the execution was serializable
 - If execution was serializable, then ok ←
 - If not serializable, then it's too late! ←
- Concurrency control:
 - during execution take measures such that a non-serializable execution can never happen prevent

graphs are just to understand

we need to prevent non-serializable executions

Concurrency Control: Locking


- No conflict: transactions can execute at the same time (e.g. all reads)
- Upon first conflict: the second transaction has to wait until the first transaction commits/aborts
- Locks: Two types, because two read operations do not conflict
- Basics of locking:
 - Each transaction T_i must obtain a S (shared) lock on object before reading, and an X (exclusive) lock on object before writing.
 - If an X lock is granted on object O, no other lock (X or S) might be granted on O at the same time. *If one is writing, everyone must wait*
 - If an S lock is granted on object O, no X lock might be granted on O at the same time. *If one is reading, writes must wait*
 - If a conflicting lock is active, the transaction must wait until the lock is released
 - Conflicting locks are expressed by the compatibility matrix:

	S	X
S	✓	--
X	--	--

at end of transaction

Strict Two Phase Locking

- Phase 1 = growing phase: acquiring locks whenever you need one
 - Each transaction T_i must request a S (*shared*) lock on object *before* reading, and an X (*exclusive*) lock on object *before* writing.
 - If no conflicting lock is active is set, the lock can be granted (and the transaction can execute the operation),
 - If a conflicting lock is active the transaction must wait until the lock is released
- Phase 2 = shrinking phase: After a transaction has released one of its lock (unlock) it may not request any further locks
- Strict: a transactions releases all its lock at the end of its execution after commit
 - Shrinking phase happens in one shot at end of transaction.



2PL allows only serializable schedules
 Strictness: No dirty reads / no dirty writes


Strict 2PL: another example

T1: R(A), W(B)

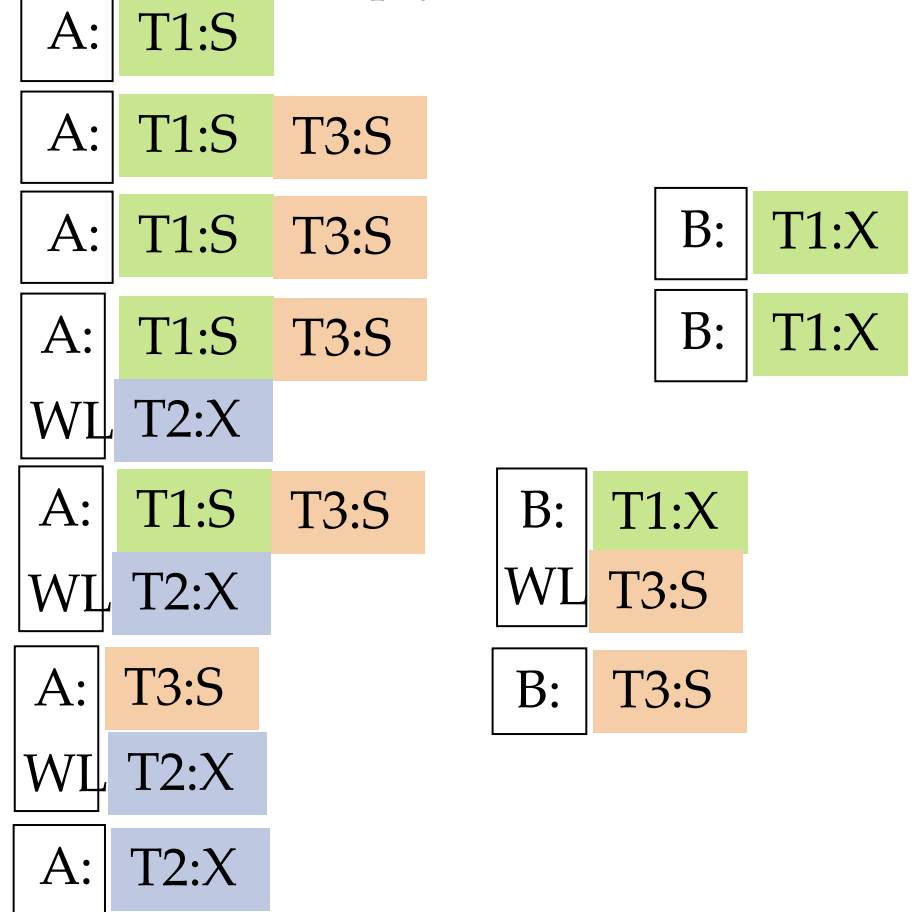
T2: W(A)

T3: R(A), R(B)

Order of submission: R1(A), R3(A), W1(B), W2(A), R3(B)



Lock Table: empty at start



Lock Table: empty at end

Implementing strict 2PL

- Lock request
 - If lock is S, no X lock is active and the request queue is empty:
 - add the lock to the granted lock queue and set lock type to S
 - If lock is X and no lock active (\Rightarrow the request queue is also empty):
 - add the lock to the granted lock queue and set lock type to X
 - Otherwise:
 - add the lock to the request lock queue
 - In the first two cases, the transaction can continue immediately. In the last case the transaction is blocked until the lock is granted
- Lock release (at end of transaction)
 - Remove the lock from the granted lock queue
 - If this was the only lock granted on the object:
 - grant one write lock (if the first lock in the request queue is a write) or
 - n read locks (if the first n locks in the request queue are reads) as described above.

Details

- A transaction does not request the same lock twice.
- A transaction does not need to request a S lock on an object for which it already holds an X lock.
- If a transaction has an S lock and needs an X lock it must wait until all other S locks (except its own) are released

Implementation Details

- Locks are managed using a **lock table**
- The lock table has a lock table entry for each object that is currently locked
 - Pointer to queue of granted locks (or simply the number of transactions currently holding a lock)
 - Type of lock held (shared or exclusive)
 - Pointer to queue of lock requests (waiting transactions)
 - A transaction T contains only one lock per object
 - if a T has an S lock and requests an X lock, the S lock is upgraded to an X lock (may have to wait/block for other S locks to be released)
- Locking and unlocking have to be atomic operations
 - Set latch/semaphore when accessing lock table
- Transaction table:
 - For each transaction T contains pointer to a list of locks held by T

Why does 2PL work?

- When is a schedule not serializable?
 - If there are operations of transactions T1 and T2 such that T1 should be ordered before T2 AND after T2 in the schedule T1 -> T2 -> T1


Submission order 1

R1(A) W2(A) W2(B) W1(B)


Submission order 2

R1(A) W2(B) W1(B) W2(A)

R1/W2 + R2/W1
 R1/W2 + W2/W1
 R1/W2 + W2/R1
 W1/W2 + R2/W1
 W1/W2 + W2/W1
 ...
 W1/R2 + R2/W1
 ...



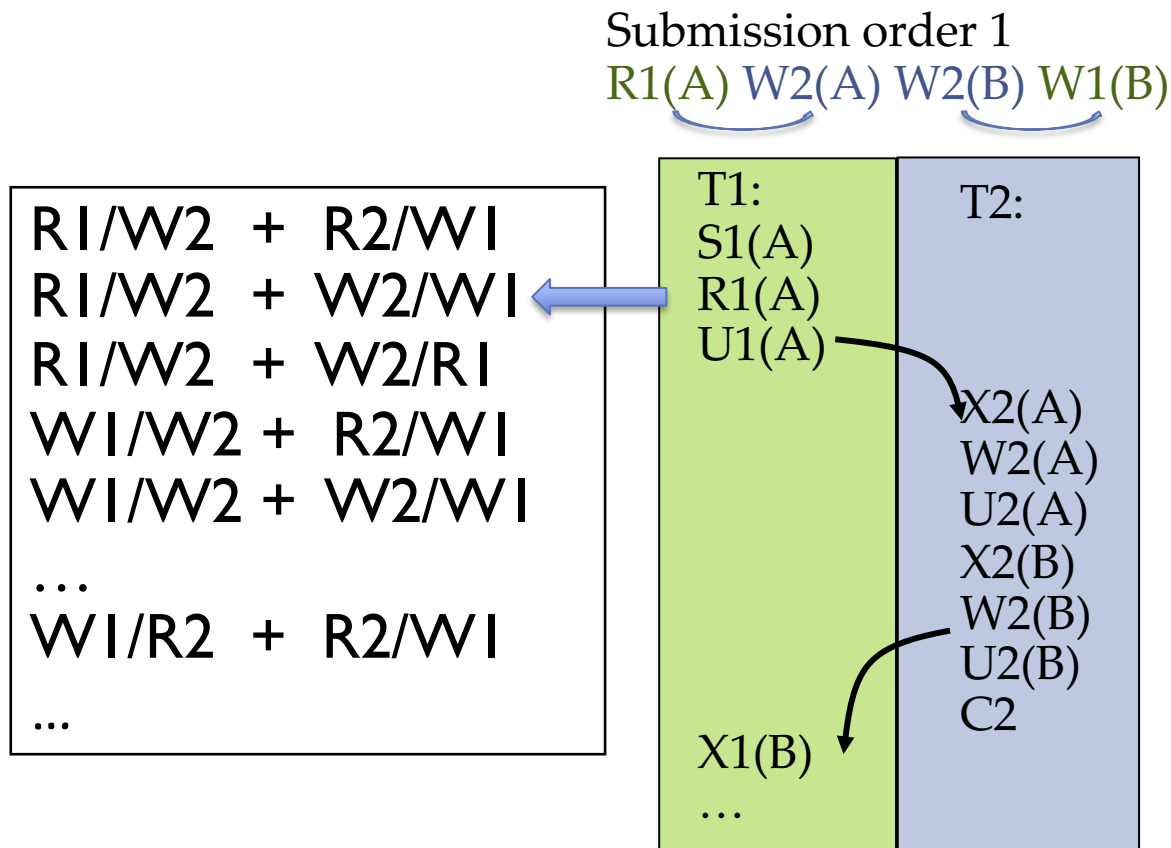
transactions ordered according to who got the first lock



Deadlock:
see later

Why does simple locking NOT work?

- It allows any order!!





Same as
setting no
locks
at all



Operations are ordered as they
are submitted

Deadlocks

- Deadlock: Cycle of transactions waiting for locks to be released by each other.
- Waits-for graph:
 - Nodes are transactions
 - There is an edge from T_i to T_j if T_i is waiting for T_j to release a lock
- Deadlock detection: look for cycles in the wait-for graph

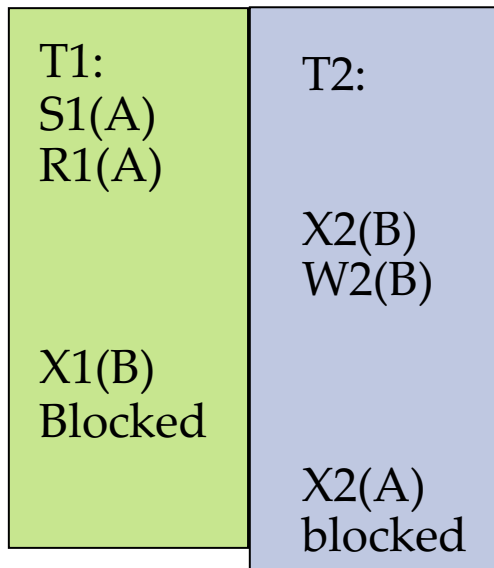


Deadlock Detection (Continued)




Dependency graph - wait-for-graph

- Note: is similar to dependency graph with the following difference
 - If there is an edge from T2 to T1 in the wait-for-graph, then T2's operation will execute after T1's operation (T2 waits for T1 to release the lock), hence, in the dependency graph there is an edge from T1 to T2
 - Deadlocks can happen because 2PL avoids unserializable schedules by locking objects!



Wait-for-graph



Depend. graph

