


# Graph Databases

# The “Others”


Complete trend, starting with January 2013



Source: [https://db-engines.com/en/ranking\\_categories](https://db-engines.com/en/ranking_categories)

# The “Others”

**Ranking scores per category in percent, April 2023**



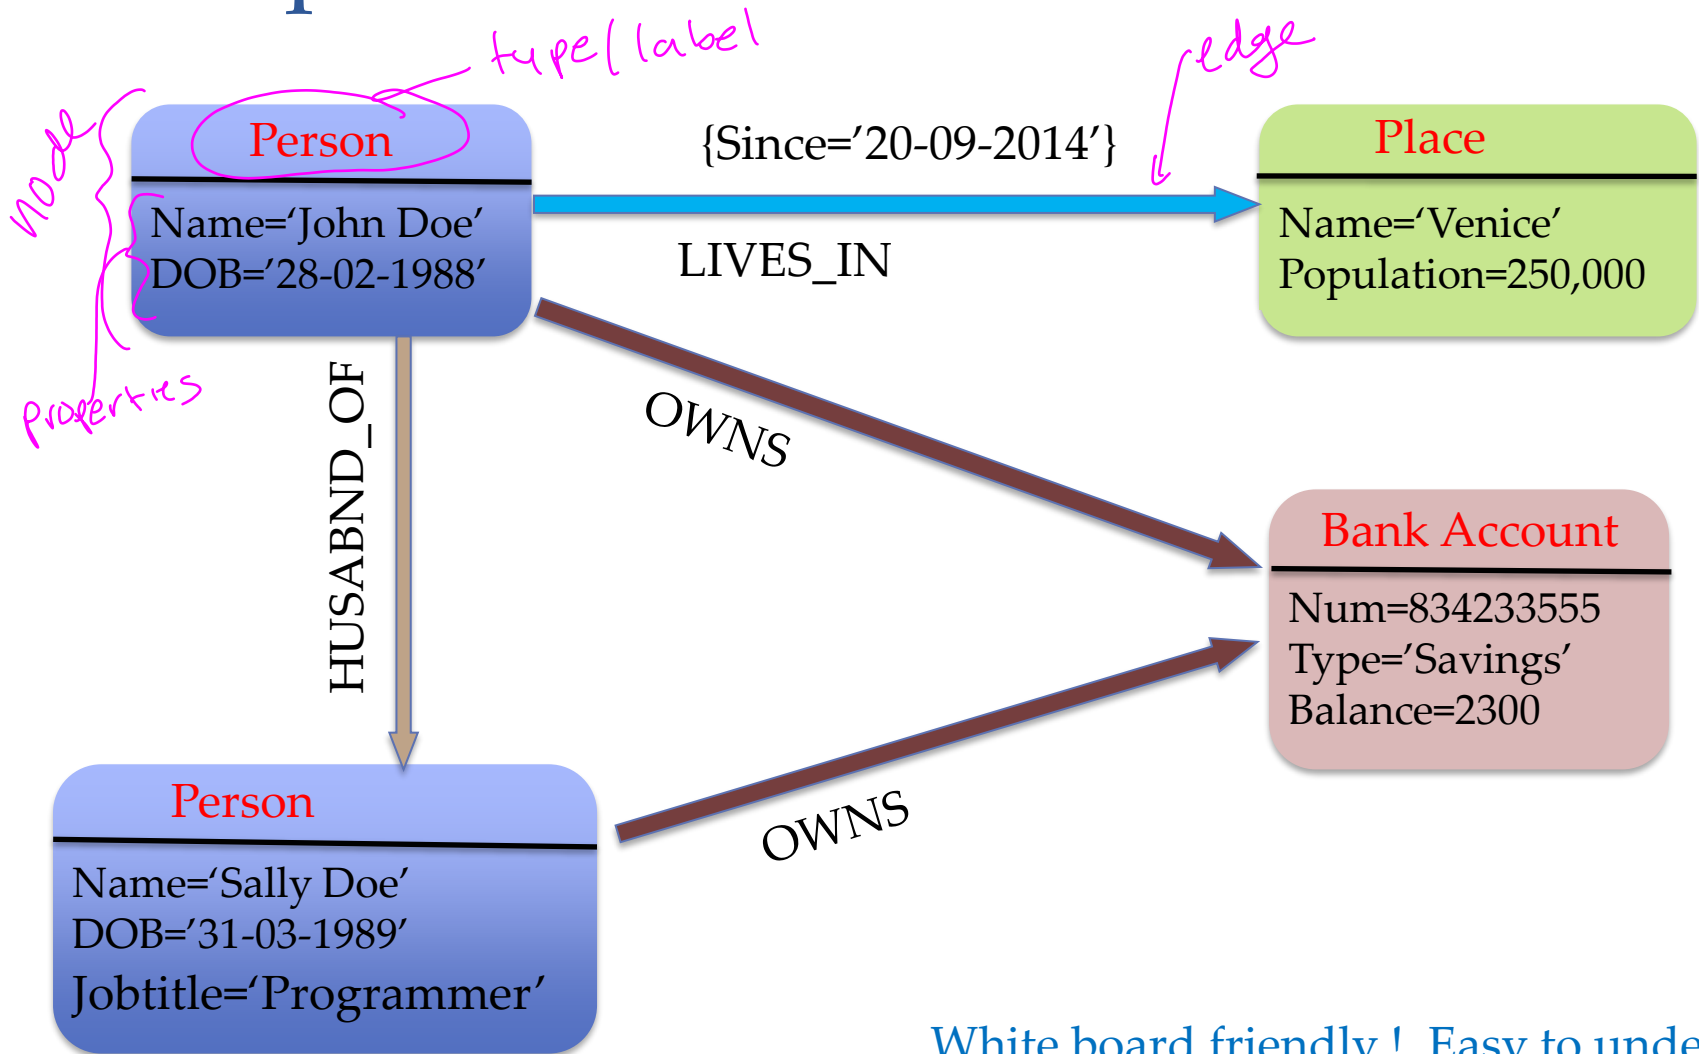
© 2023, DB-Engines.com

Source: [https://db-engines.com/en/ranking\\_categories](https://db-engines.com/en/ranking_categories)

# Introduction

- Based on Euler's graph theory
- Data Model
  - **Nodes/Vertices** of the graph → Represents real-world entities (Eg. a **Person** (**John Doe**), a **Place** (**Venice**), a **Bank account** (**834233555**), etc.)
    - Nodes may be associated with a **Label/Type** (Eg. Person, Place, etc.)
  - **Edges** between nodes in the graph → Represents relationships **between two entities** (Eg. **LivesIn**, **Owns**, etc. )
    - Neo4J:- Relationships are **directional** in nature.
  - **Properties** are **key-value** pairs that are associated with either a particular node or a particular relationship. (Eg. A person can have the following properties { **name**:**John Doe** , **dob**:**29-02-1988** }
  - Each node/edge are free to have its own set of (possibly different) set of properties for example some nodes representing a person can have the property **jobtitle** whereas others may not. → concept of a **sparse schema**.

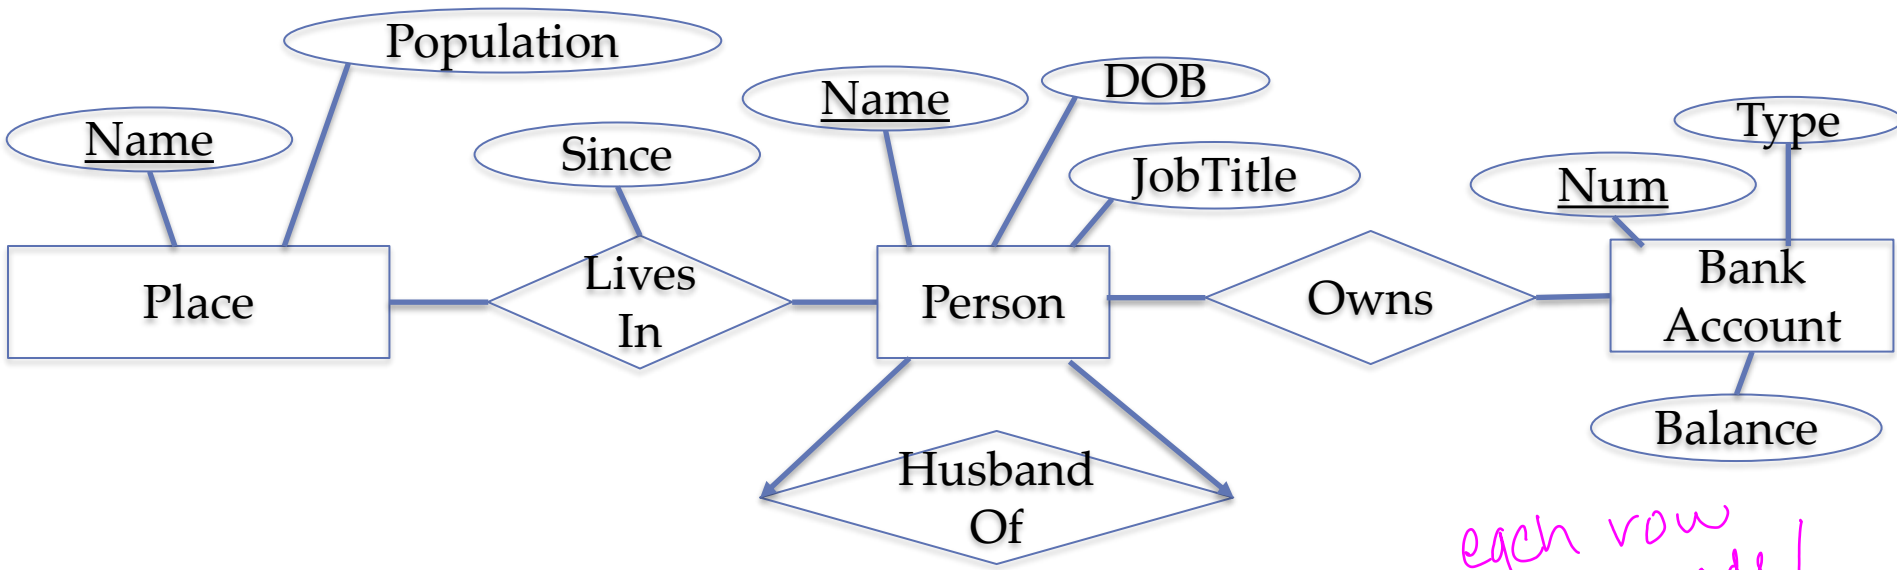
# Graph Data Model Instance Example



White board friendly !, Easy to understand.

nodes → entity sets  
edges → relationship sets

# ER / Relational



each row is a node / edge

Place

<u>Name</u>	Population
Venice	250,000

BankAccount

<u>Num</u>	Type	Balance
834233555	Savings	2300

Person

<u>Name</u>	DOB	jobTitle
John Doe	28-02-1988	NULL ←
Sally Doe	31-03-1989	Programmer

LivesIn

<u>PName</u>	<u>CityName</u>
John Doe	Venice

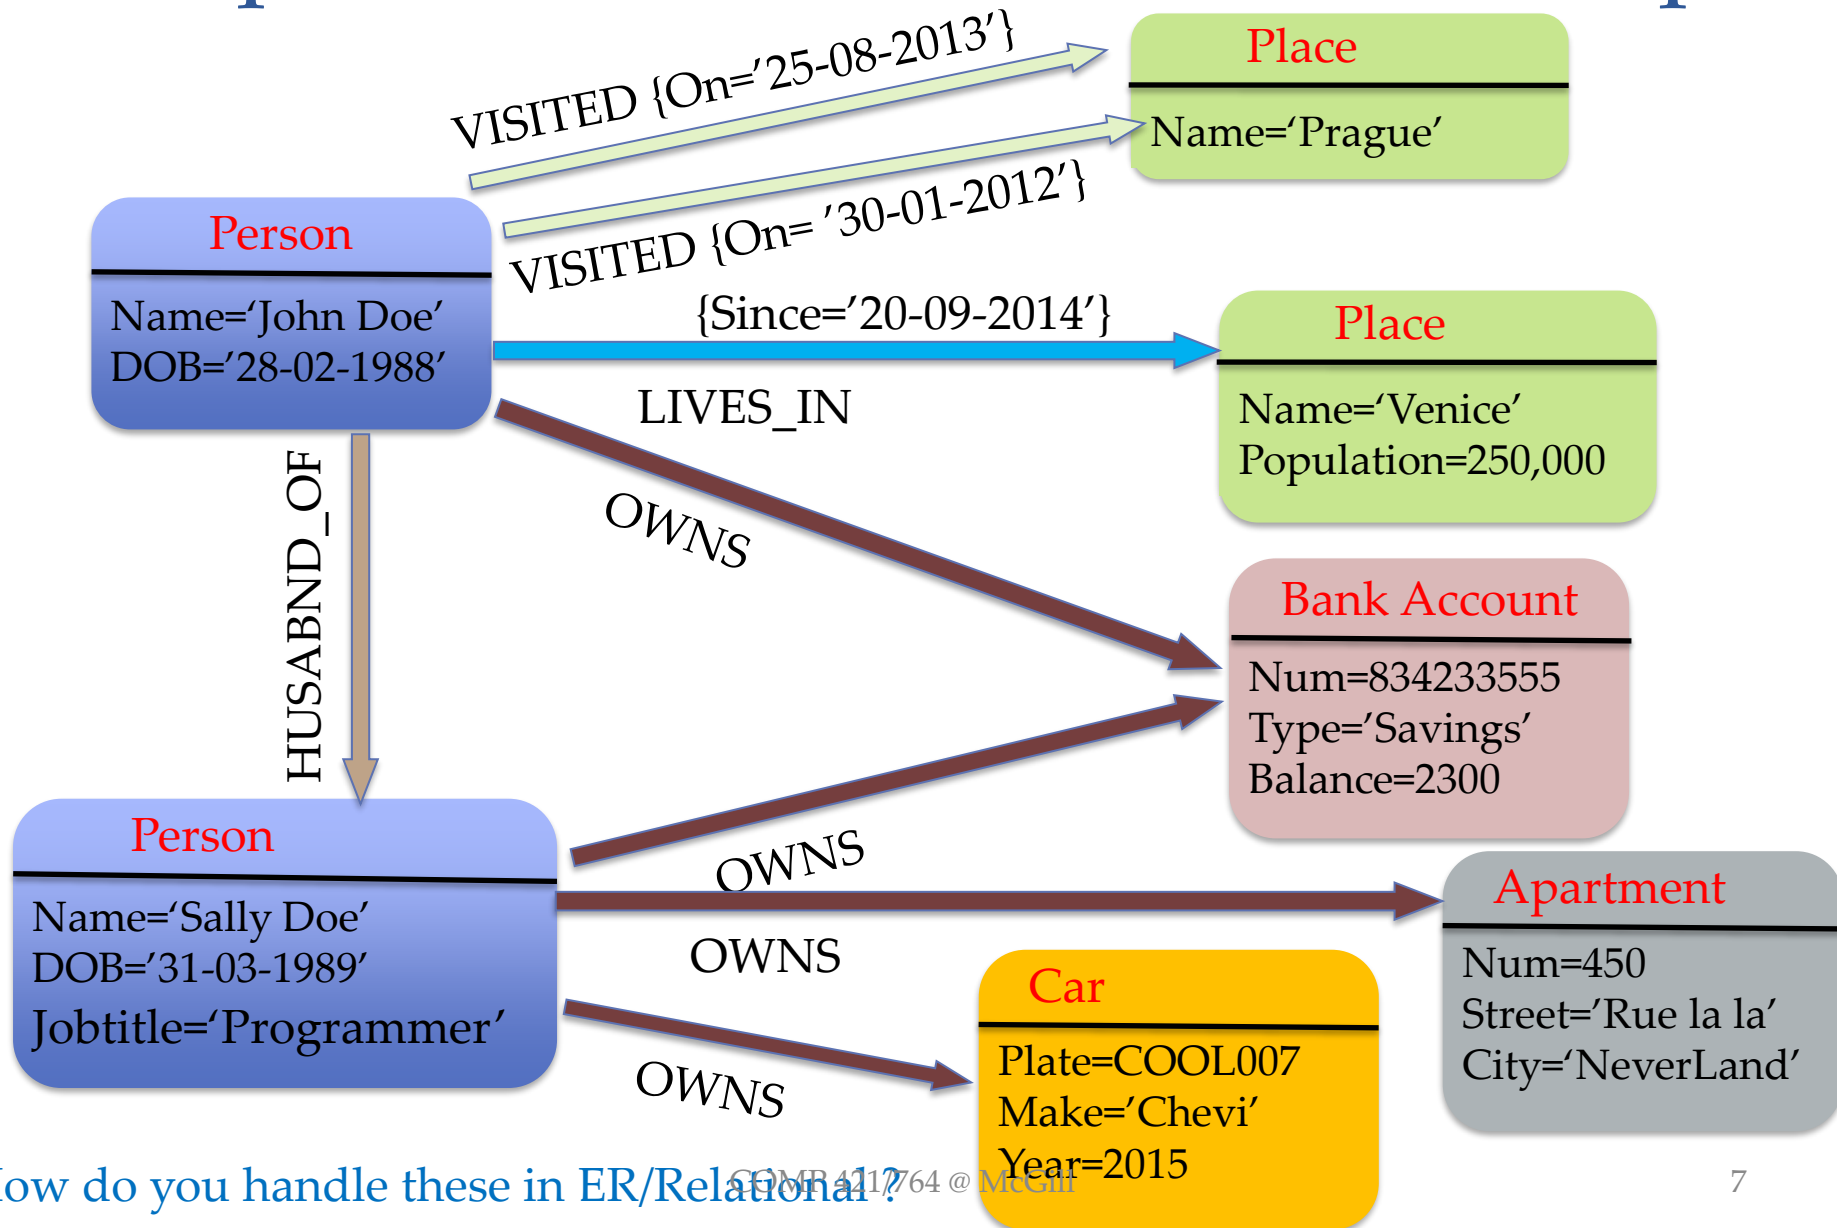
Owns

<u>PName</u>	<u>Num</u>
John Doe	834233555
Sally Doe	834233555

HusbandOf

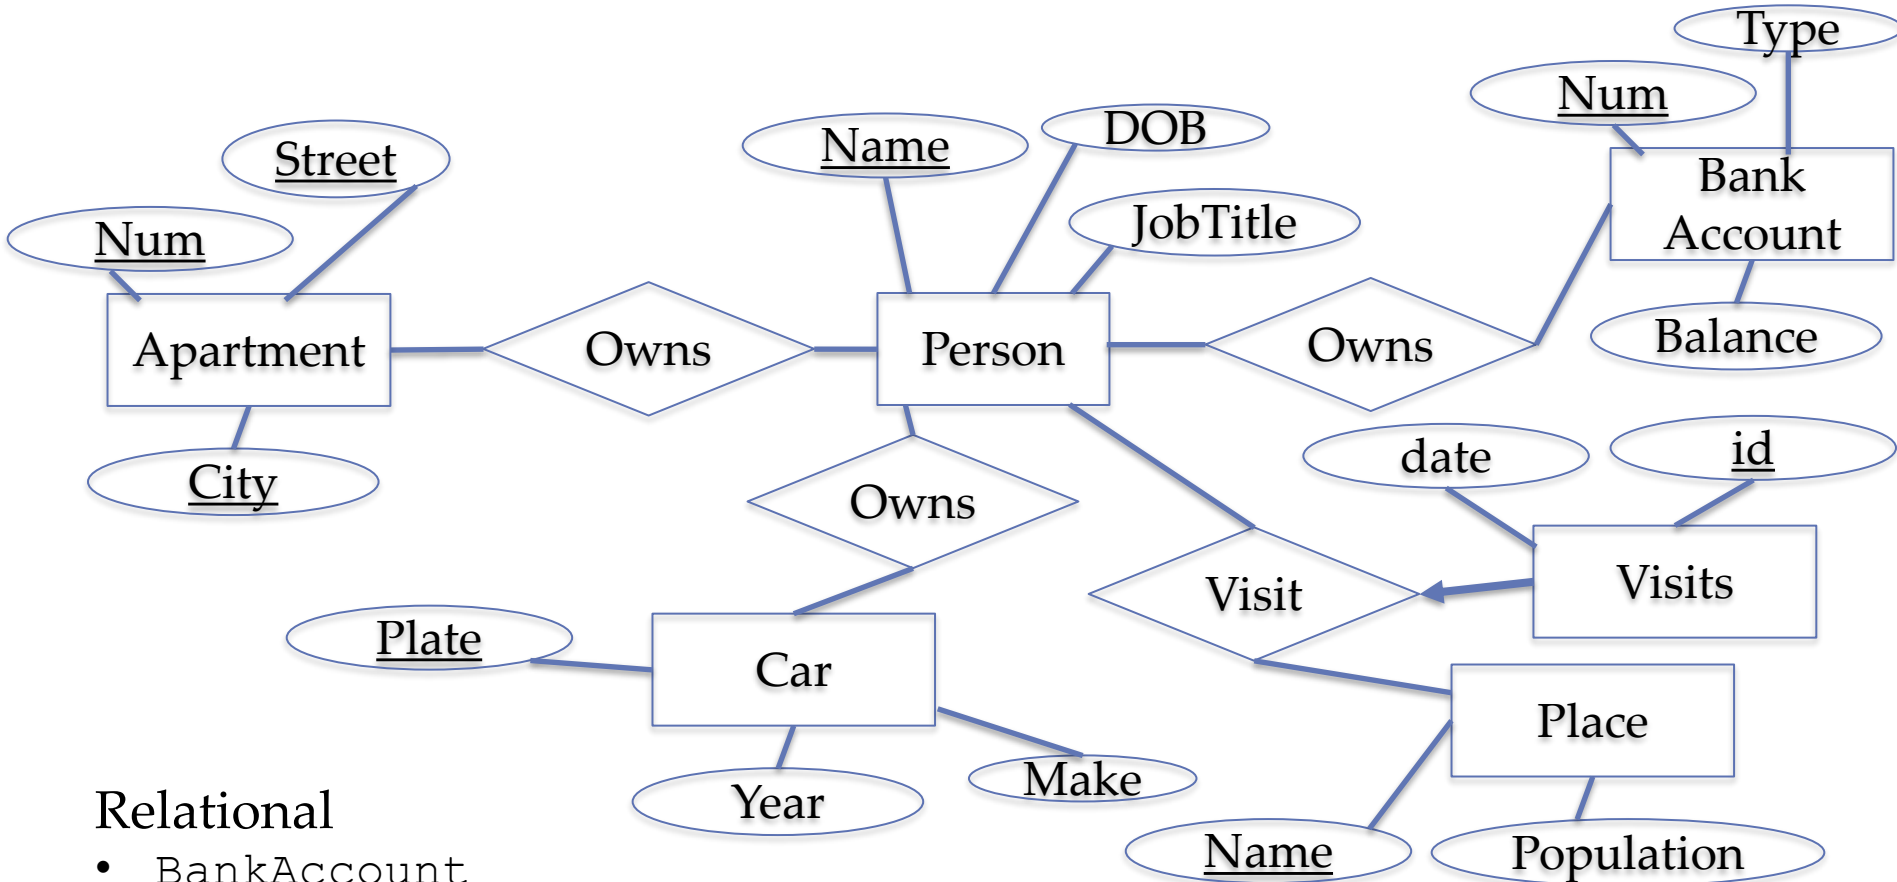
<u>HName</u>	<u>WName</u>
John Doe	Sally Doe

# Graph Data Model Instance Example



How do you handle these in ER/Relational?

# ER / Relational



## Relational

- BankAccount
- Apartment
- Car
- BankAccountOwnership
- ApartmentOwnership
- CarOwnership
- ...

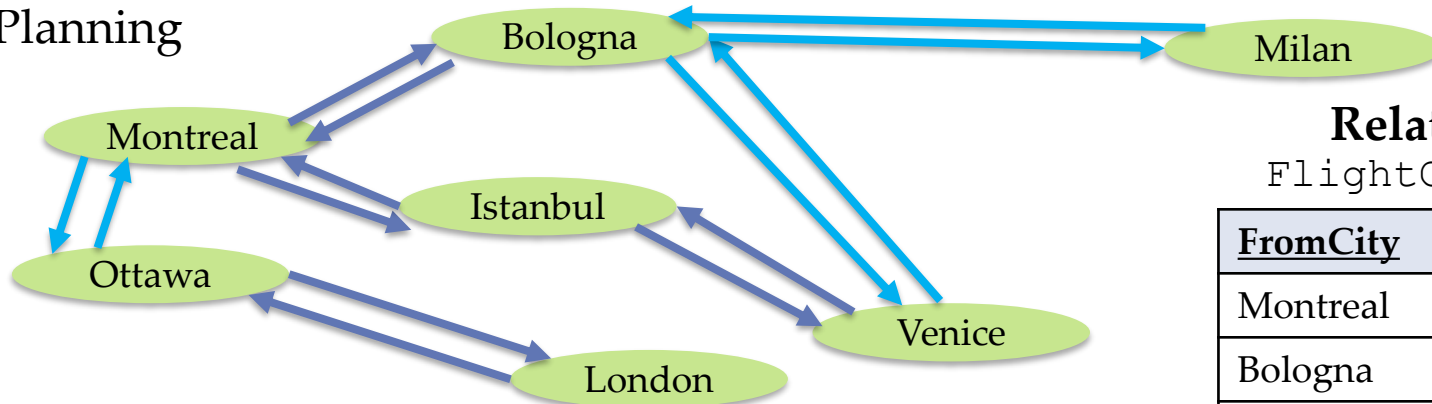
What does a SQL to retrieve all the things that Sally owns look like ?

Might have to write separate SQLs as the relations does not have compatible structures



# Working With Varying Depth in Relationships

## Trip Planning



To Get to Venice From Montreal

Direct Flight Connection ?

```
SELECT * FROM FlightConnected
WHERE FromCity = 'Montreal' AND ToCity='Venice'
```

One Stop Flight Connection ?

```
SELECT F1.FromCity, F1.ToCity, F2.ToCity
FROM FlightConnected F1, FlightConnected F2
WHERE F1.ToCity = F2.FromCity
AND F1.FromCity = 'Montreal' AND F2.ToCity='Venice'
```

- Two Stop Flights ?
  - One more self join
- Flight & Rail ?
  - Unions + Joins

## Relational

FlightConnected

FromCity	ToCity
Montreal	Bologna
Bologna	Montreal
Montreal	Istanbul
Istanbul	Montreal
Istanbul	Venice
....	....

RailConnected

FromCity	ToCity
Milan	Bologna
Bologna	Milan
Venice	Bologna
Bologna	Venice
....	....

lots of joins?  $\Rightarrow$  graph database

# Why Graph Databases

- Schema flexibility → new entity types, properties, relationships, etc. can emerge without significant impact on existing data model. I.e, the data model need not be completely developed ahead.
- Graph data model can accommodate the concept of relationship between entities a lot more efficiently than a relational model.

# Use Cases

- Social/Professional Networks
- Computer Networks
- Complex Hierarchies
- Geo-Spatial Data (Maps, Flight Reservation, etc.)
- Relationship Between Webpages ? (Web Graph)
- Maintain Knowledge Base
- Maintain IT Infrastructure
- Real-time Recommendation
- Fraud Detection


# How to interact with the Graph ?

- Custom Programming language APIs
- Query Languages
  - Cypher (Most Prominent and Open)
    - Simple, ASCII art type of queries.
  - Gremlin
  - SPARQL
  - XPath/XQuery (For XML based databases)
- Graph Query Language (GQL) – proposed standard
  - Characteristics from Cypher, SQL, etc.
- Neo4j <https://neo4j.com/download> [ Desktop Version ]
  - Programming language APIs in various host languages (Java, Ruby, Python ...)
  - Cypher (In this class...)

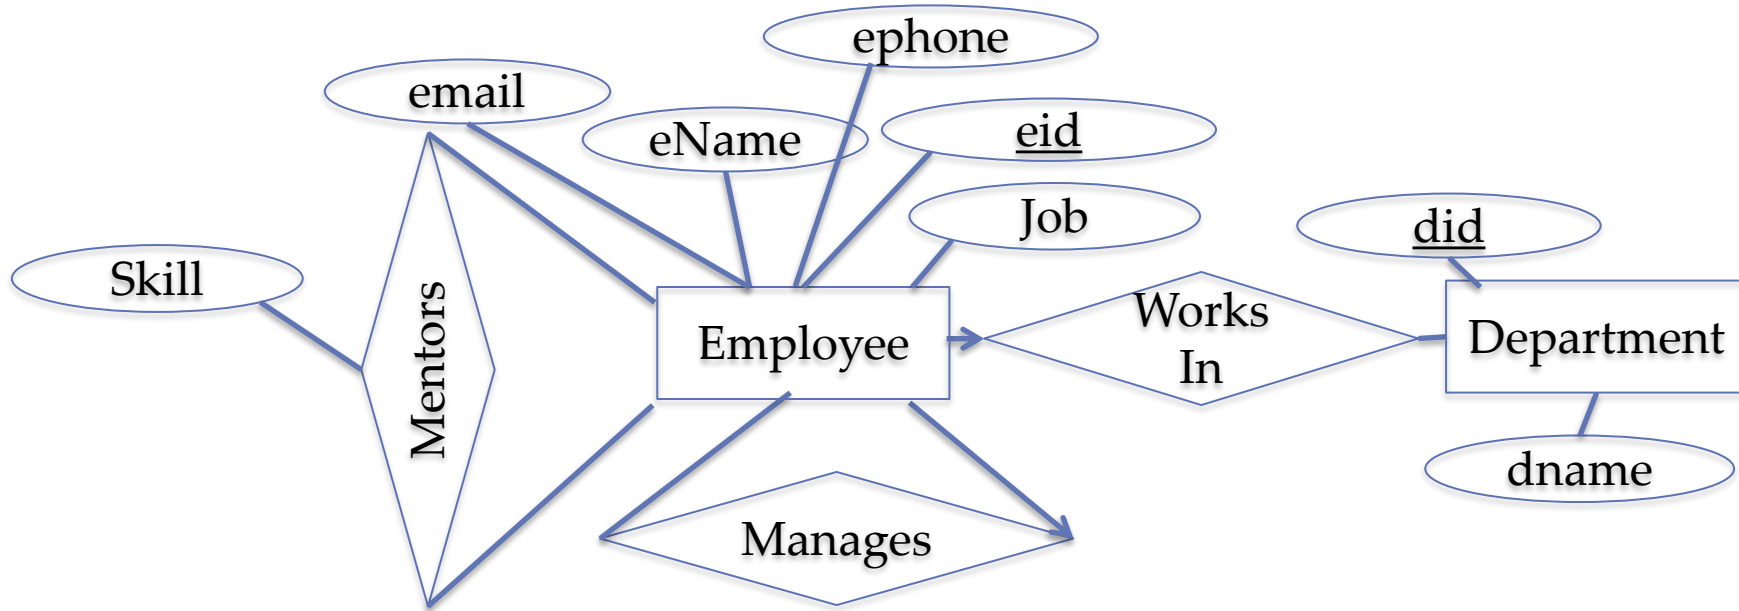
# Cypher

- Declarative language for Graph Database Systems, similar to SQL for Relational Database Systems.
  - Specifies what data to retrieve, not how to retrieve them (Similar concept to SQL).
  - Heavily influenced by SQL and SPARQL.
- Data Types
  - Standard data types
    - Integers, Floating point, Strings, Boolean
  - Extended Graph data types
    - Nodes, Relationships, Paths, Maps, Lists

# Sample Graph Database



# ER / Relational



## Relational

Department(did, dname)

Employee(eid, ename, ephone, email, job, deptid) deptid is FK to Department(did)

Manages(mgreid, empeid) both mgreid and empeid is FK to Employee(eid)

Mentorship(mentor eid, mentee eid) both mentor\_eid, mentee\_eid FK to employee



# "SELECT"



Return All Employee Records  
(all columns)

Return All Employee Nodes  
(all properties)

Variable names are required only if the node is being created or its contents needs to be referred elsewhere again. They are case sensitive

```
SELECT *  
FROM Employees
```

*type*  
MATCH(e:Employee)

RETURN e

*match any  
node of  
type employee*

Returns the entire node



# “PROJECTION”

SQL

```
SELECT email, ephone  
FROM Employees
```

Cypher

```
MATCH(e:Employee)  
RETURN e.email, e.ephone
```

Returns only specific  
fields

## ORDERING Output

```
SELECT email, ephone  
FROM Employees  
ORDER BY email
```

```
MATCH(e:Employee)  
RETURN e.email, e.ephone  
ORDER BY e.email
```

# “SELECT” with a Condition

SQL

```
SELECT *  
FROM Employees  
WHERE ename = 'Janet'
```

Find the Employee info of Janet

Cypher

```
MATCH(e:Employee)  
WHERE e.ename = 'Janet'  
RETURN e
```

or

```
MATCH(e:Employee {ename:'Janet'})  
RETURN e
```

Search for multiple employees at the same time

```
SELECT *  
FROM Employees  
WHERE ename IN ( 'Janet',  
'Steve' )
```

```
MATCH(e:Employee)  
WHERE e.ename IN [ 'Janet', 'Steve' ]  
RETURN e;
```

Pattern matching names starting with S

```
WHERE ename LIKE 'S%'
```

```
WHERE e.ename =~ 'S.*'
```

# "SELECT" with a Condition

SQL

Cypher

Find Employees without a department assigned

```
SELECT *  
FROM Employees  
WHERE deptid IS NULL
```

```
MATCH(e:Employee)  
WHERE NOT (e)-[:WORKS_IN]->()  
RETURN e
```

*employees with a works\_in edge to any other node*

*negated*



Find Employees who are not mentoring anyone

```
SELECT *  
FROM Employees  
WHERE eid NOT IN  
(SELECT mentor_eid FROM Mentorships)
```

```
MATCH(e:Employee)  
WHERE NOT (e)-[:MENTORS]->()  
RETURN e
```

Find Employees who are not mentored by anyone

```
SELECT *  
FROM Employees  
WHERE eid NOT IN  
(SELECT mentee_eid  
FROM Mentorships)
```

```
MATCH(e:Employee)  
WHERE NOT (e)<-[:MENTORS]-()  
RETURN e
```

*anyone mentors e*

*remove direction from edge*

# NULL



A non existent property in the node is treated as NULL

```
MATCH(e:Employee)
WHERE e.job IS NULL
RETURN e
```

# Operators (Not a Complete List)

General	DISTINCT
Math	+, -, *, /, %, ^
Comparison	=, <>, <, >, <=, >=, IS NULL, IS NOT NULL
String comparison	STARTS WITH, ENDS WITH, CONTAINS
Boolean	AND, OR, XOR, NOT
String operators	+ (Concatenation), =~ (regex matching)

# Modifying a Graph

- **CREATE / DELETE**
  - Create / Delete nodes/relationships
- **SET/REMOVE**
  - Set values to properties Add/Remove labels to nodes and relationships
- **MERGE**
  - Finding an existing node / create a new node

# "INSERT"

SQL

Cypher

INSERT INTO Department ('PR', 12)

INSERT INTO Employee

(201, 'Jane', '111-333-9999', 12)

*No restrictions*

CREATE (d:Department {dname:"PR", did:12}) <-[WORKS\_IN]-

(e:Employee {ename:"Jane", eid:201, ephone:"111-333-9999"})

OR

CREATE (e:Employee {ename:"Jane", eid:201, ephone:"111-333-9999"})

-[WORKS\_IN]-> (d:Department {dname:"PR", did:12})

INSERT  
Followed by  
SELECT

CREATE (e:Employee {ename:"Jane", eid:201, ephone:"111-333-9999"})

-[r:WORKS\_IN]-> (d:Department {dname:"PR", did:12})

RETURN e,d,r

# “INSERT”



Add a new relationship between two existing nodes.

```
MATCH(n1:Employee {eid:101}), (n2:Employee {eid:201})  
CREATE (n1) -[:MANAGES]-> (n2);
```

Inserts can have multiple nodes of same or different types *(no restrictions)*

```
CREATE  
(n1:Employee {ename:'Janet', eid:101, email:'ja@comp.com', ephone:'123-456-1111', job:'CEO'})  
,(n2:Employee {ename:'Steve', eid:102, email:'st@comp.com', ephone:'123-456-1112', job:'VP IT'})  
,(n3:Employee {ename:'Louis', eid:103, email:'lo@comp.com', ephone:'123-456-1113', job:'VP HR'})  
, (d1:Department {dname:'IT', did:10})  
,(d2:Department {dname:'HR', did:11});
```



# DELETE

## SQL

Delete the records from the referencing table first, followed by the records from the referenced table.

```
DELETE FROM Employee WHERE empid = 201;  
DELETE FROM Department WHERE deptid = 12;
```

```
MATCH(e:Employee{empid:201})-[r:WORKS_IN]->(d:Department{deptid:12})  
DELETE e,d,r
```

Delete Only the Employee

```
DELETE FROM Employee WHERE empid = 201;
```

Deletes any edges as well as the Node

```
MATCH(e:Employee{empid:201})
```



```
DETACH DELETE e
```

Could be a better way of deleting a node, as we do not have to know about all the edges that involves this node

## Cypher

Delete the relationships (edges) interacting with that node before/while deleting the nodes themselves.

# “JOINS” / Traversals

Find Employees working in HR department.

```
SELECT e.*  
FROM Employees e, Dept d  
WHERE e.deptid = d.deptid  
      AND dname = 'HR'
```

```
SELECT e.*  
FROM Employees e  
      INNER JOIN Dept d  
      ON e.deptid = d.deptid  
WHERE dname = 'HR'
```

```
MATCH(e:Employee) -[w:WORKS_IN]->  
(d:Department{dname:"HR"})  
RETURN e
```

```
MATCH(e:Employee) -[w:WORKS_IN]->  
(d:Department)  
WHERE d.dname = "HR"  
RETURN e
```

# Traversals in Various Forms



Returns information on Paul and all the nodes he has relationships with.

MATCH(e:Employee) -- (n) WHERE e.ename = 'Paul' RETURN e,n

MATCH(e:Employee) -[]- (n) WHERE e.ename = 'Paul' RETURN e,n

Returns information on Paul and all the outgoing relationships he has along with the related nodes

MATCH(e:Employee) -[r]-> (n) WHERE e.ename = 'Paul' RETURN e,r,n

Returns information names of all the employees that Steve is managing.

MATCH(e:Employee) -[:MANAGES]-> (n:Employee)

WHERE e.ename = 'Steve'

RETURN n.ename

*-[]- {any edge  
any type, any direction}*

*any type*

*need variable name to return*

# Traversals in Various Forms

## Cypher

Returns information of people who works for or is mentored by Katie

```
MATCH(e:Employee) -[:MANAGES|MENTORS]-> (n)
WHERE e.ename = 'Katie'
RETURN n
```

*Handwritten note: A pink circle highlights the pipe character in the pattern, with the word "or" written below it.*

Returns information of people who neither works for nor is mentored by Katie

```
MATCH(e:Employee), (n:Employee)
WHERE e.ename = 'Katie' AND NOT (e) -[:MANAGES|MENTORS]-> (n)
RETURN n;
```

*Handwritten notes: A pink circle highlights the word "NOT" in the query. A curved arrow points from this circle to the pipe character in the pattern, with the text "not + or = nor" written below. Another pink circle highlights the pipe character in the pattern.*

→ NOT goes in where clause only

# Traversals in Various Forms

## Cypher

Returns information of people who is reporting to managers who themselves report to Steve

```
MATCH(e:Employee) -[:MANAGES]->( ) -[:MANAGES]->(n)
```

```
WHERE e.ename = 'Steve'
```

```
RETURN n
```

```
MATCH(e:Employee) -[:MANAGES*2]->(n)
```

```
WHERE e.ename = 'Steve'
```

```
RETURN n
```

# Traversals in Various Forms



<code>(e)-[*]-&gt;(n)</code>	<code>// All the way (outgoing edges)</code>
<code>(e)-[*..5]-&gt;(n)</code>	<code>// Up to a depth of 5 edges (outgoing)</code>
<code>(e)-[*3..]-&gt;(n)</code>	<code>// 3 or more edges (outgoing)</code>
<code>(e)-[*3..5]-&gt;(n)</code>	<code>// 3 to 5 edges (outgoing)</code>
<code>(e)&lt;-[*3..5]-(n)</code>	<code>// 3 to 5 edges (incoming)</code>
<code>(e)-[*3..5]-(n)</code>	<code>// 3 to 5 edges (incoming or outgoing)</code>

# Traversals in Various Forms

## Cypher

Employees under Steve who are not managers.

```
MATCH(e:Employee) -[:MANAGES]-> (e2:Employee)
WHERE e.ename = 'Steve' AND NOT (e2)-[:MANAGES]->()
RETURN e2;
```

*Steve manages someone  
and that someone does  
not manage anyone*

Employees who are mentored by a direct report of Steve.

```
MATCH(e:Employee) -[:MANAGES]->() -[:MENTORS]->(n)
WHERE e.ename = 'Steve'
RETURN n
```

*↑  
person managed  
by Steve*

# Traversals in Various Forms

## Cypher

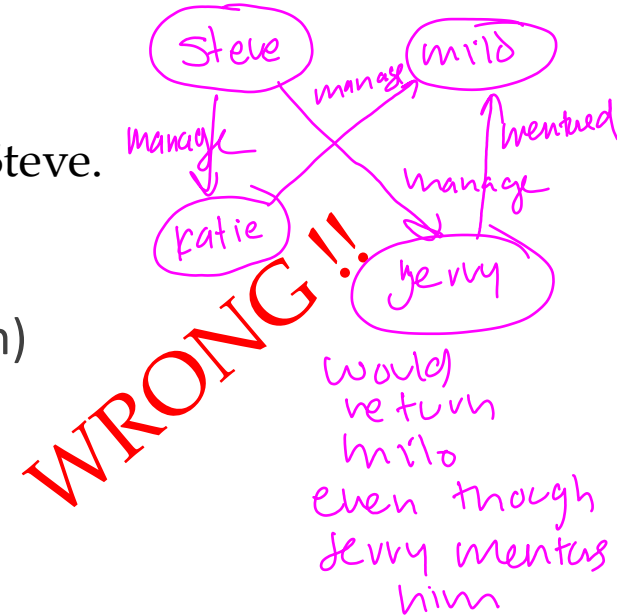
Employees who are NOT mentored by a direct report of Steve.

```
MATCH(e:Employee) -[:MANAGES]->(dr), (n:Employee)
WHERE e.ename = 'Steve' AND NOT (dr)-[:MENTORS]->(n)
RETURN n;
```

## Correct approaches

```
MATCH(n:Employee)
WHERE NOT (:Employee{ename:'Steve'})-[:MANAGES]->()-[:MENTORS]->(n)
RETURN n;
```

```
MATCH(e:Employee), (n:Employee)
WHERE e.ename = 'Steve' AND NOT (e)-[:MANAGES]->()-[:MENTORS]->(n)
RETURN n;
```





# Traversals in Various Forms



Important !!

For a given path output of a pattern, each edge is traversed only once !

For example if (Janet)-[:FRIEND\_OF]->(Sue) and (Sue)-[:FRIEND\_OF]->(Janet)

```
MATCH (e1:Employee(ename:'Janet')-[:FRIEND_OF*]->(e2:Employee)
RETURN e1,e2
```

Will return only two paths.

(Janet)->(Sue)

(Janet)->(Sue)->(Janet)

— back to the beginning → stops  
other friends of sue won't be returned

# Traversals in Various Forms



How to find if Janet is Managing an employee, who manages an employee who is friends with Janet ?

```
MATCH
(n:Employee{ename:'Janet'})-[:MANAGES]->(e1:Employee)-[:MANAGES]->(e2:Employee)
, (n)-[:FRIEND_OF]->(e2)
RETURN n,e1, e2
```

How to find a list of people who manages someone who mentors more than one employee ?

```
MATCH (b:Employee)-[:MANAGES]->(m:Employee)
      ,(m)-[:MENTORS]->(e1:Employee)
      ,(m)-[:MENTORS]->(e2:Employee)
WHERE e1 <> e2
RETURN DISTINCT b
```

<> not eval

# Other Aspects

- Constraints

- CREATE CONSTRAINT ON (e: Employee) ASSERT e.eid IS UNIQUE
- CREATE CONSTRAINT ON (e: Employee) ASSERT exists(e.ename) *like not null*
- CREATE CONSTRAINT ON ()-[m:MENTORS]-() ASSERT exists(m.skill)

# Other Aspects

- Transactions
- Dynamic property matching
- Multiple labels on the same node (not for relationships).
- Aggregation
- WITH clause
- Multiple MATCH clauses in a single statement

# Other Resources

- Download Neo4j ( Desktop )
  - <https://neo4j.com/download>
- Cypher Query Language
  - <http://neo4j.com/docs/developer-manual/current/cypher/>
- Neo4j webinar videos (If you get addicted to graph databases)
  - [https://www.youtube.com/channel/UCvze3hU6OZBk\\_B1vkhH2IH9Q](https://www.youtube.com/channel/UCvze3hU6OZBk_B1vkhH2IH9Q)